

# Architecting Stream Processing Solutions Using Google Cloud Pub/Sub

---

GETTING STARTED WITH CLOUD PUB/SUB



**Vitthal Srinivasan**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Messaging middleware**

**Decouples senders and receivers**

**Global, replicated, autoscaling**

**Publishers publish to topics**

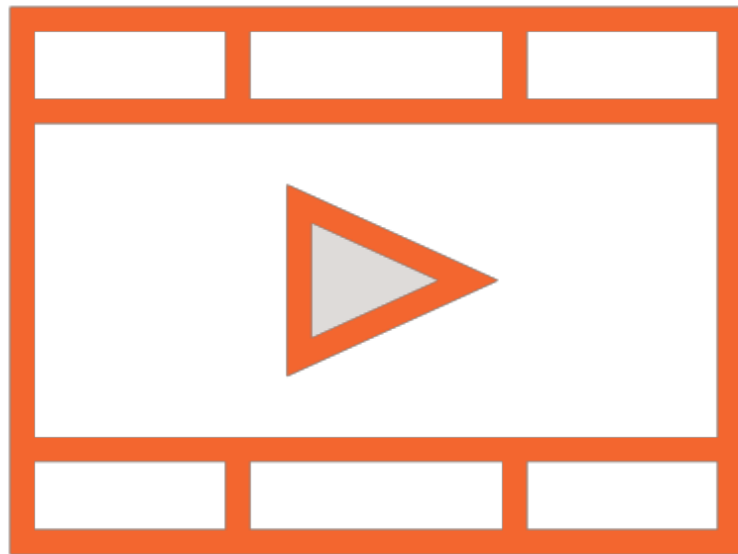
**Subscribers listen on subscriptions**

**Complex pricing**

# Prerequisites and Course Outline

---

# Prerequisites: Basic Cloud Computing



**Choosing and Implementing Google  
Cloud Compute Solutions**

**Architecting Google Cloud Storage  
Configurations**

# Course Outline



## Getting started

- Terminology: publishers, subscribers, messages
- Concepts and message flow
- Architectural overview

## Working with Pub/Sub

- Publishers, topics, subscribers, subscriptions
- Web console and gcloud
- Syncing to timestamp and snapshots

## Using client libraries

- Python client libraries for publishers and subscribers
- Async callbacks, batching and flow control
- Creating a Hangouts bot using Pub/Sub

# Scenarios: SpikeySales.com



## **Hypothetical Online Retailer**

- Flash sales of trending products
- Spikes in user traffic

## **SpikeySales on the GCP**

- Cloud computing fits perfectly
- Pay-as-you-go
- No idle capacity during off-sale periods

Pub/Sub


---

# Cloud Pub/Sub

Simple, global, reliable real-time messaging and streaming ingestion service on the GCP.



# Pub/Sub



**Many-to-many asynchronous messaging**


**Decouples senders and receivers**

**Reliable, scalable, secure**

**At-least-once delivery**

**Exactly-once processing (in Dataflow)**

# Pub/Sub



**Serverless**


**Autoscaling along all axes**

- Senders
- Receivers
- Number and size of messages

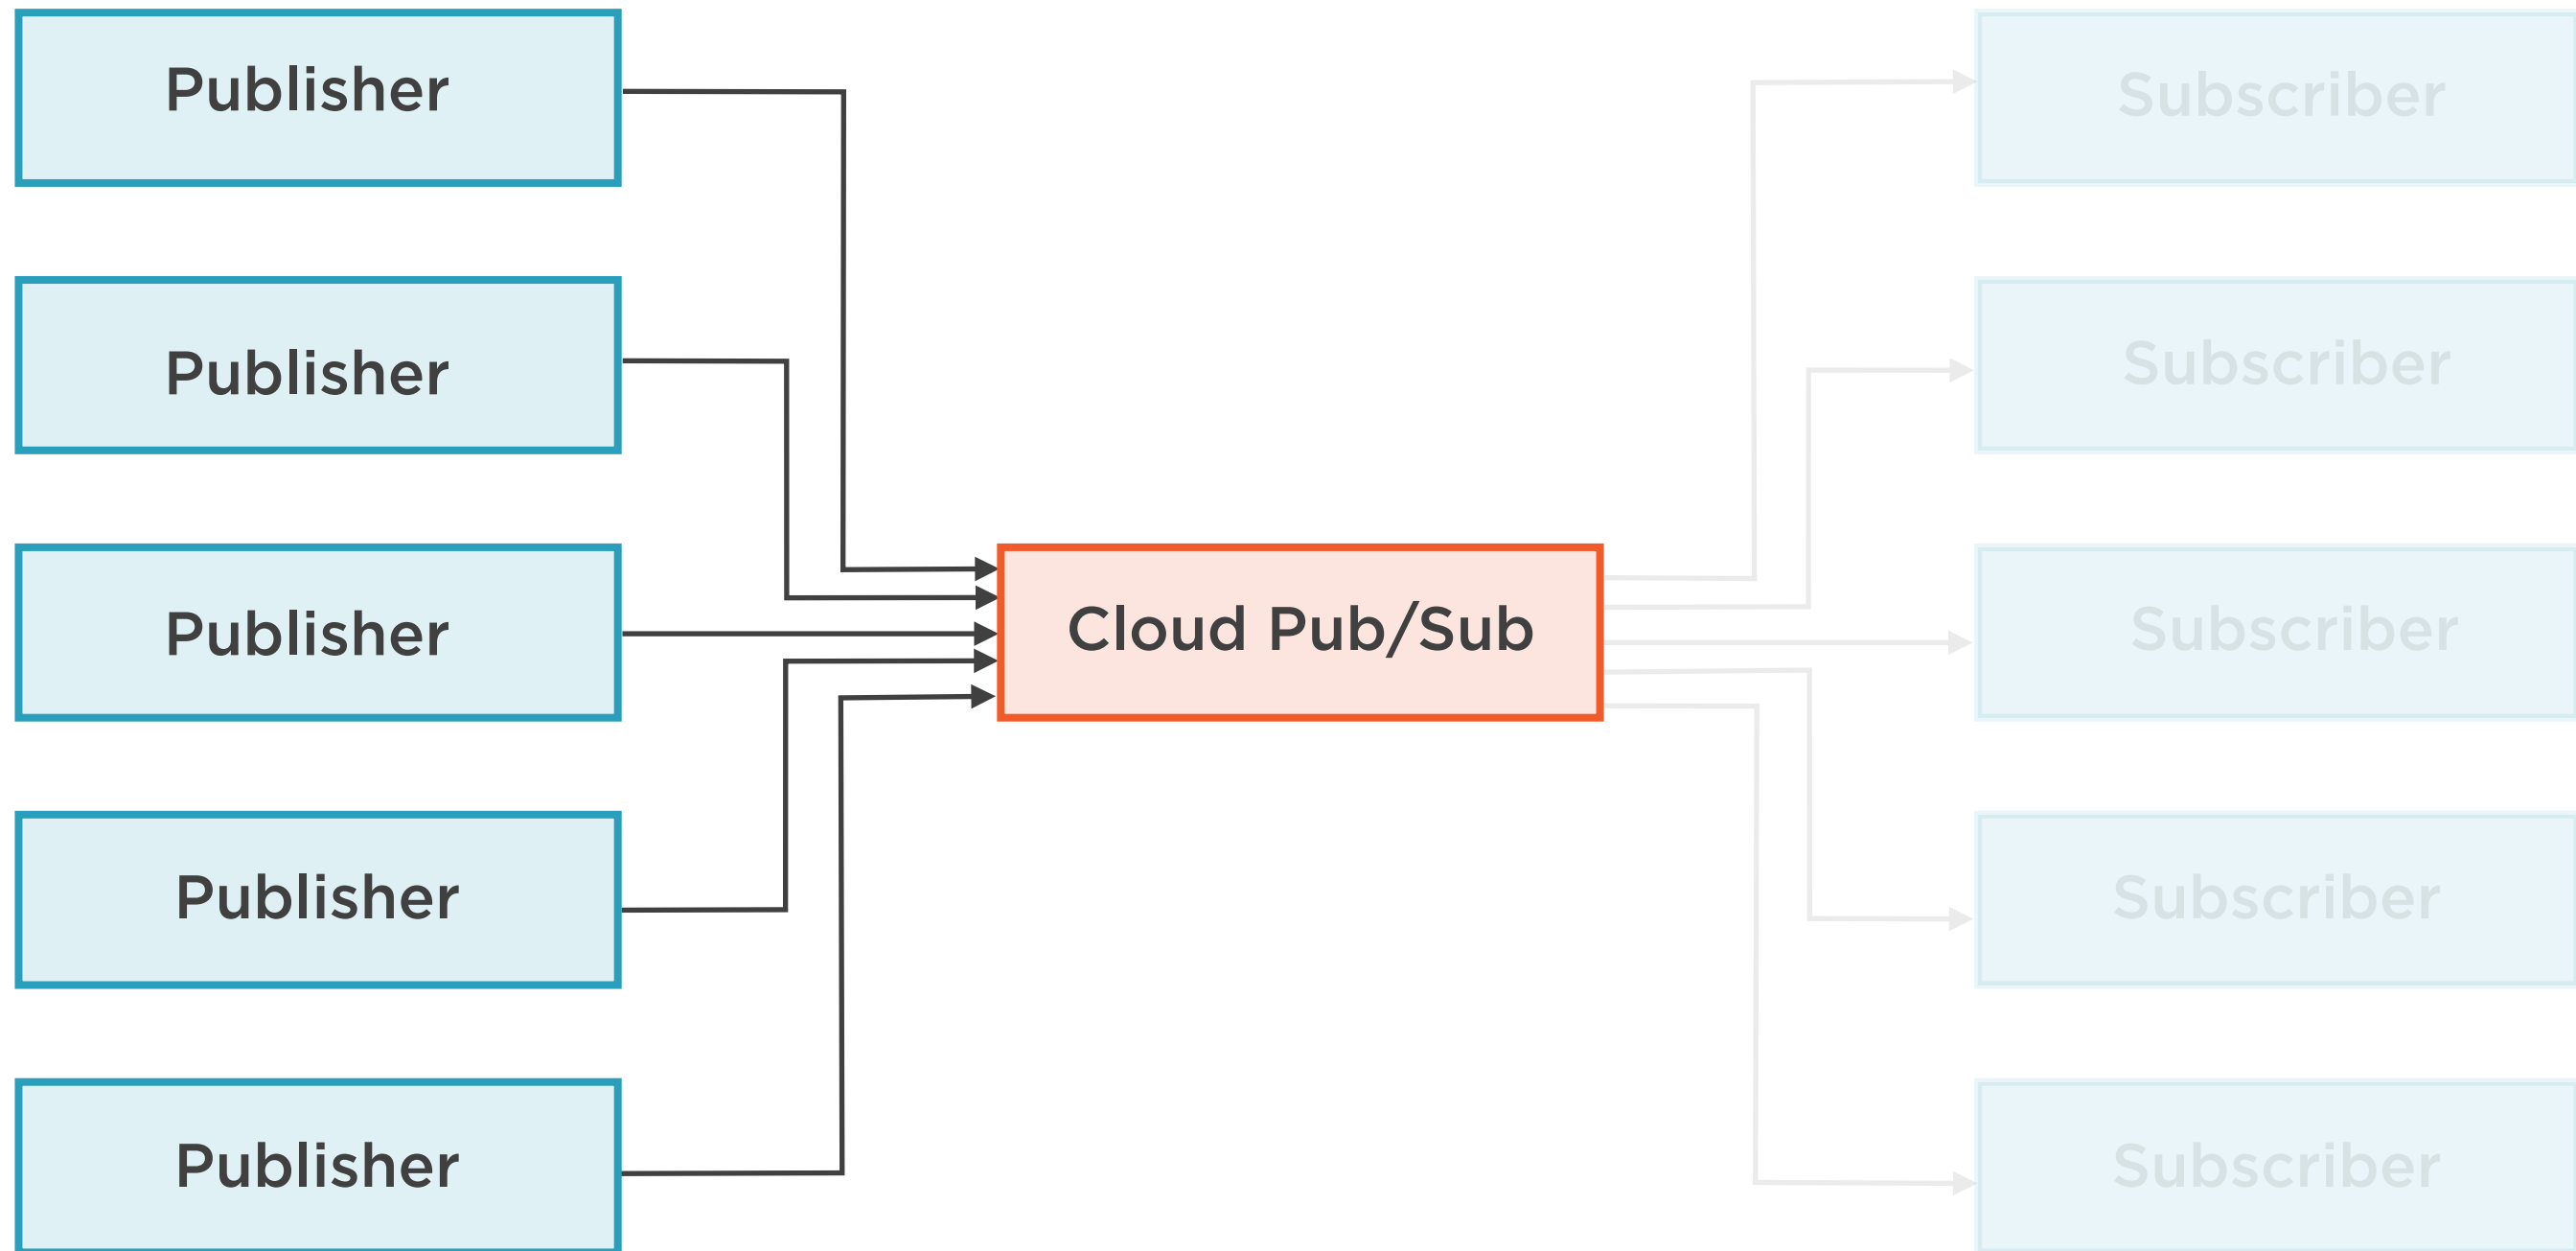
**Replicated**

**Load-balanced internally**

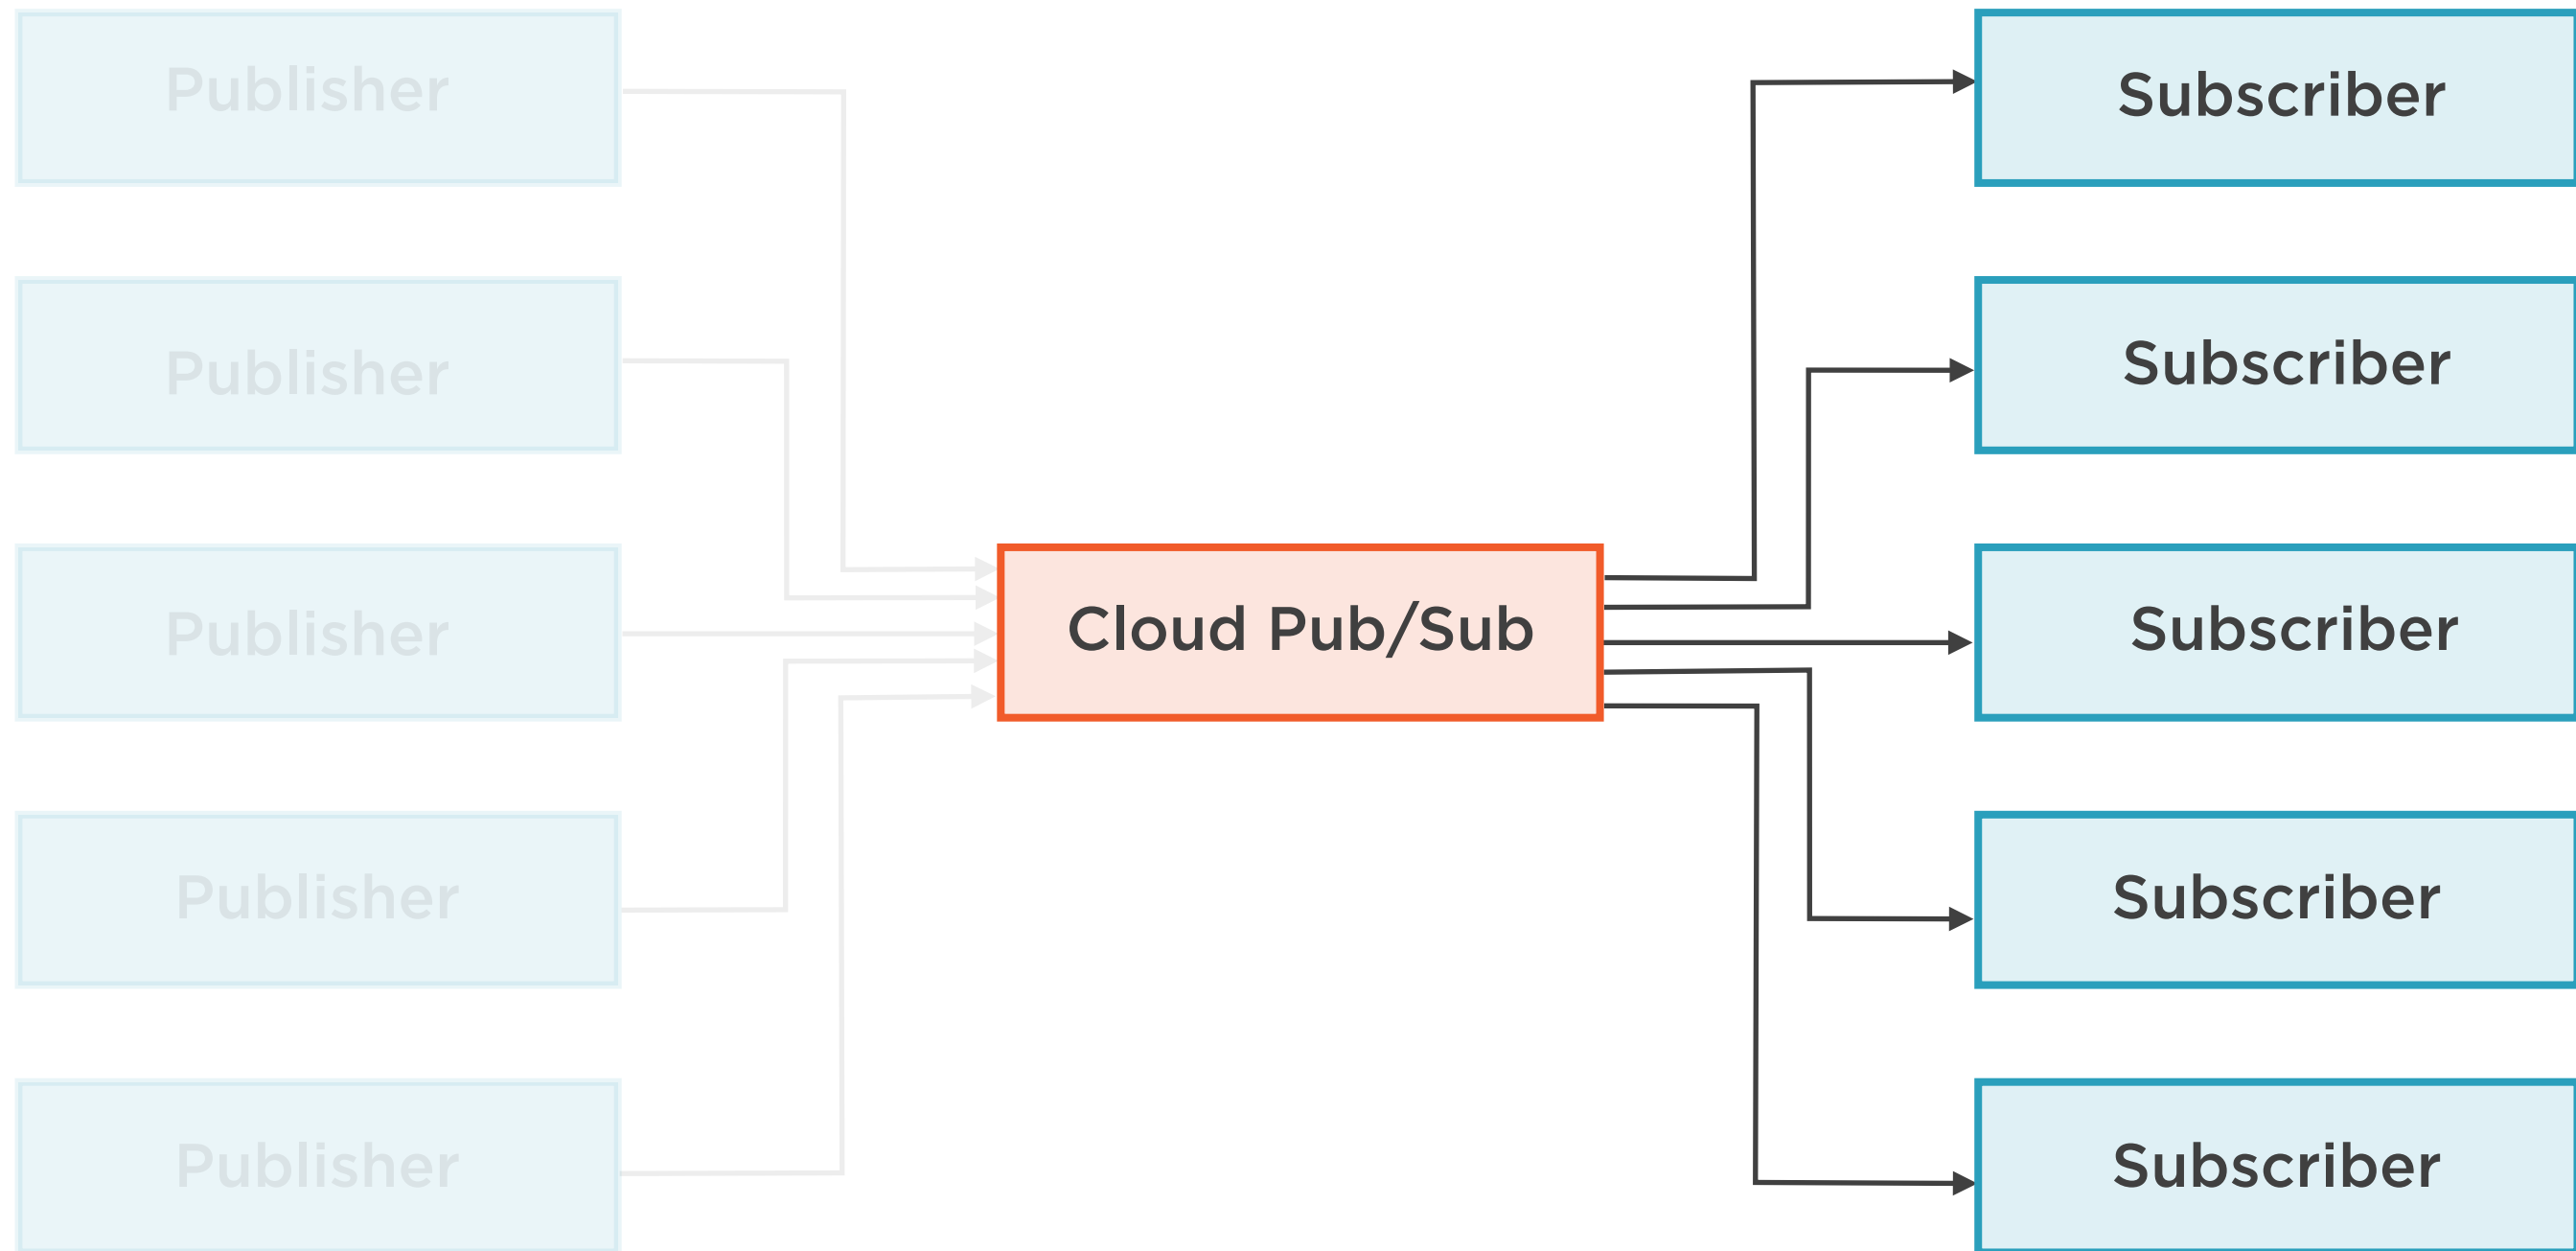
# Messaging Middleware



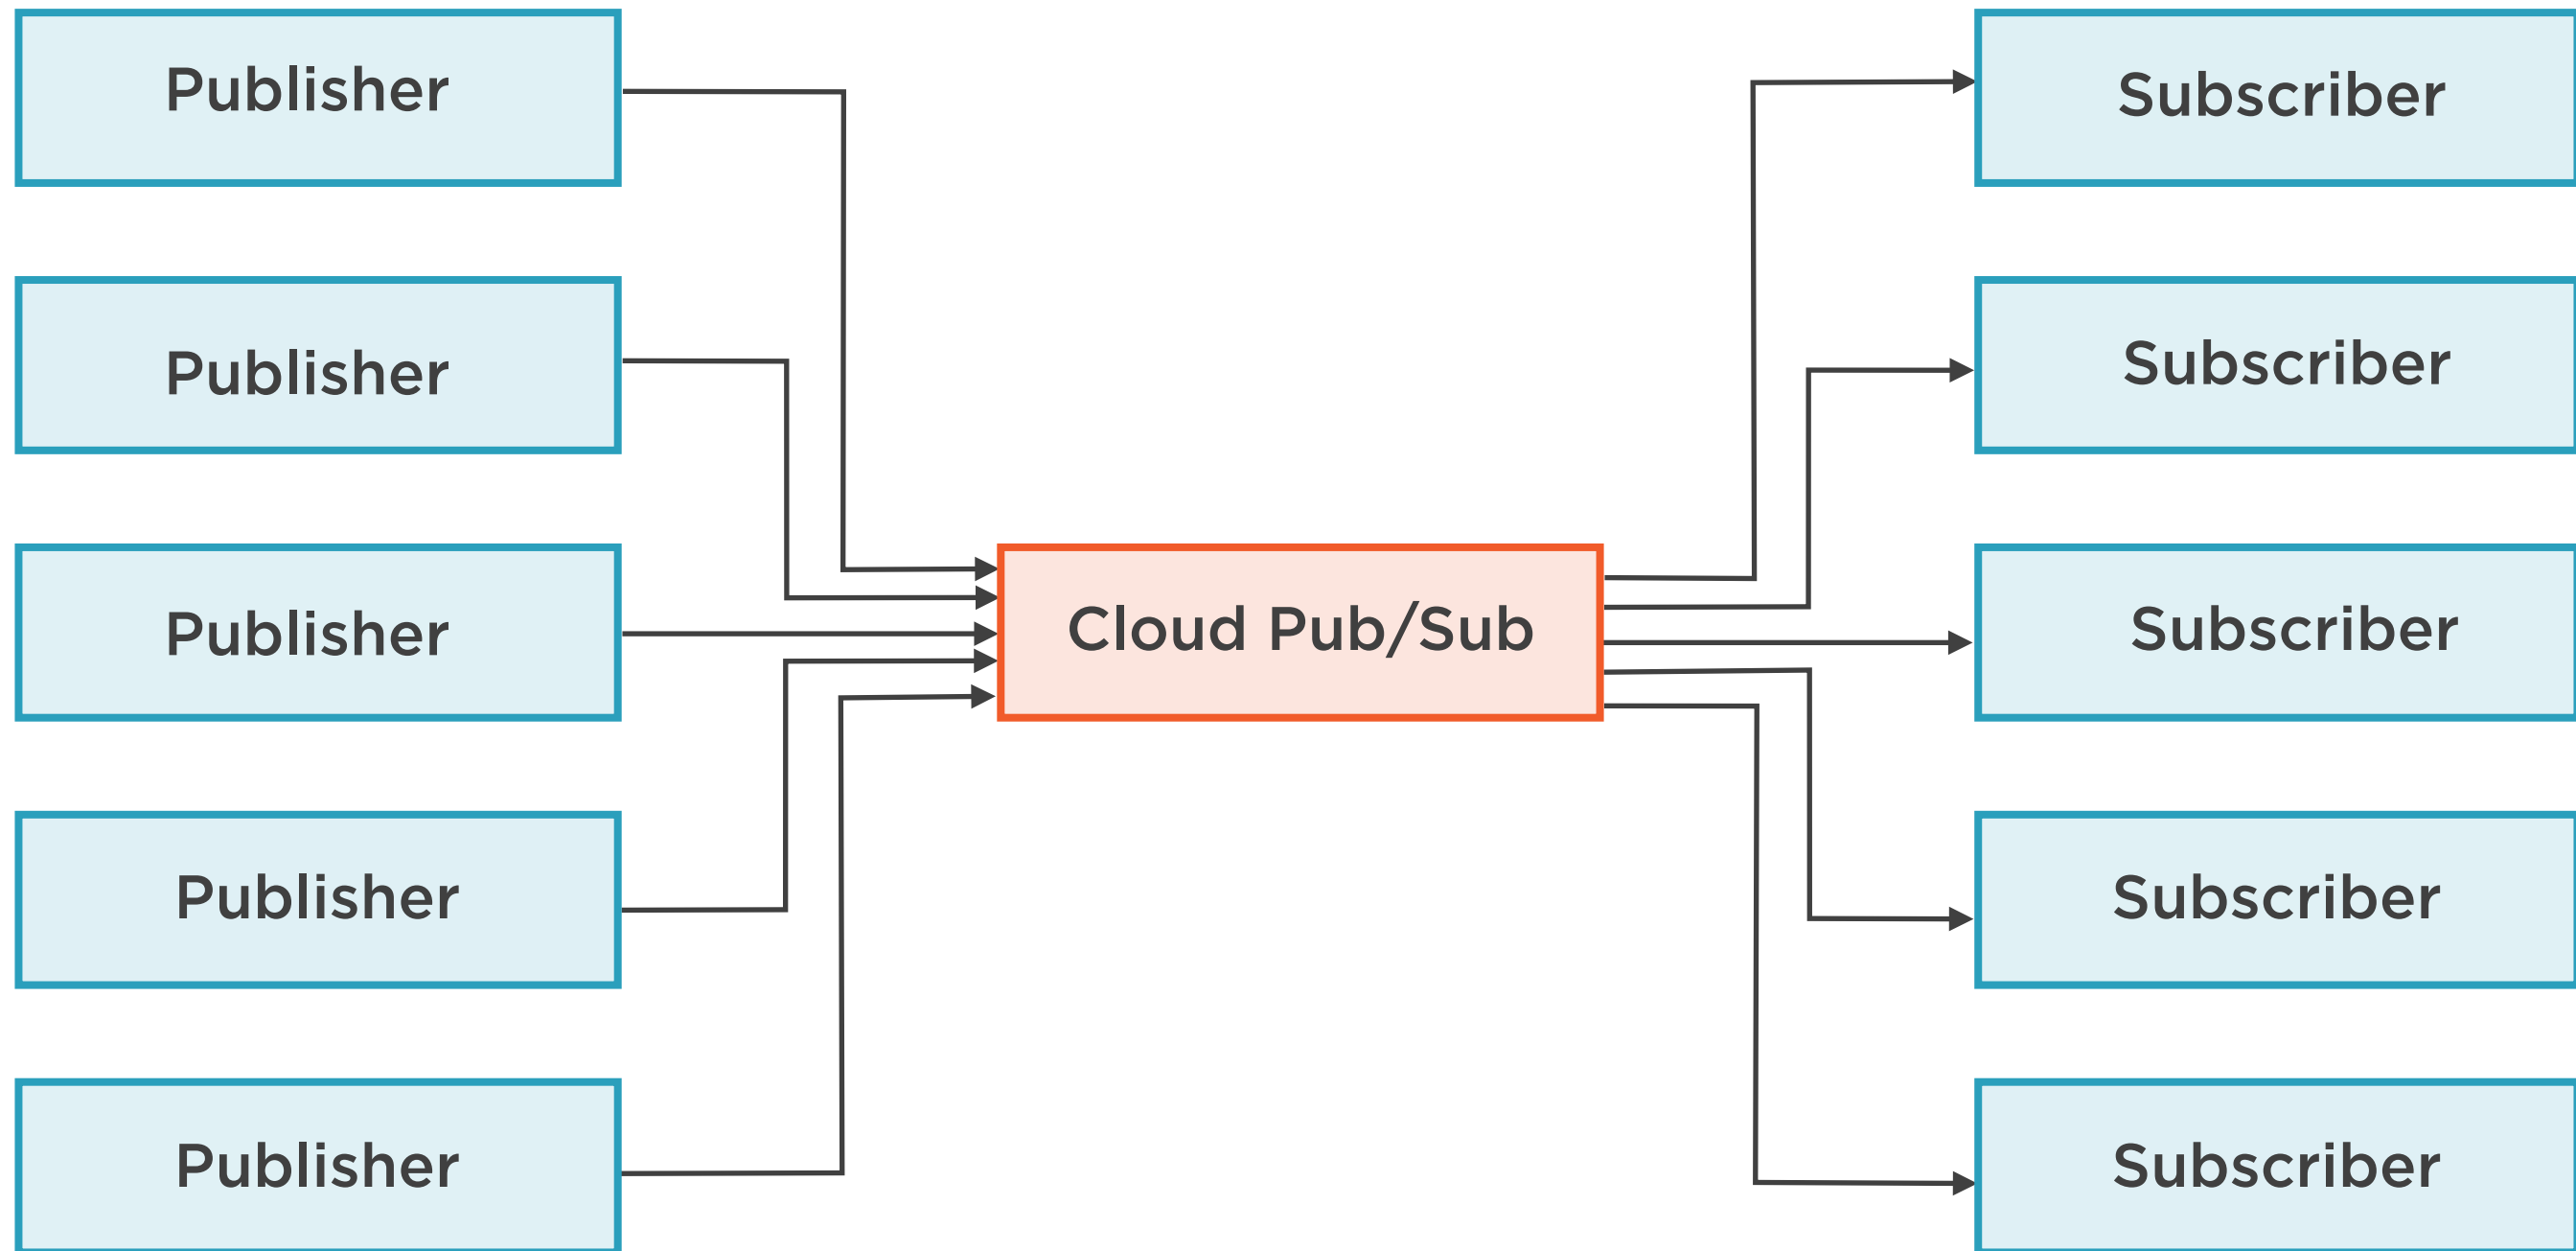
# Messaging Middleware



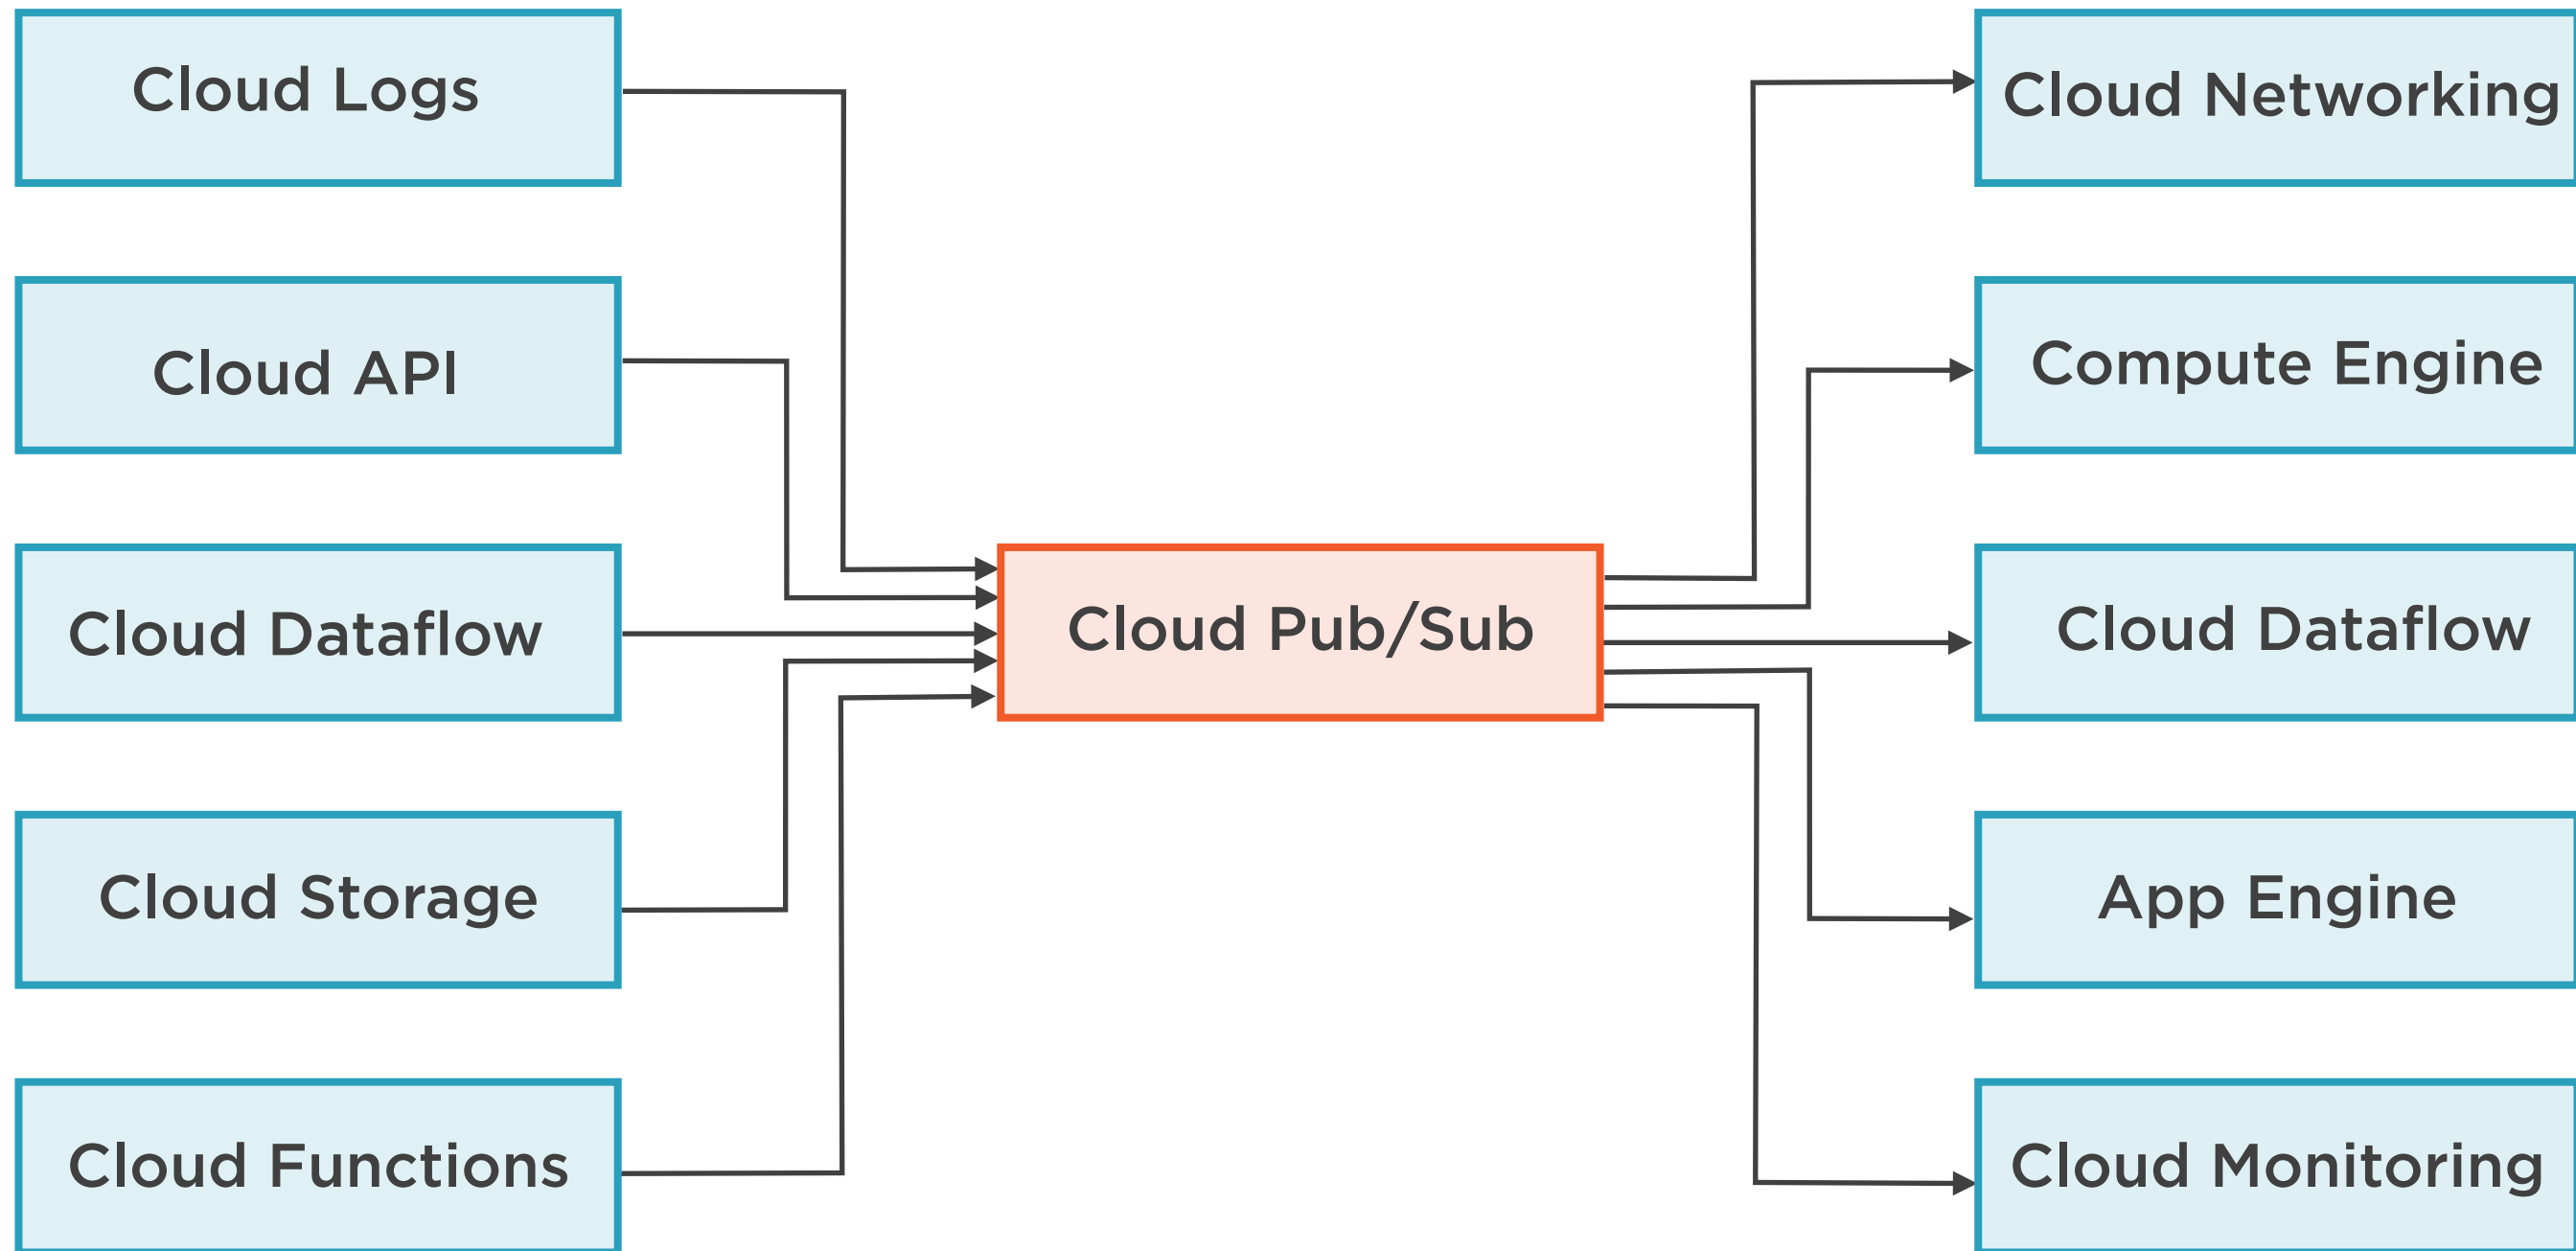
# Messaging Middleware




# Messaging Middleware



# Works with All GCP Services



# Pub/Sub




**Publisher** apps create and send messages on a **Topic**

**Subscriber** apps subscribe to a topic to receive messages


**Subscription** is a queue (message stream) to a subscriber




# Publishers, Topics, Subscribers, Subscriptions




# Publishers Publish to a Topic




# A Topic Can Have Multiple Publishers



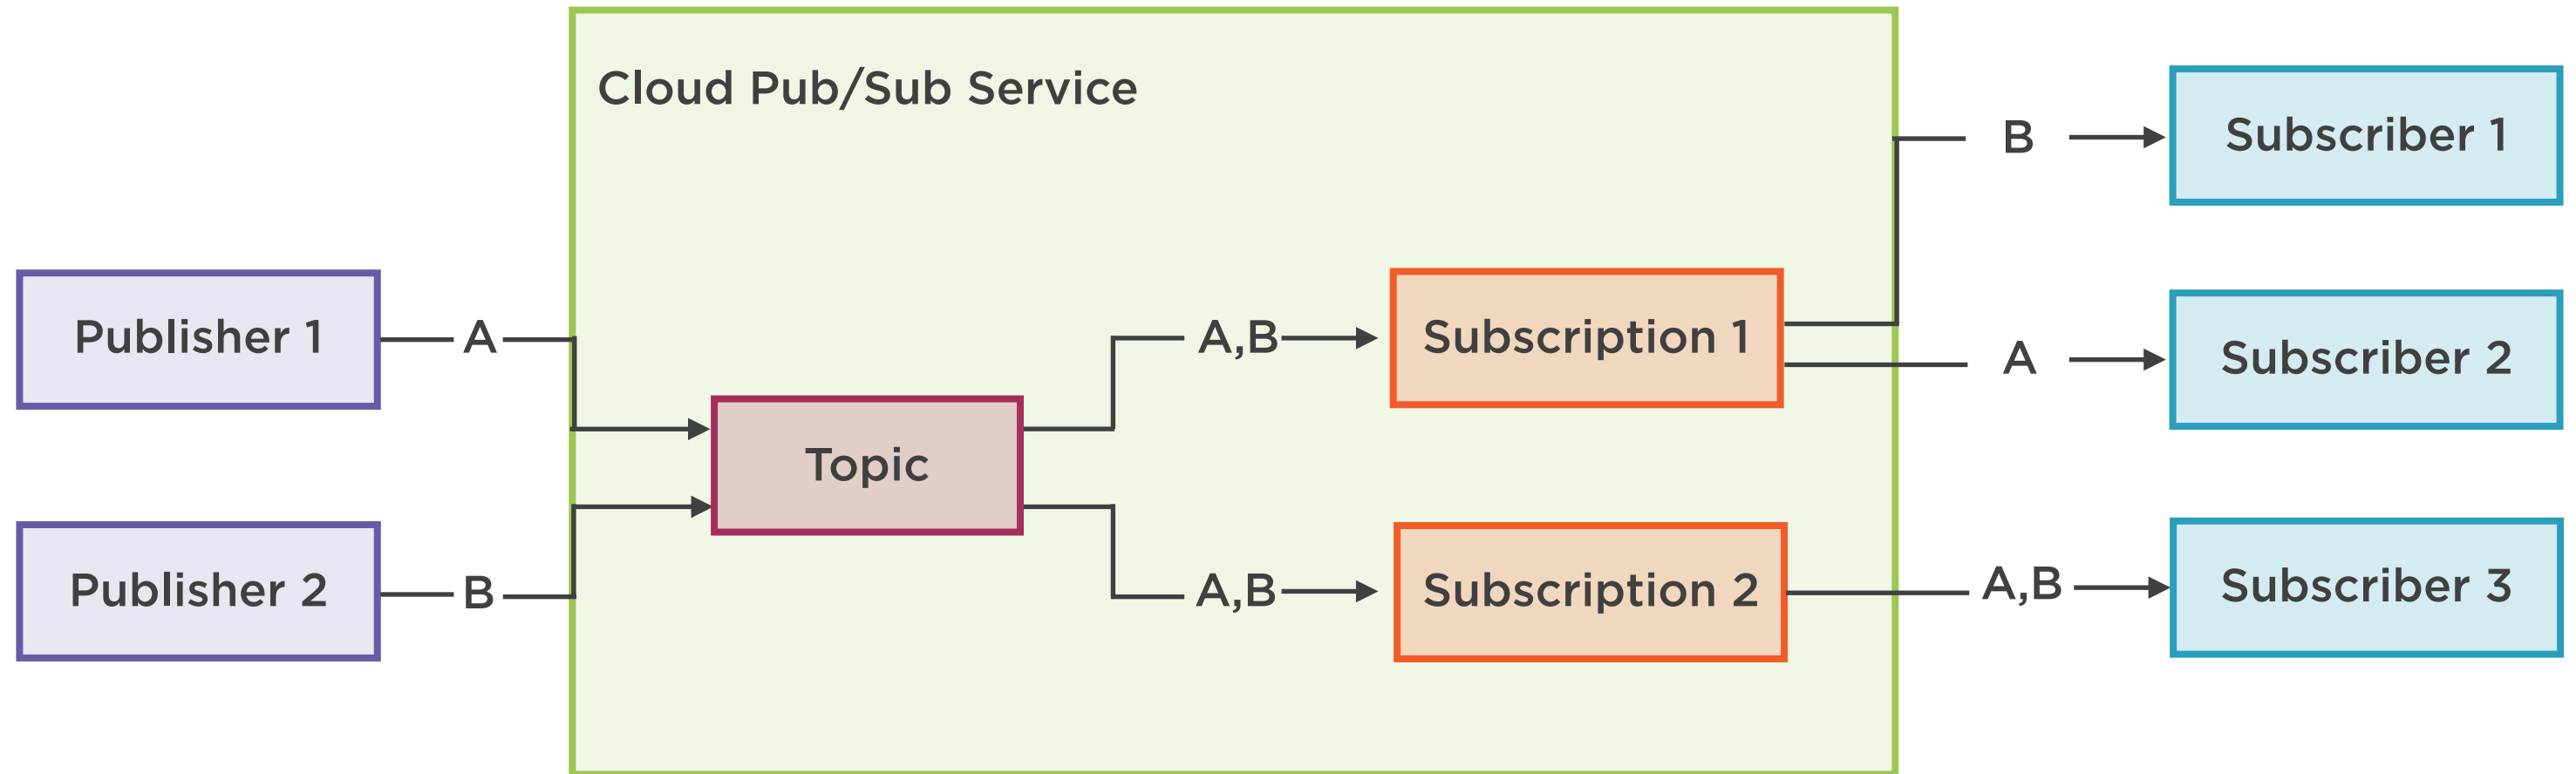
# A Publisher Can Publish to Multiple Topics



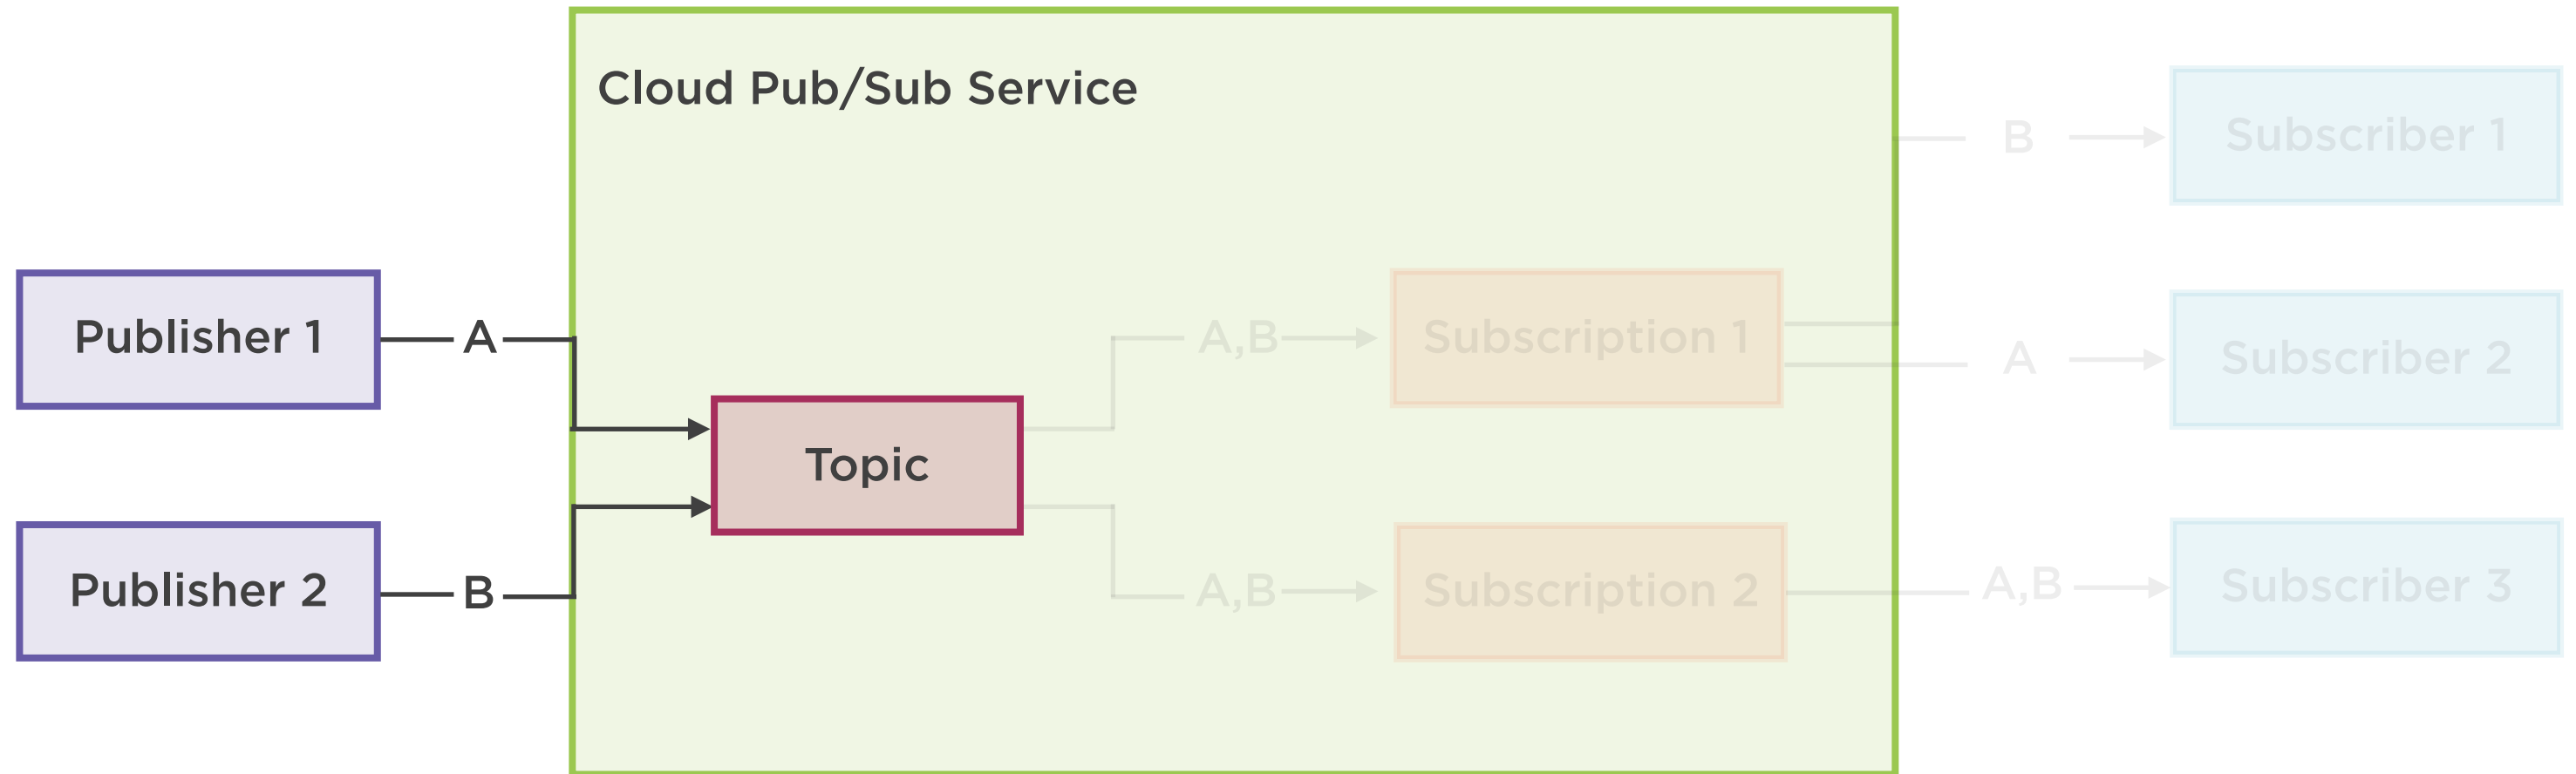
# Subscribers Receive Messages from Subscriptions




# Topics and Subscriptions



# Two Messages on a Topic

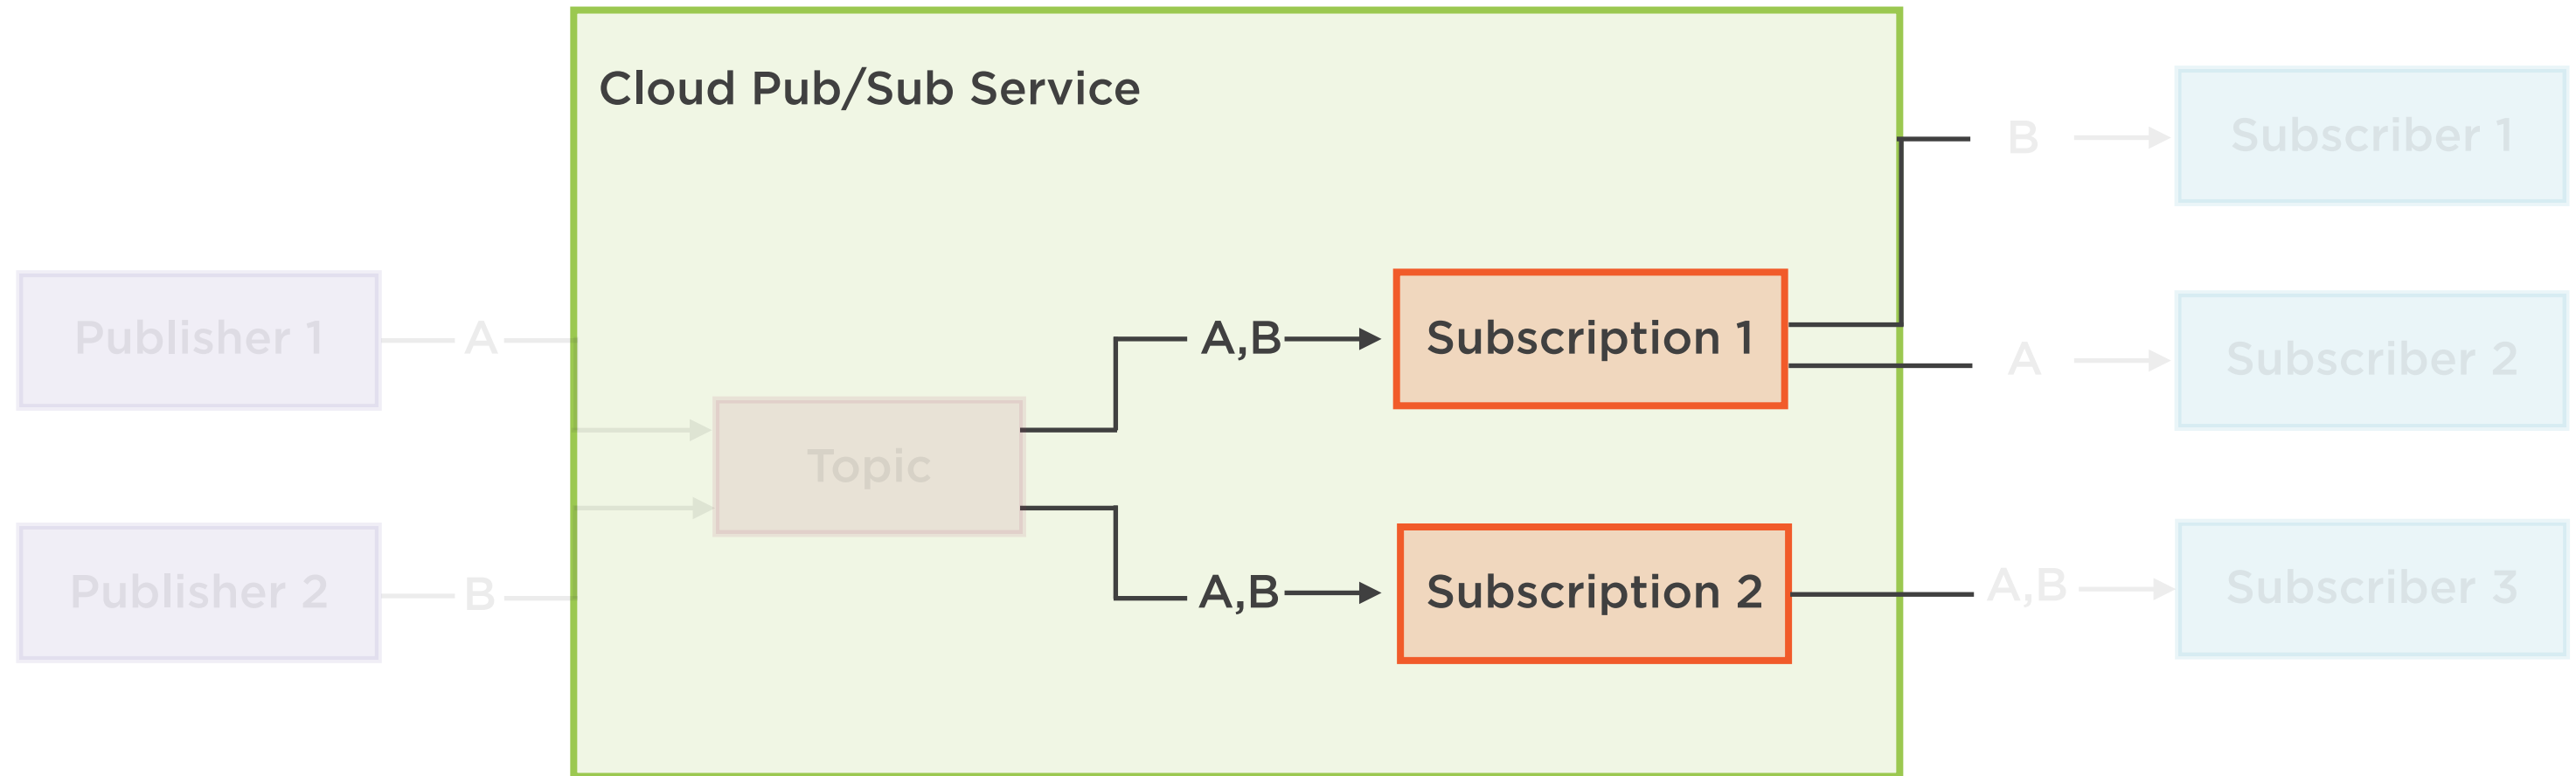


# Two Subscriptions for the Topic

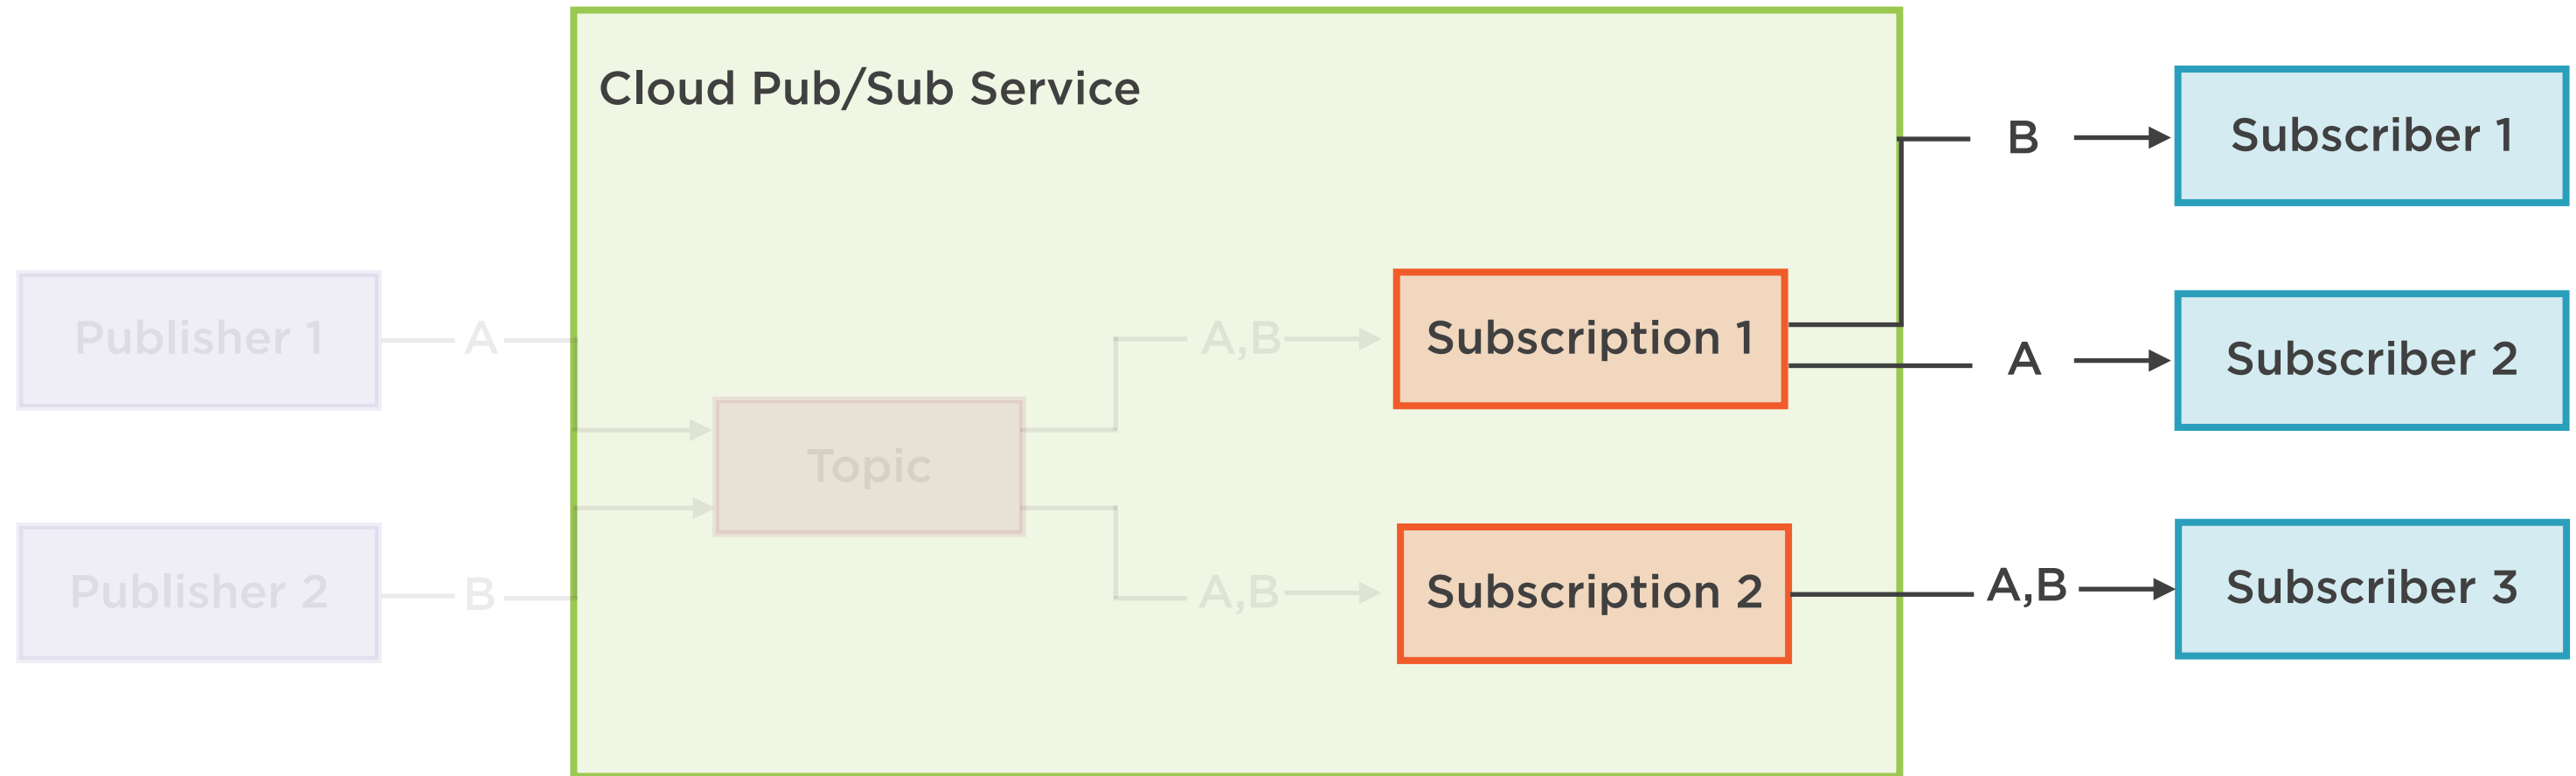




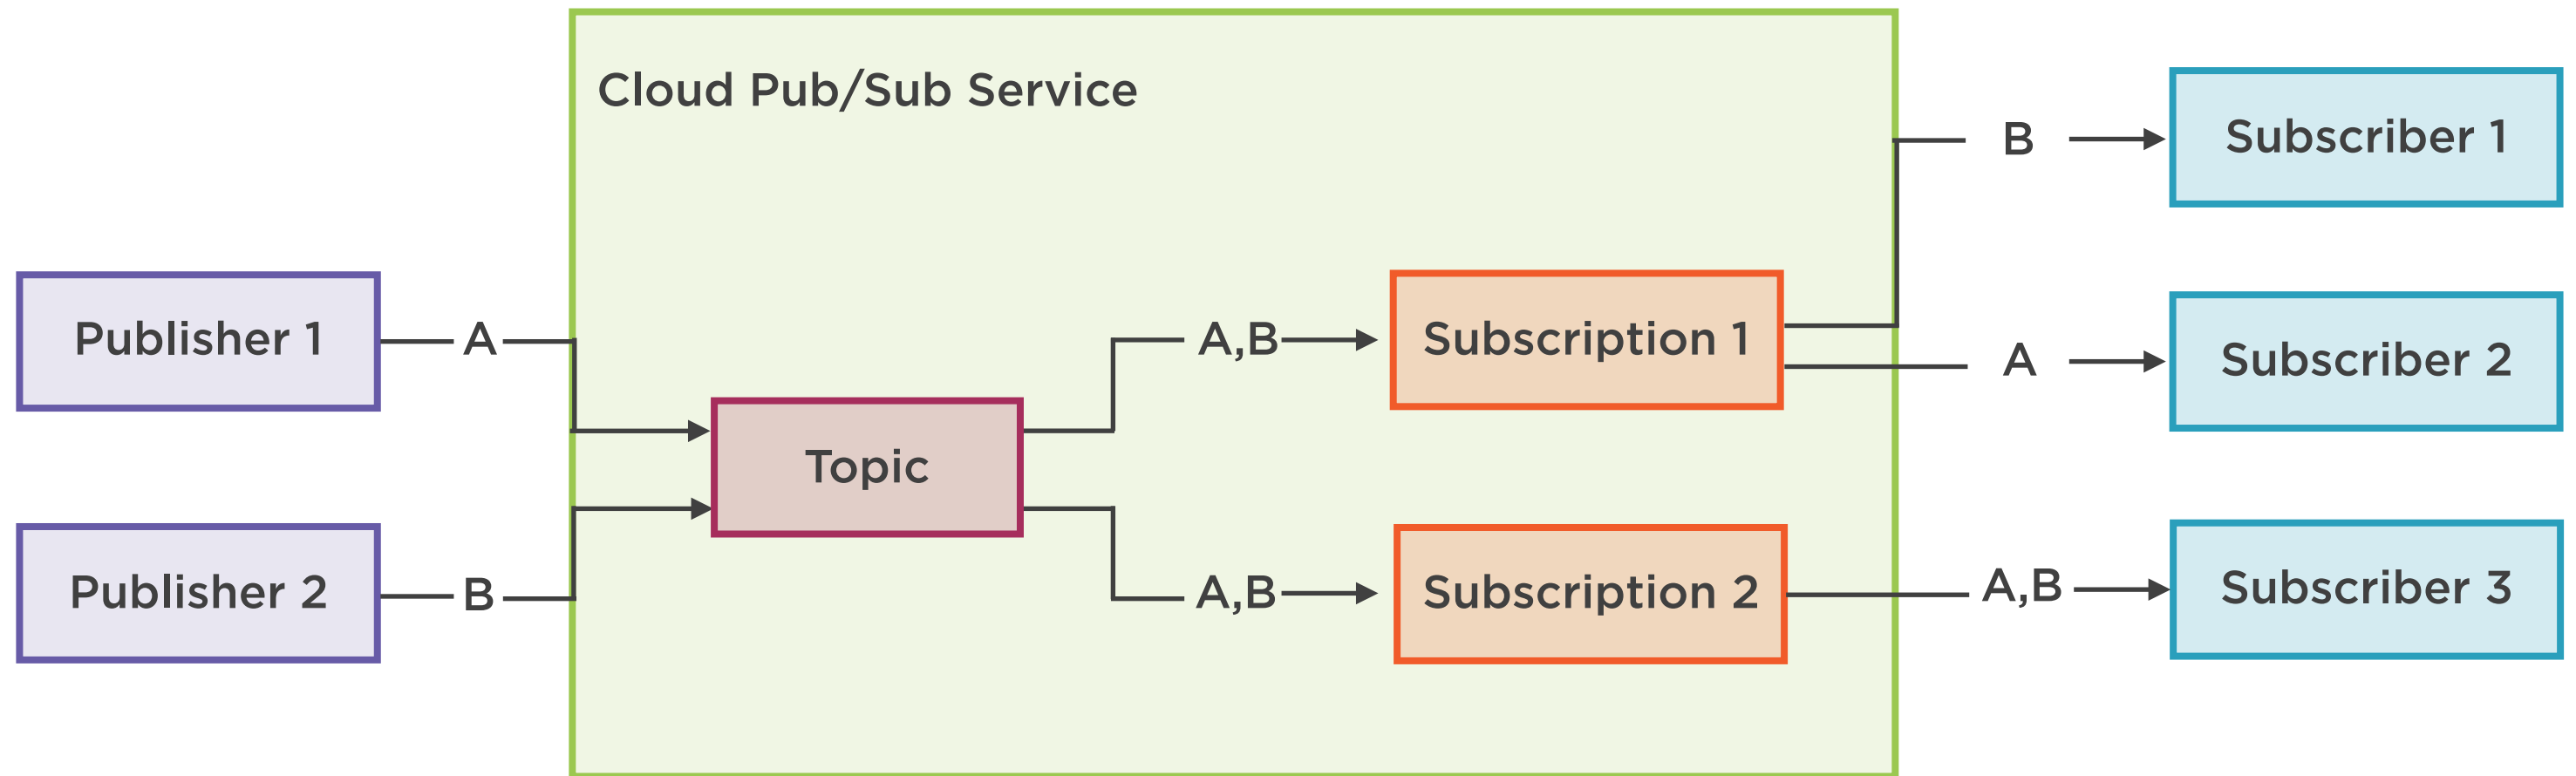
# Same Messages Available on Both Subscriptions



# Multiple Subscribers Share a Subscription



# Fan-in and Fan-out



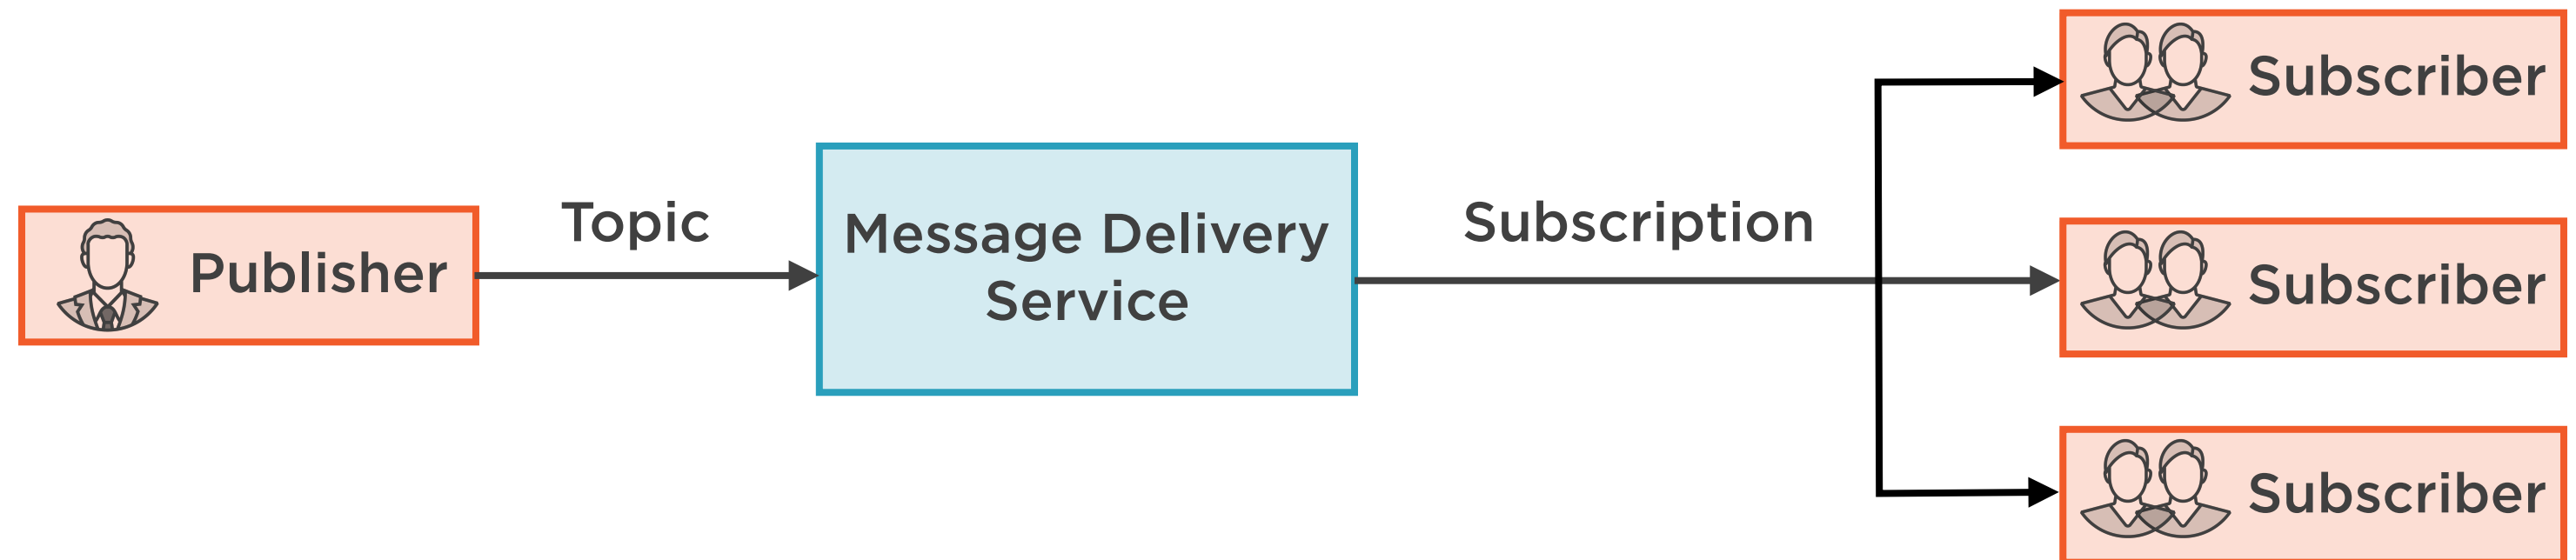
# Publishers and Subscribers

---

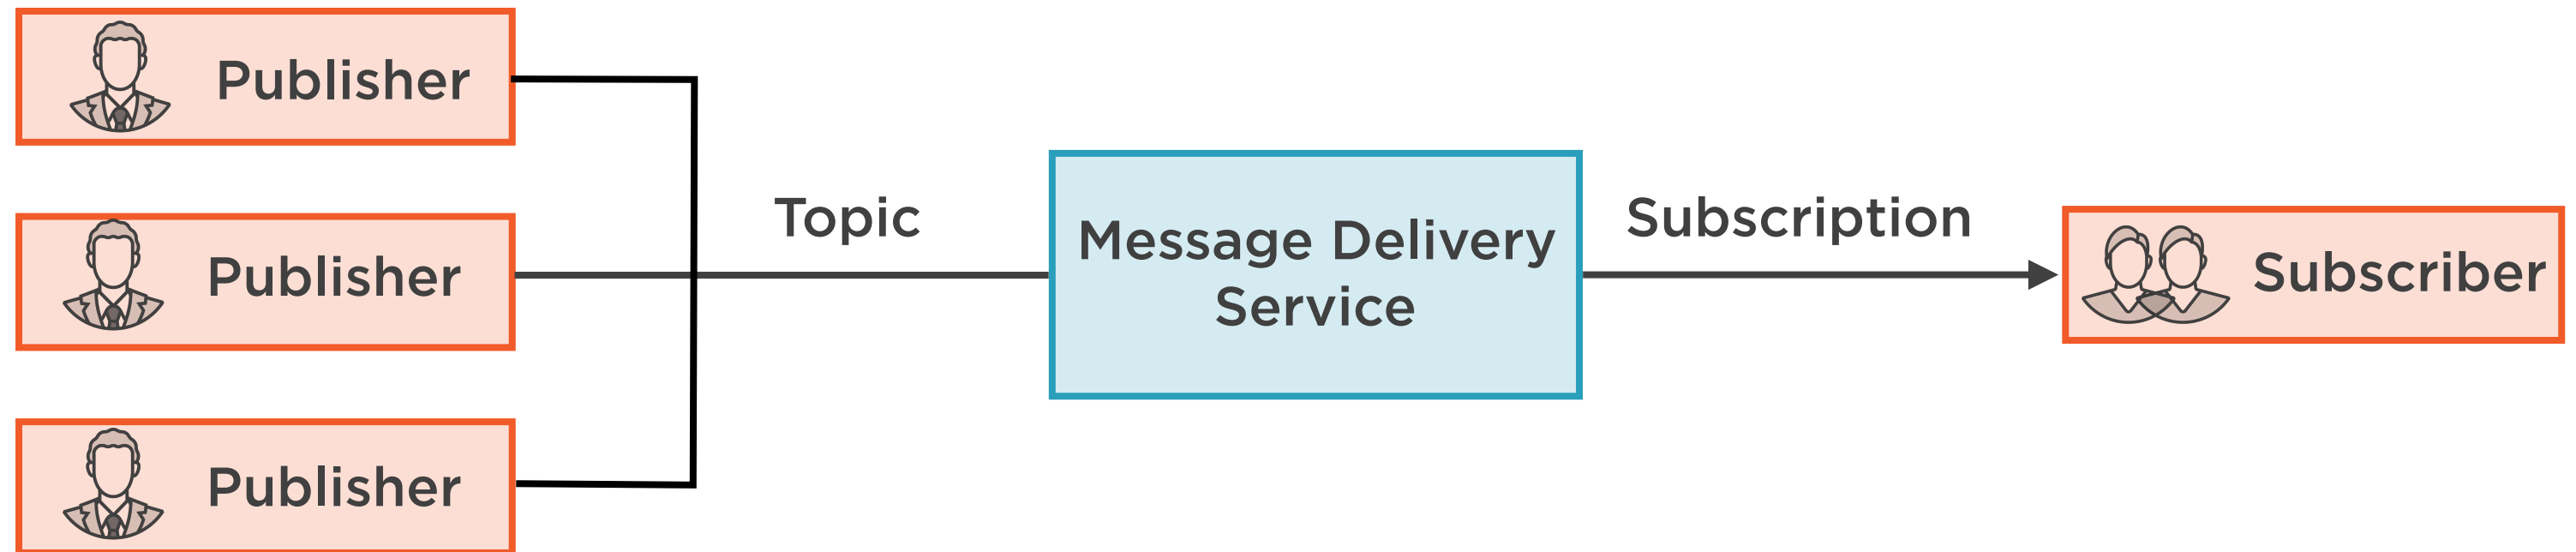
# Single Publisher, Single Subscriber



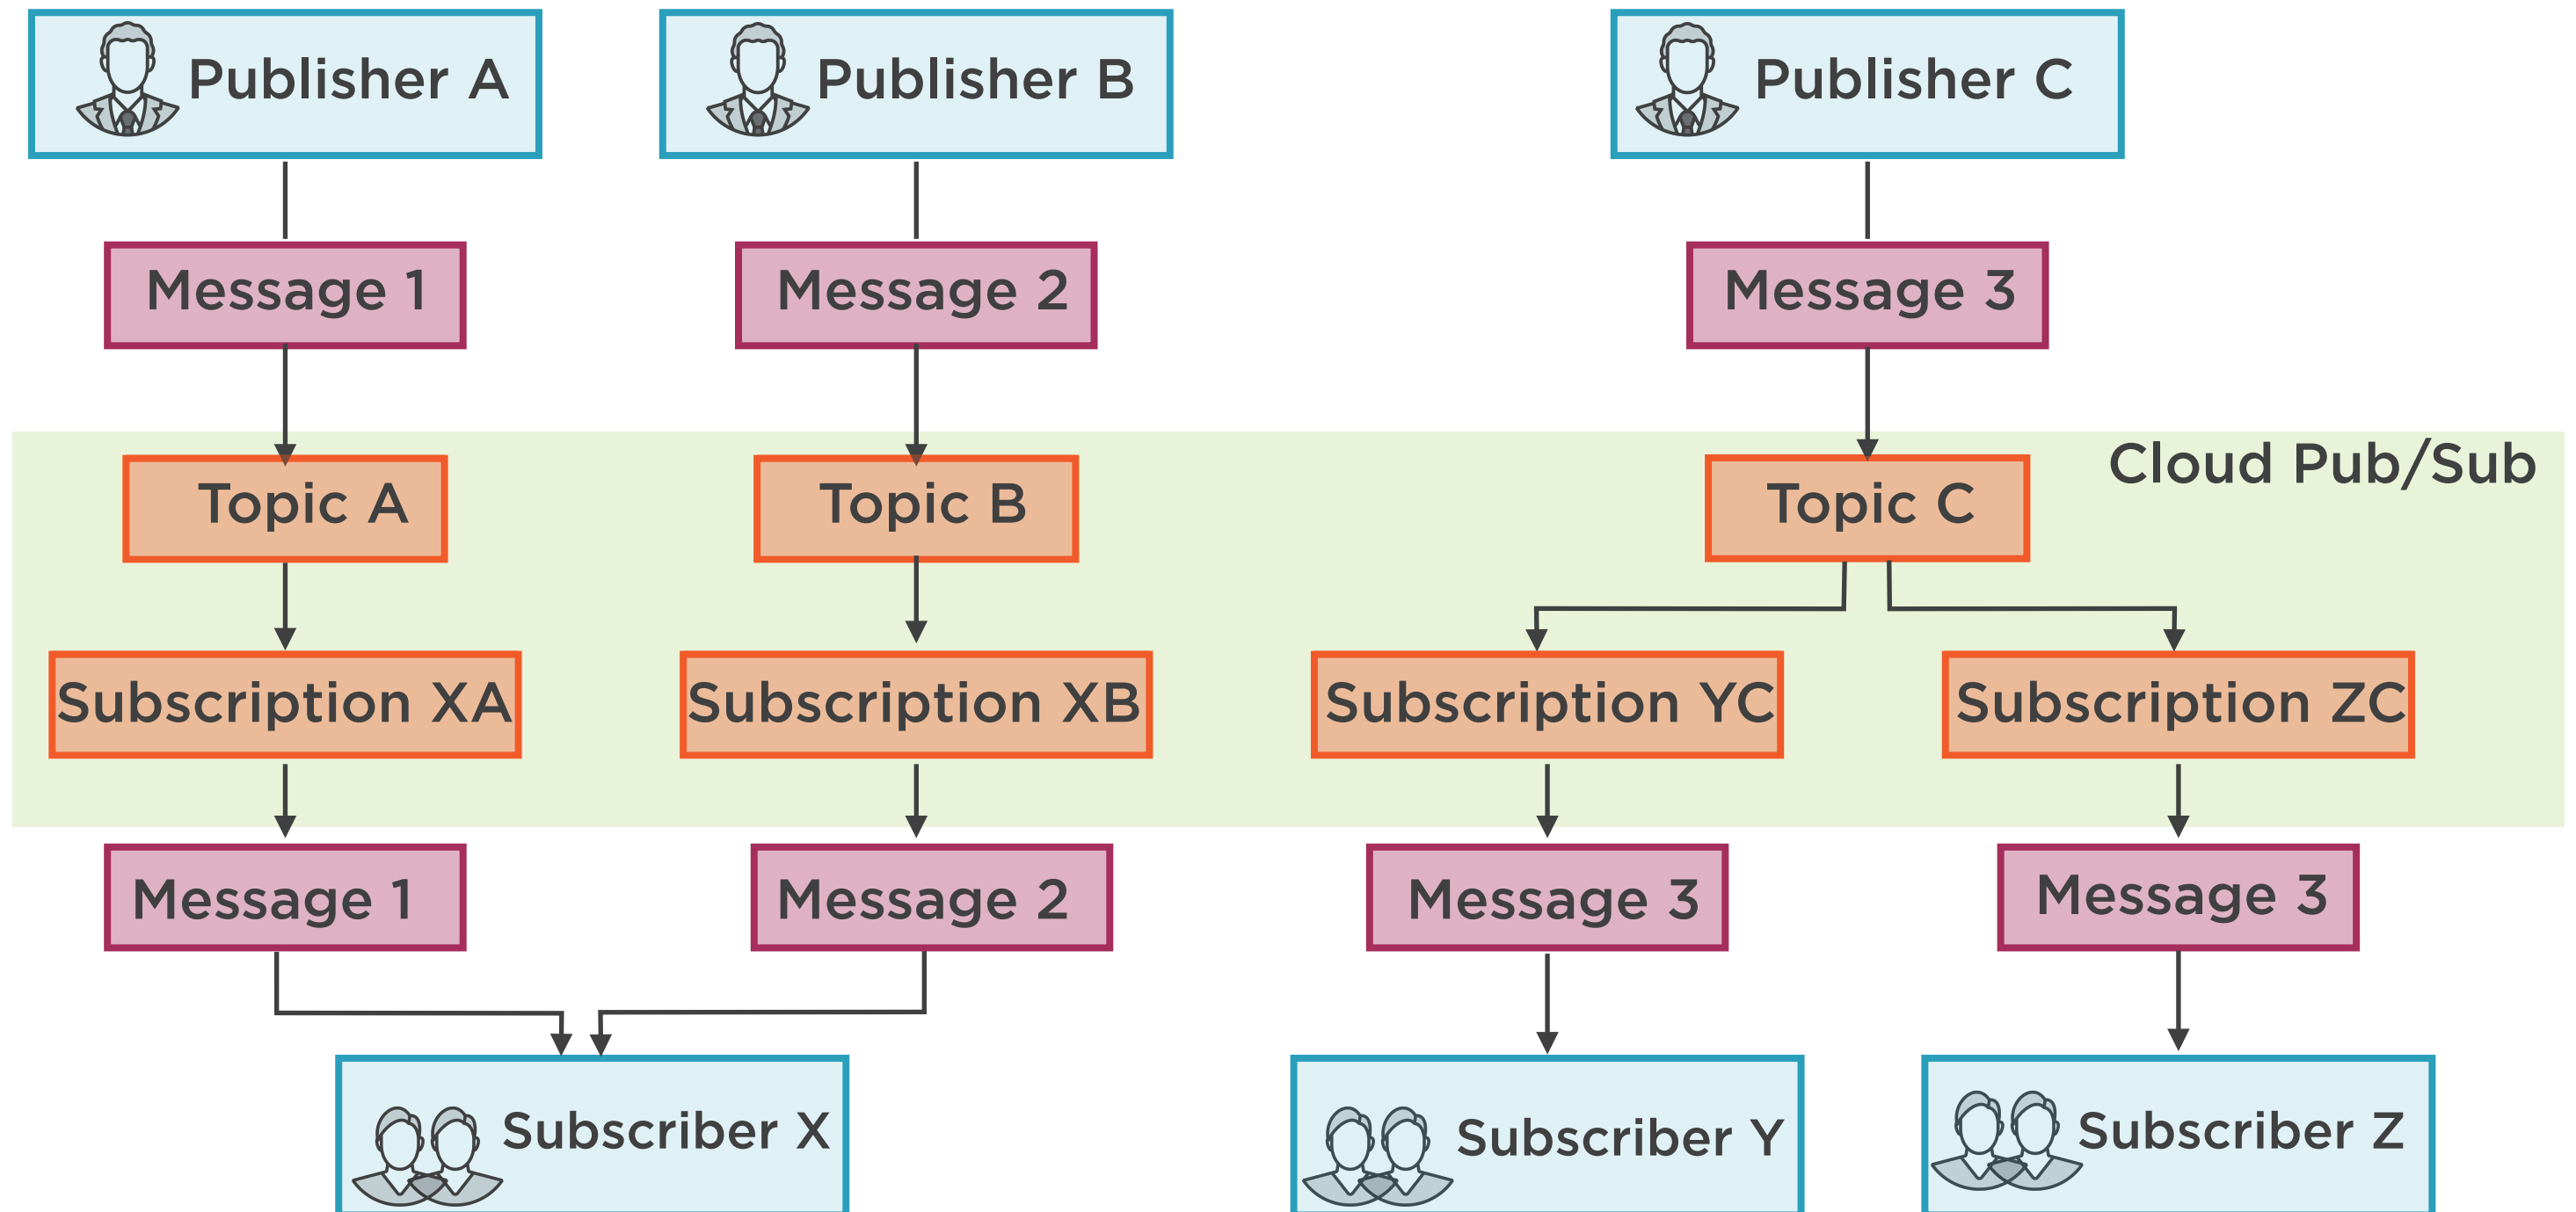
# Single Publisher, Multiple Subscribers



# Multiple Publishers, Single Subscriber




# Publishers and Subscribers






# Publishers and Subscribers



# Publishers



App Engine and Cloud Function apps


GCE, GKE apps

App running on third party network

Any mobile or web app

Any application that can make HTTPS requests to [googleapis.com](https://googleapis.com)

# Publishers

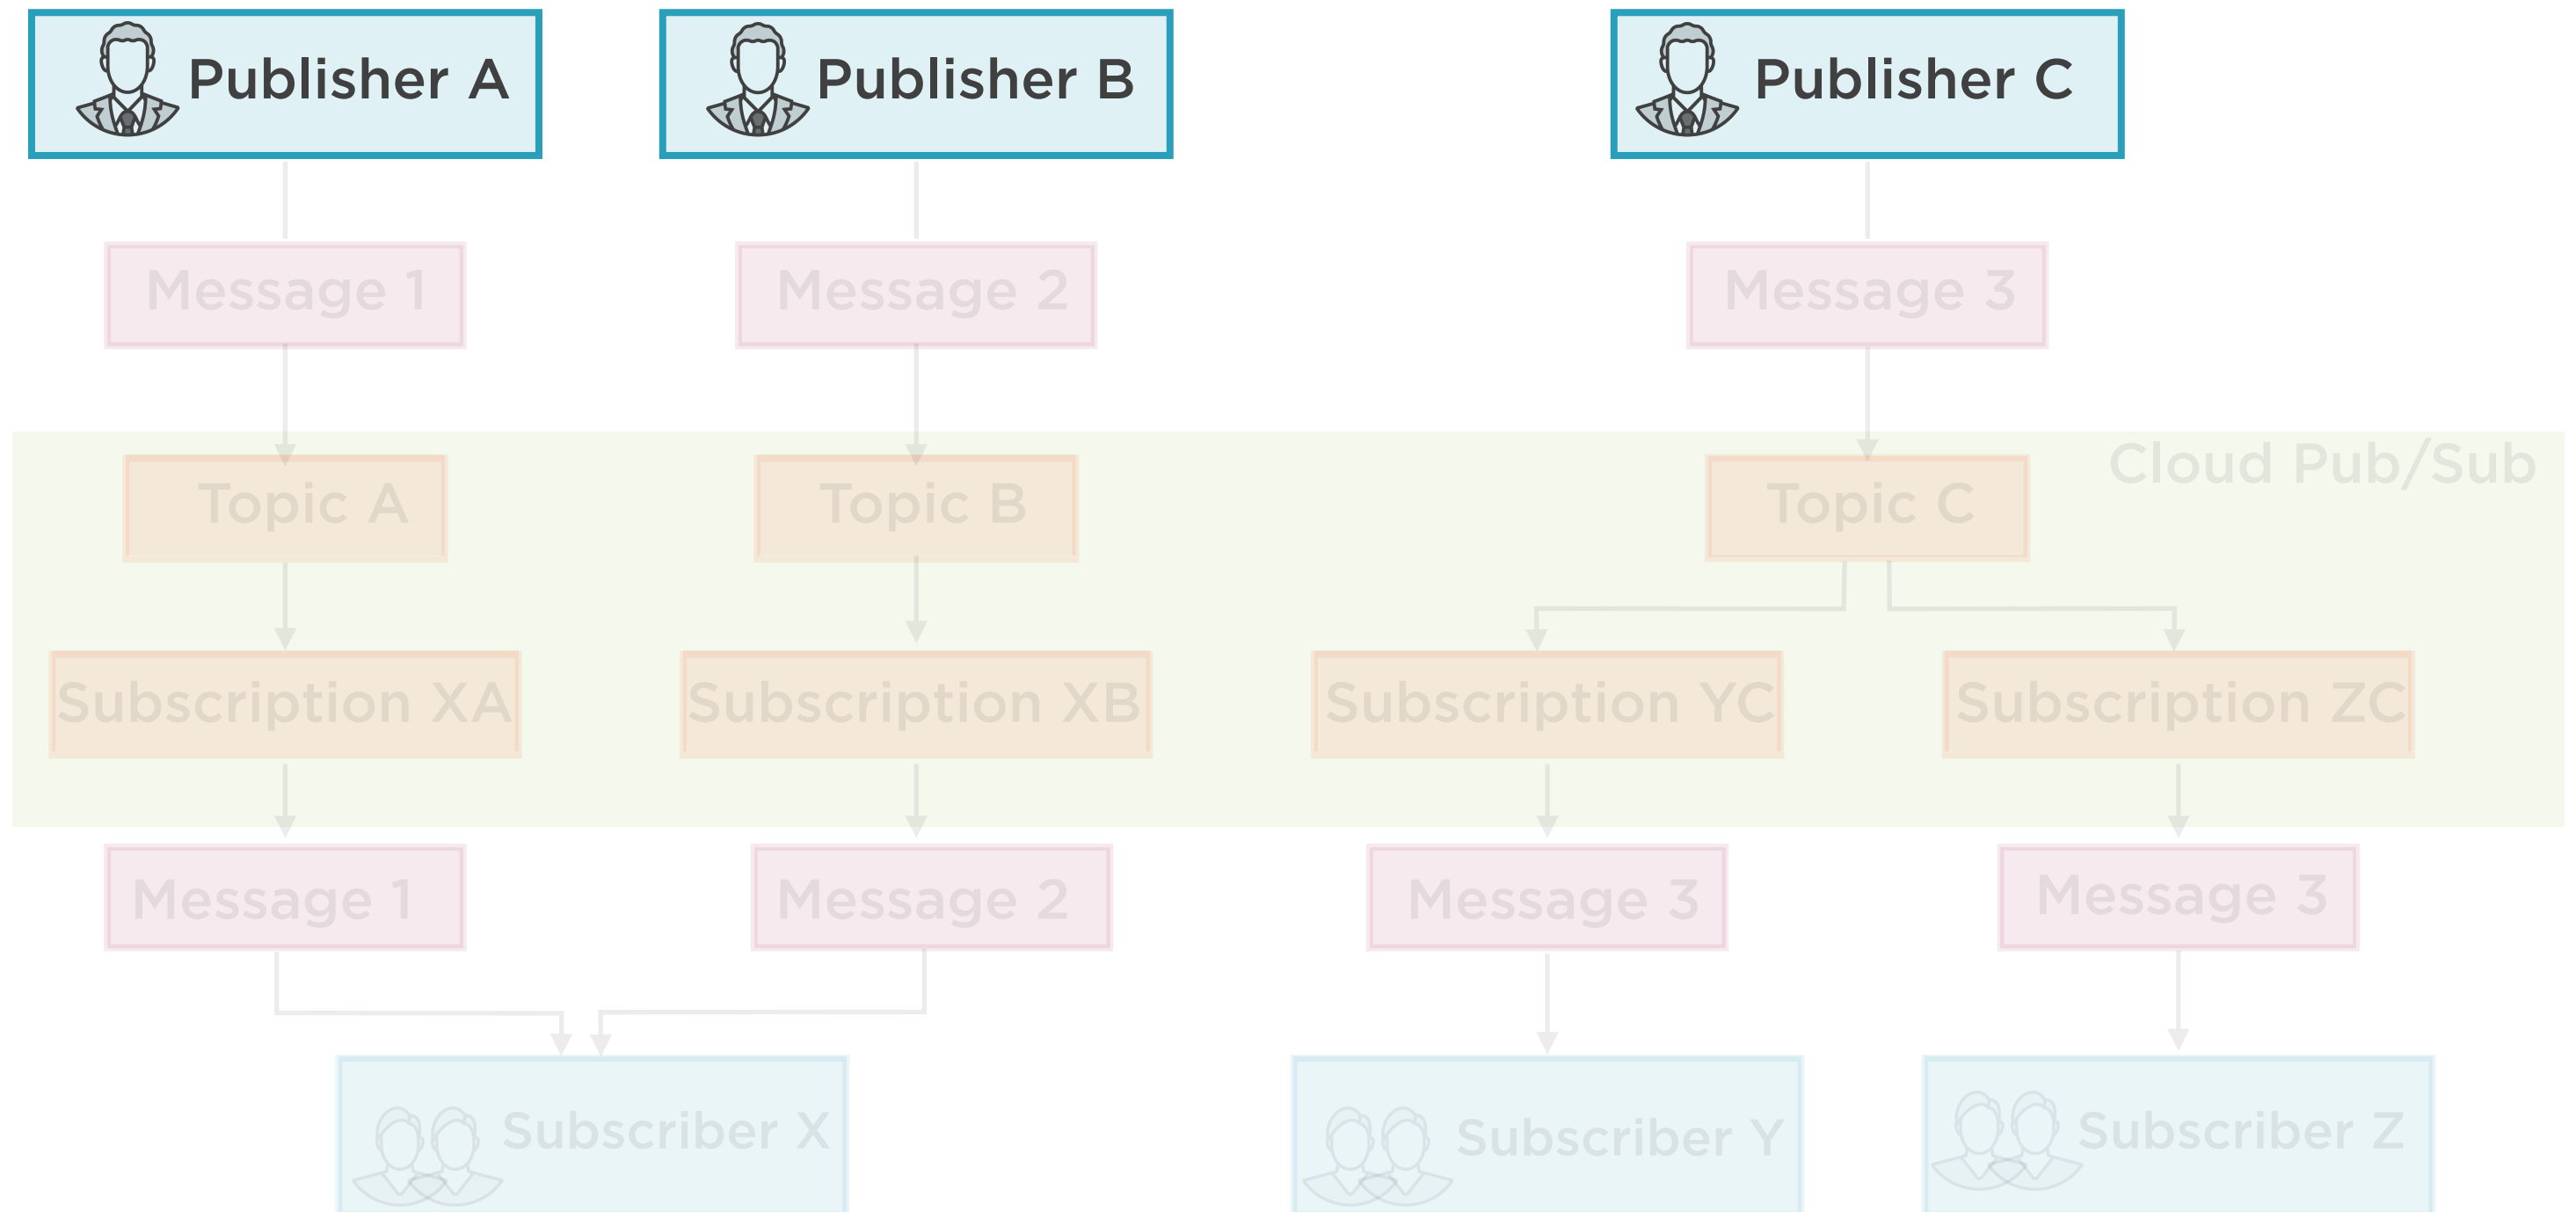


**Publisher** apps create and send messages on a **Topic**


**Messages** persisted in a message store - until delivered/acknowledged

One queue per subscription


# Publishers and Subscribers



# Publishers and Subscribers



# Subscribers



**Subscriber** apps subscribe to a topic to receive messages

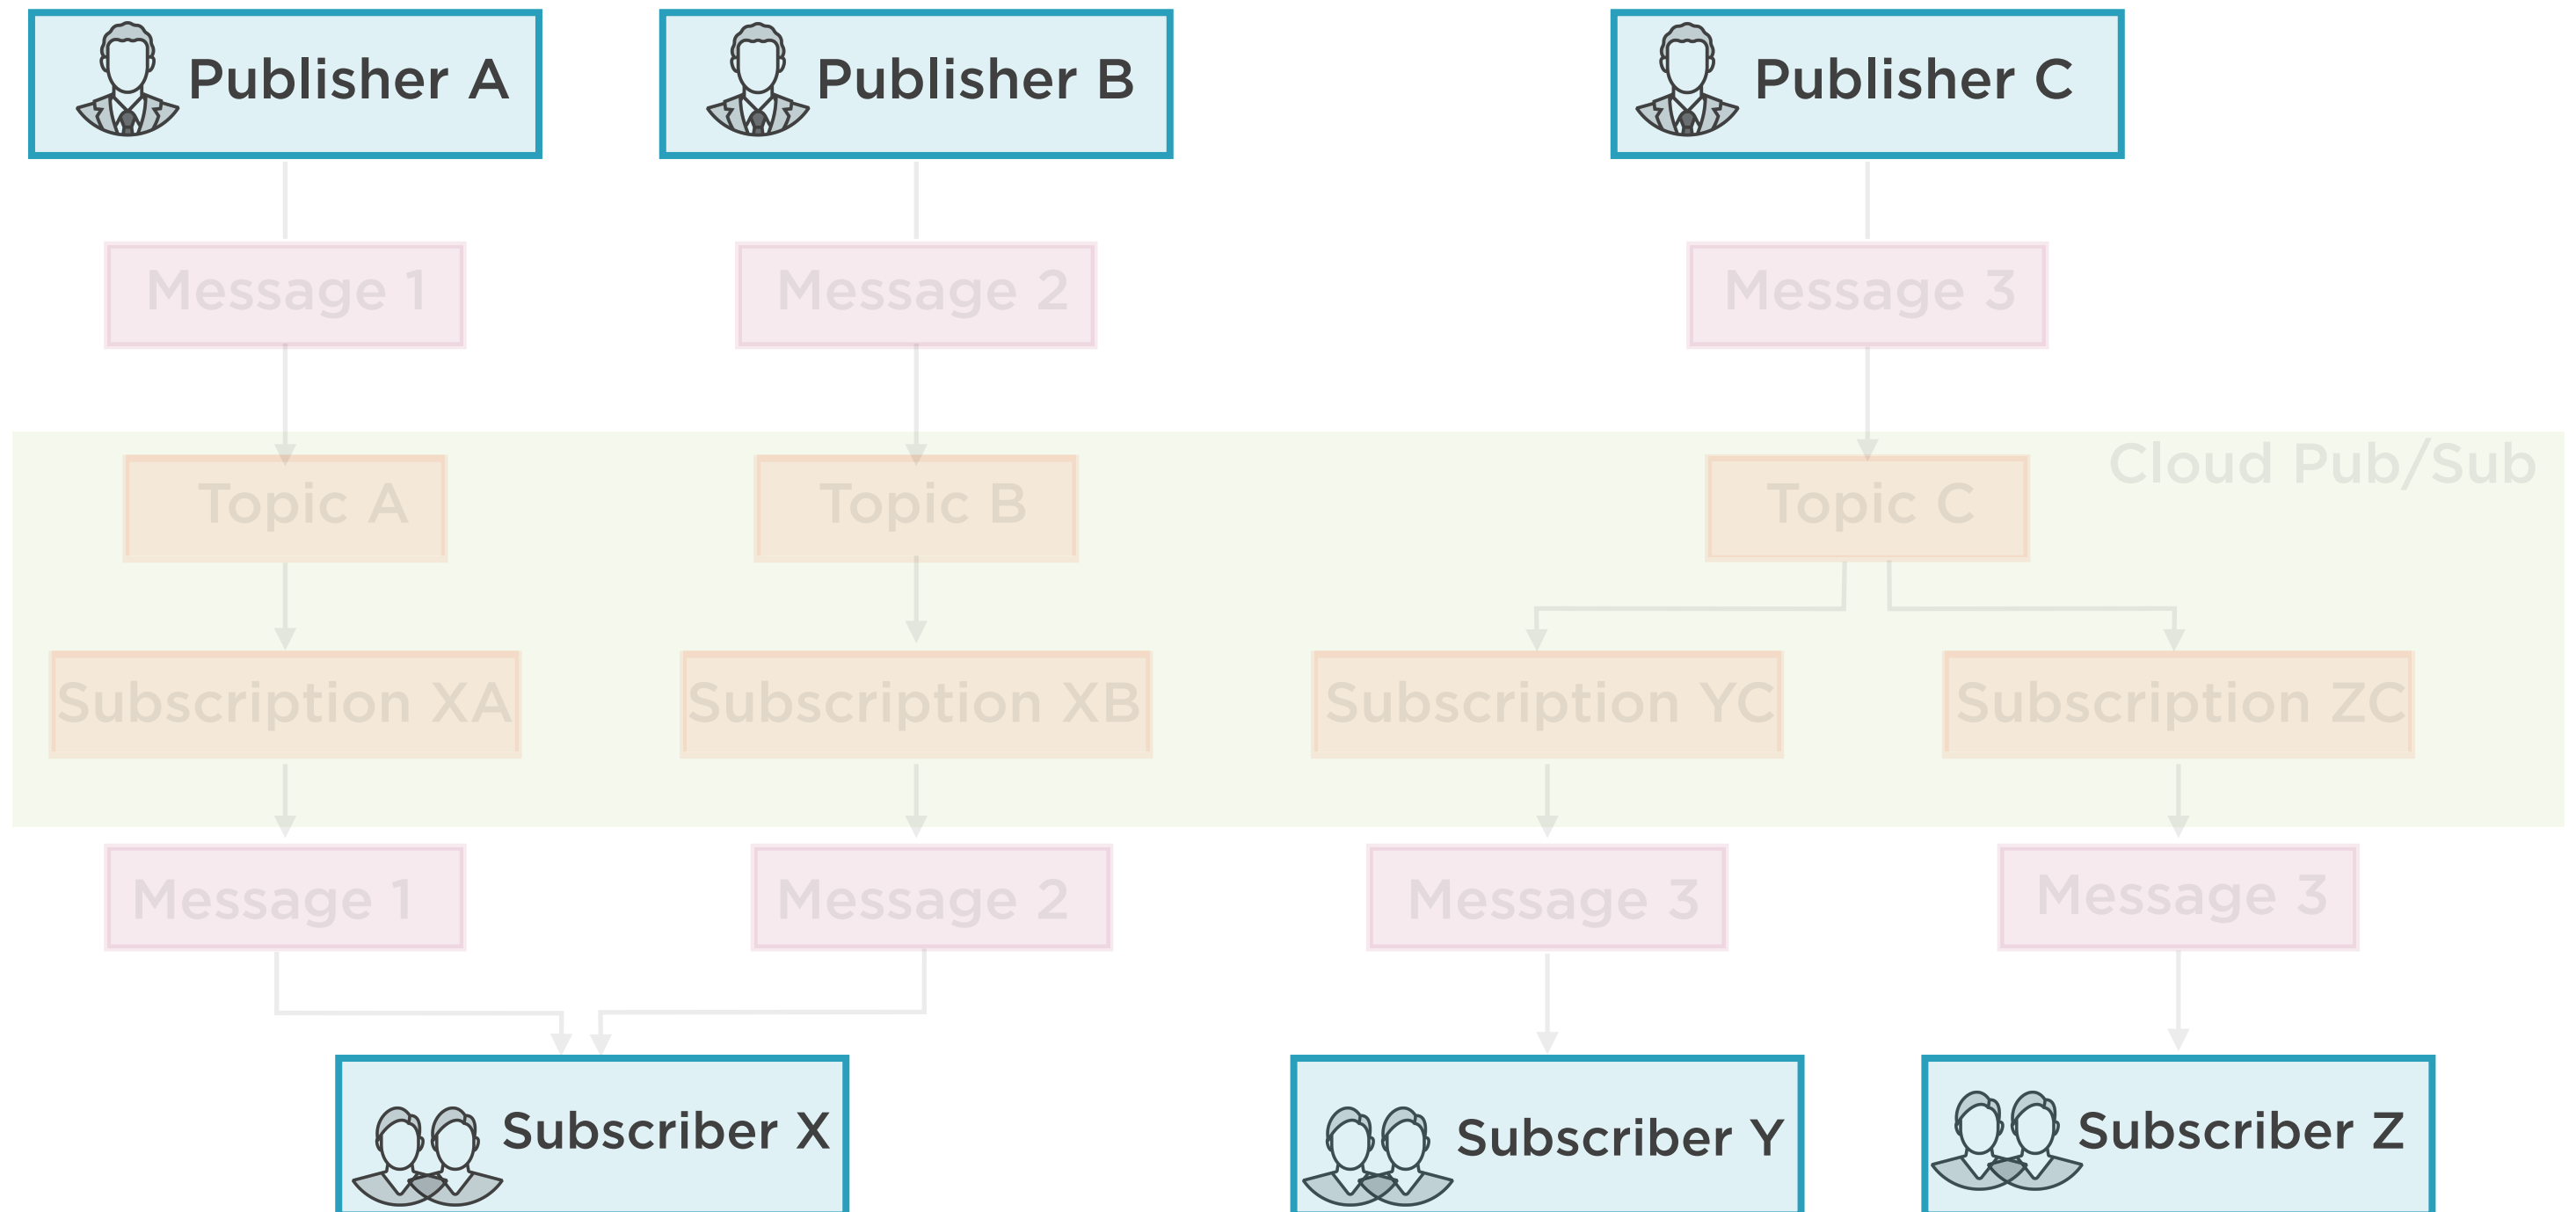
Once acknowledged by subscriber, message deleted from queue

Push and Pull subscriptions

# Messages, Topics and Subscriptions


---

# Publishers and Subscribers





# Publishers and Subscribers



# Messages



Message data

Attributes: Key-value pairs set by publisher

Organized into topics

Stored until delivered and acknowledged

# Messages



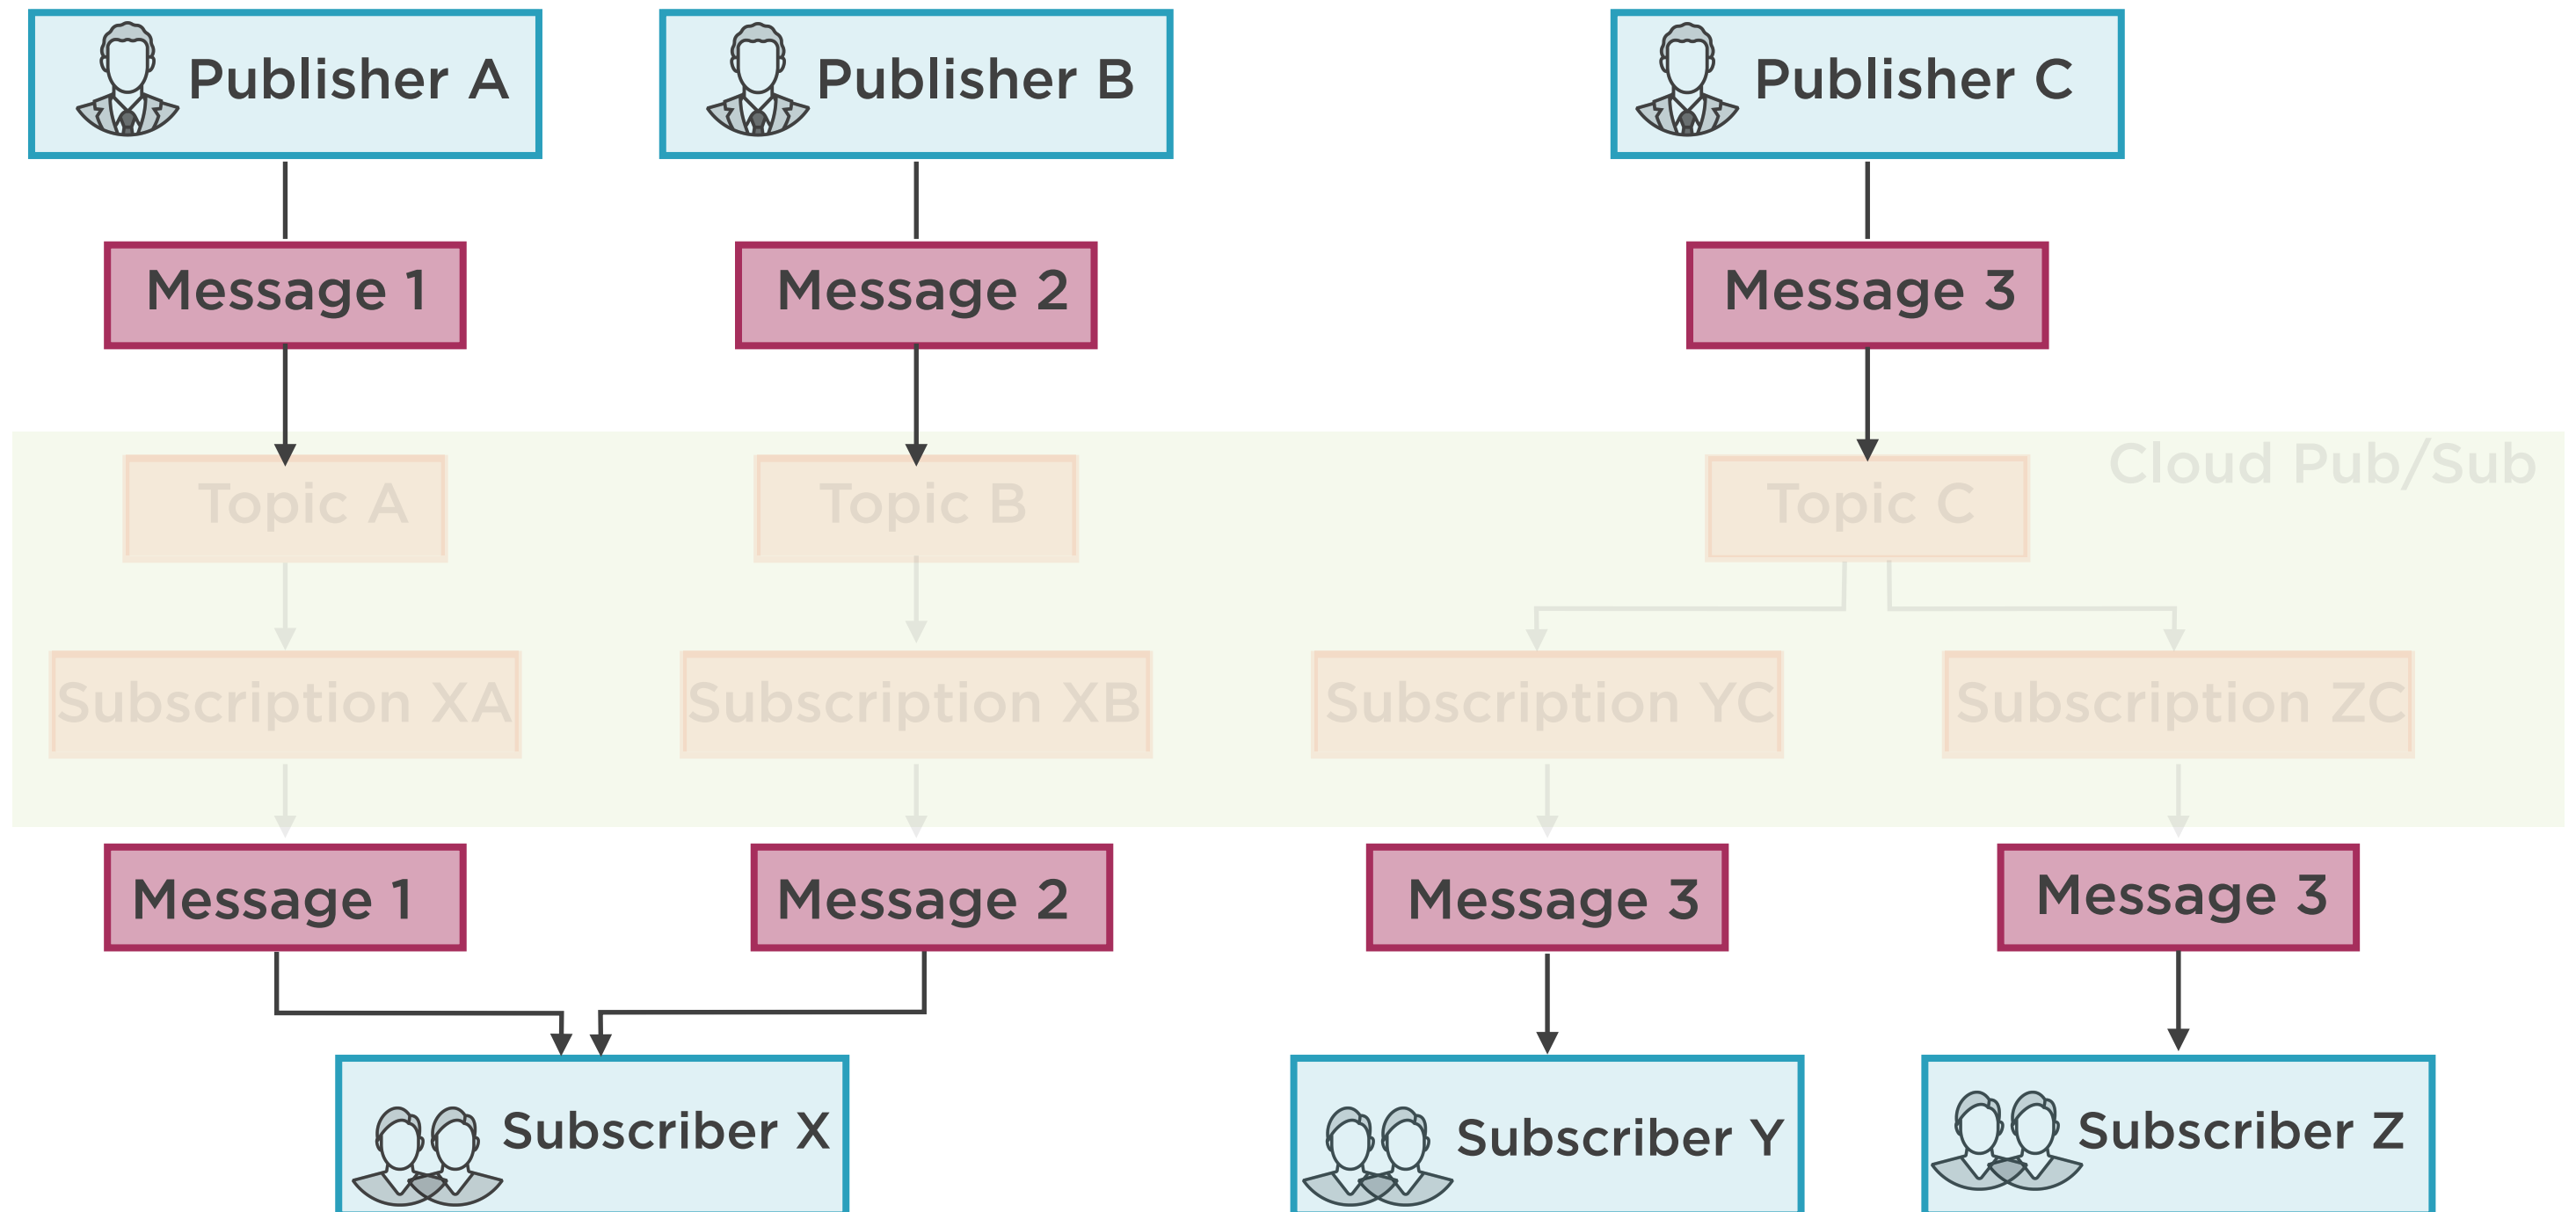
## **Pub/Sub provides**

- At-least-once delivery
- No guaranteed ordering


## **Use Dataflow with Pub/Sub**

- Ordering
- De-duplication

# Publishers and Subscribers



# Publishers and Subscribers



# Topics




Named resources to which publishers send messages

Used to organize messages

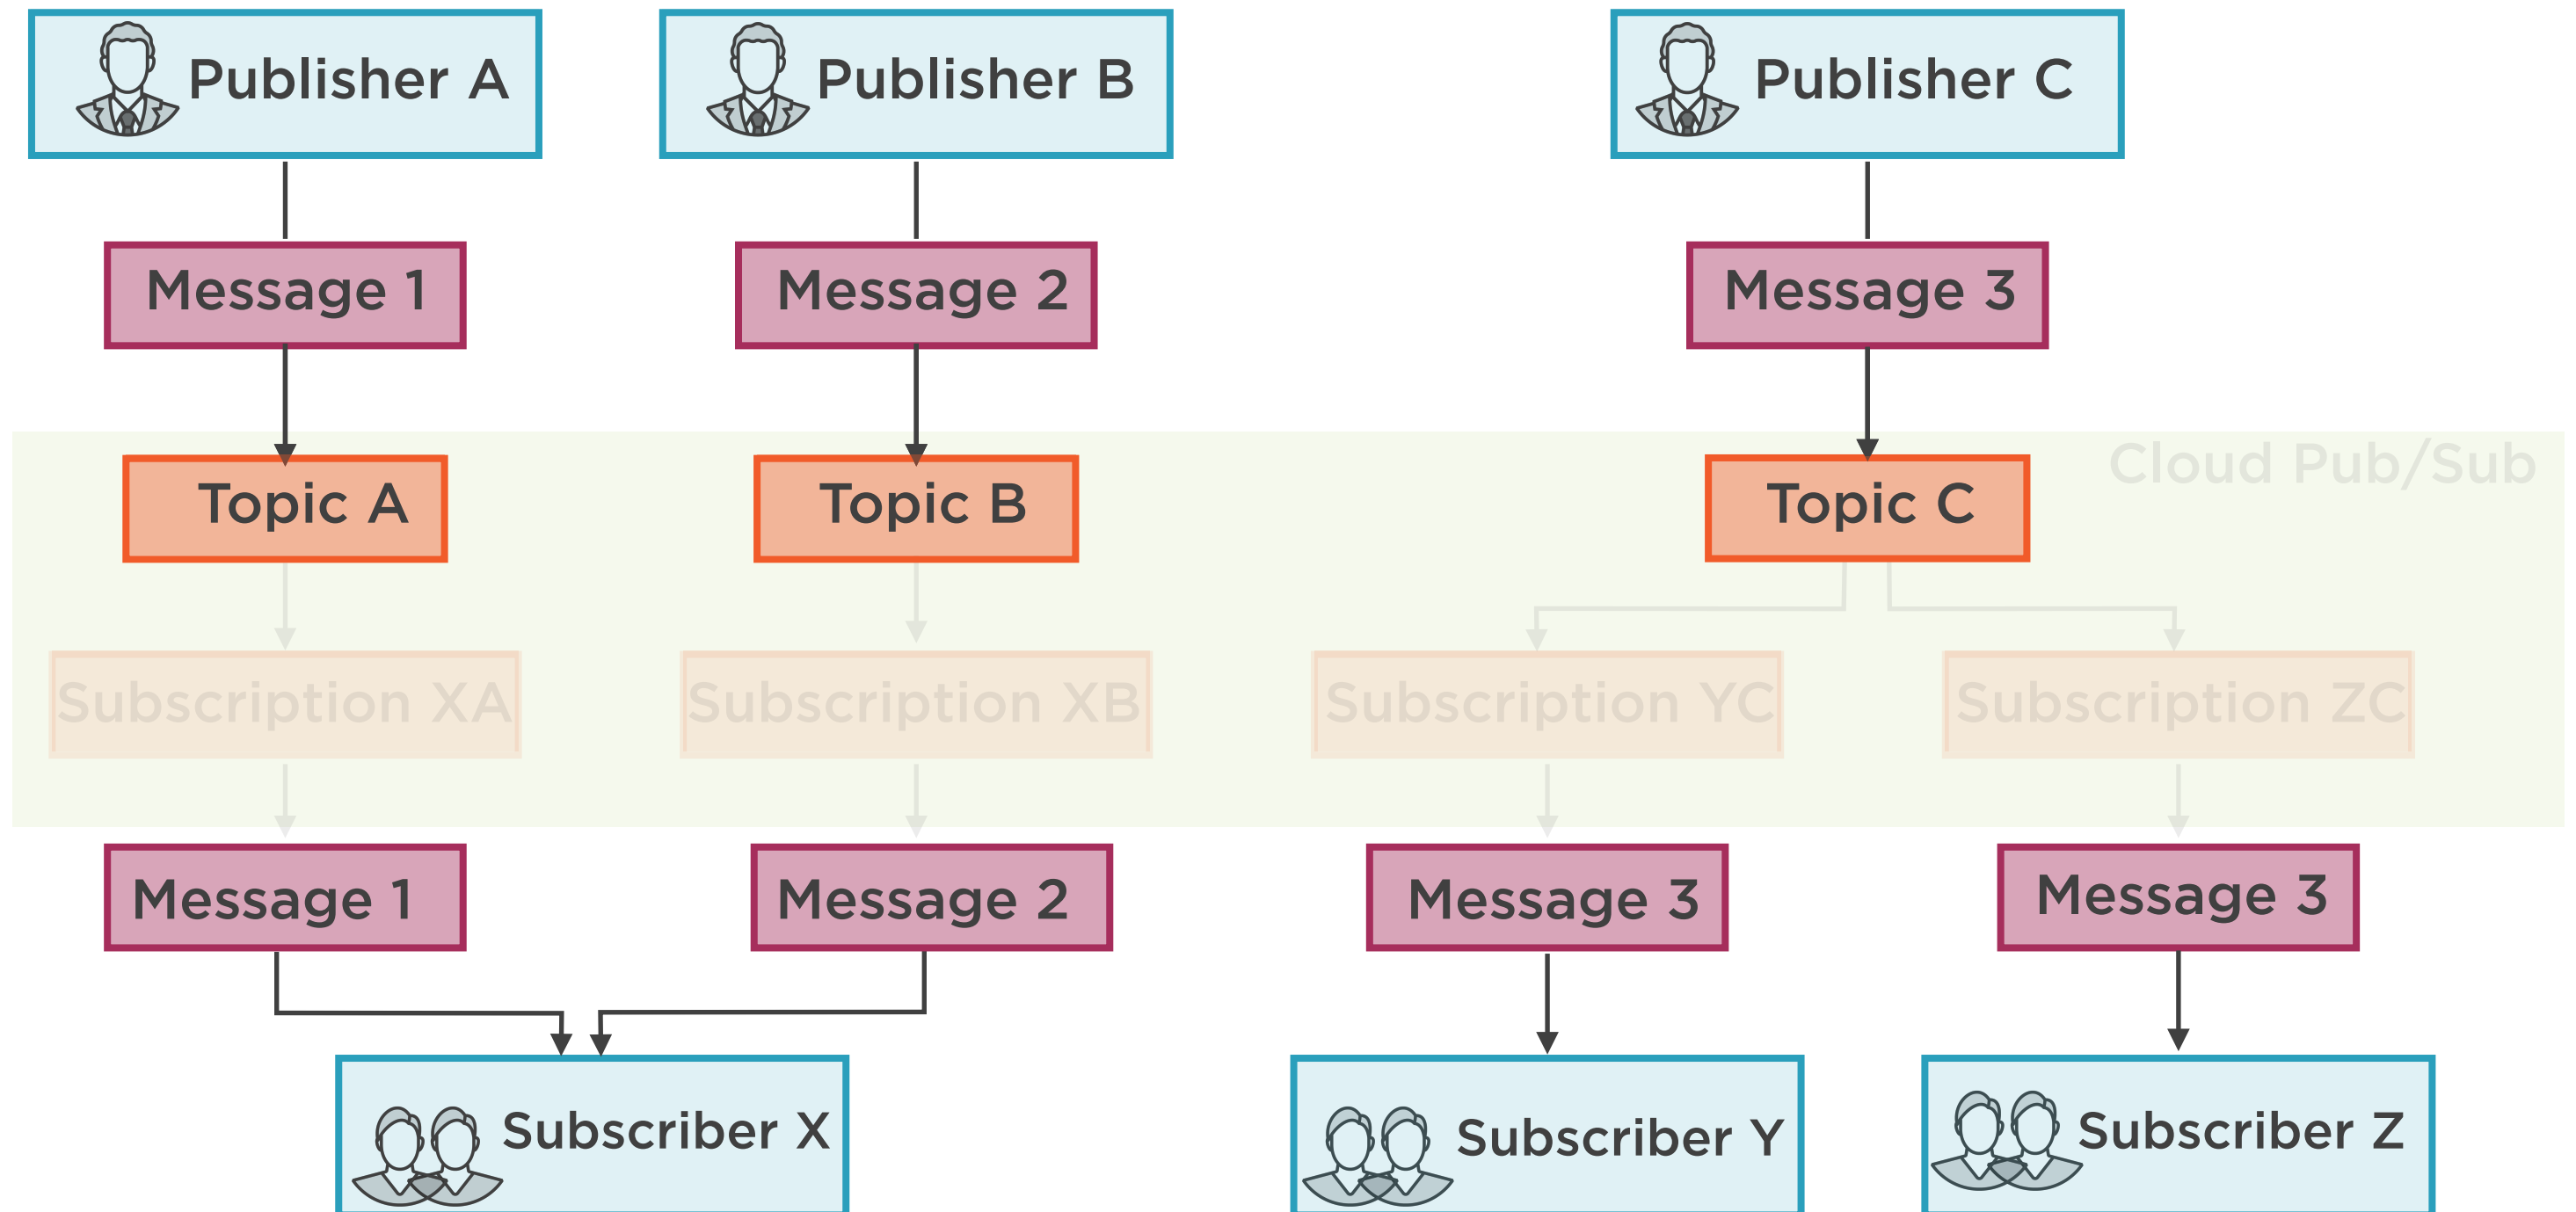
Publishers send messages to topics

Subscribers subscribe to topics

# Publishers and Subscribers

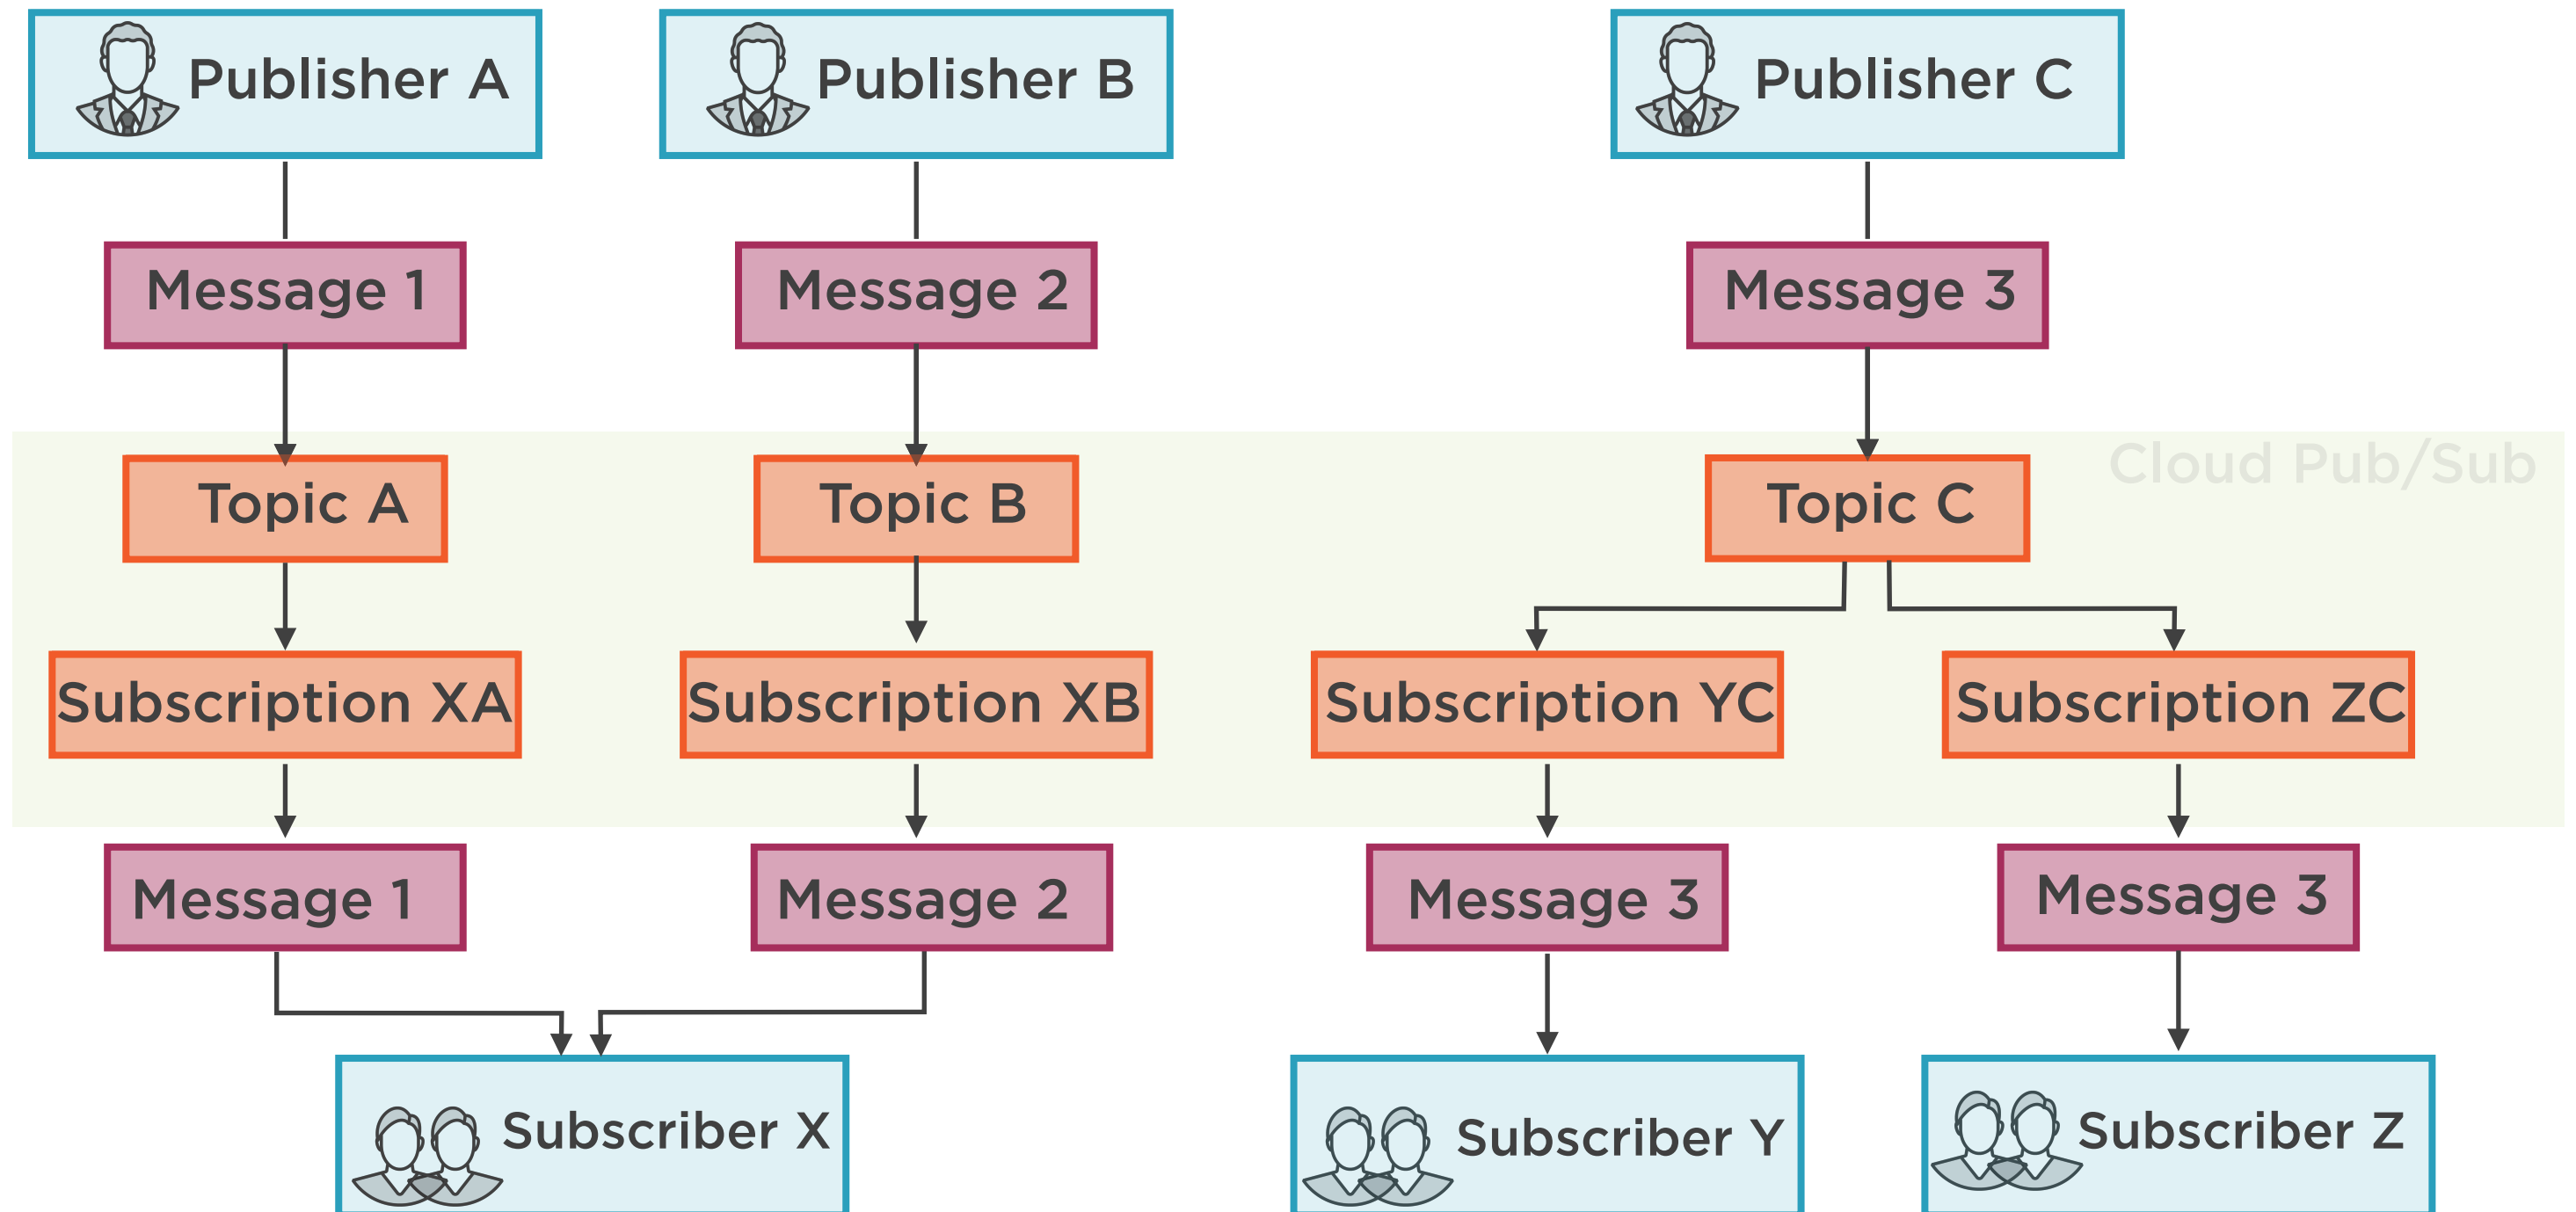


# Publishers and Subscribers





# Publishers and Subscribers



# Subscriptions



Pub/Sub sends each message on a topic to each subscription individually

Message persisted inside Pub/Sub

Not deleted until **at least one subscriber per subscription** has acknowledged

# Subscriptions



Pub/Sub sends each message on a topic to each subscription individually

Each subscription then might

- Either **push** messages to subscriber endpoint
- Or wait for subscriber to **pull** them at a later point

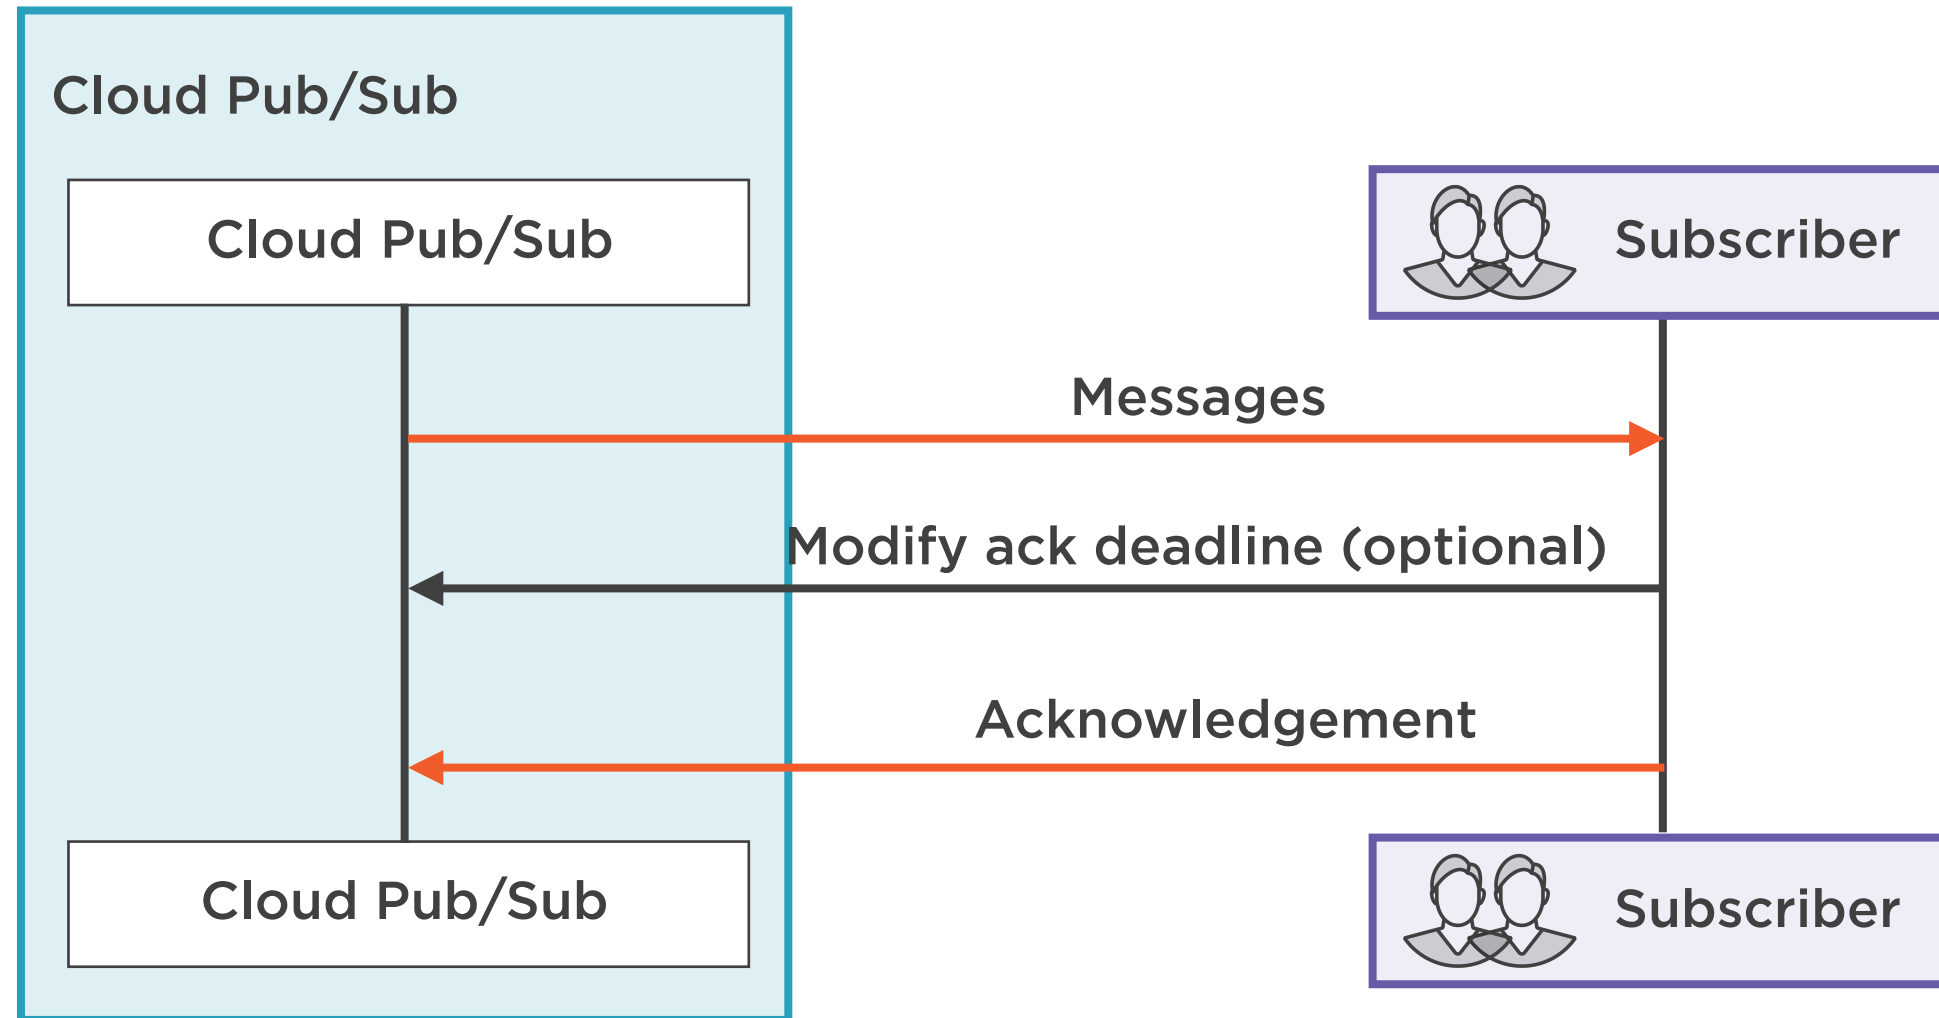
# Subscriptions



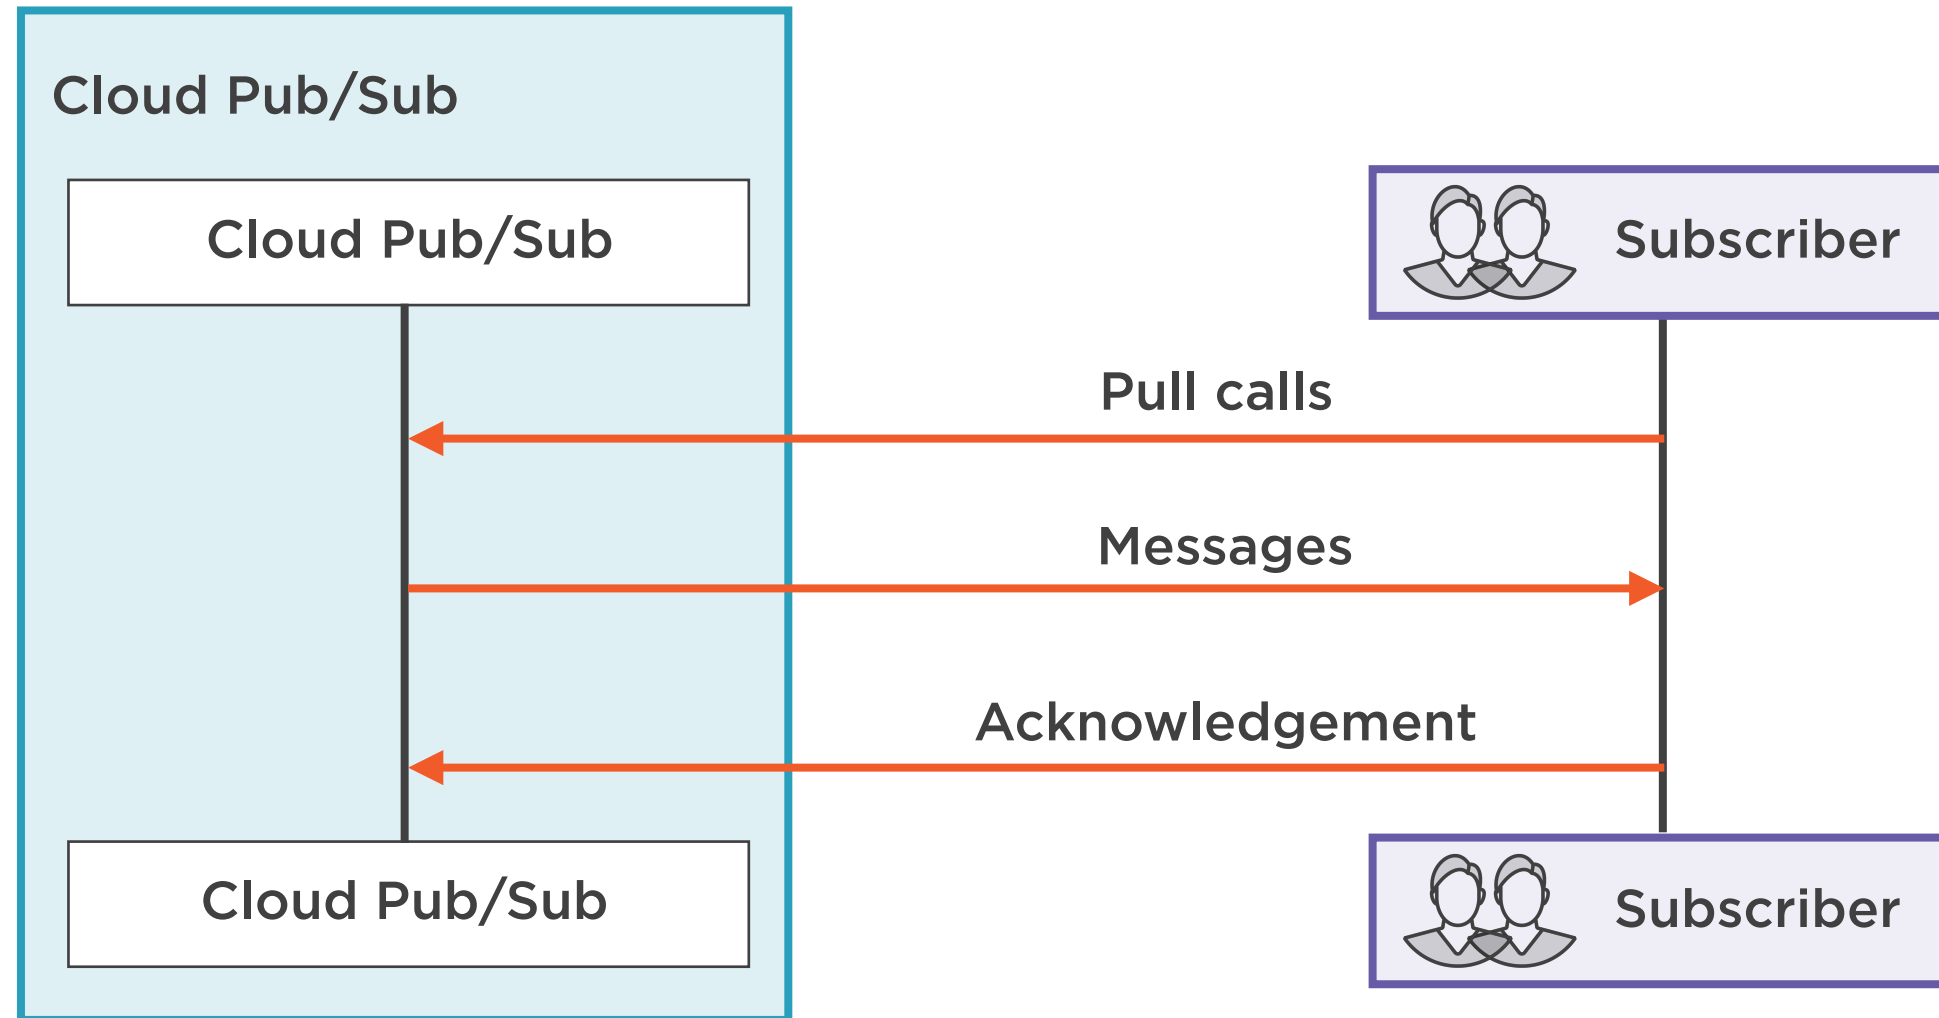
**Push** subscriptions: Any app that can make HTTPS requests to [googleapis.com](https://googleapis.com)

**Pull** subscribers: Webhook endpoint that can accept HTTPS POST requests


# Push Subscriptions



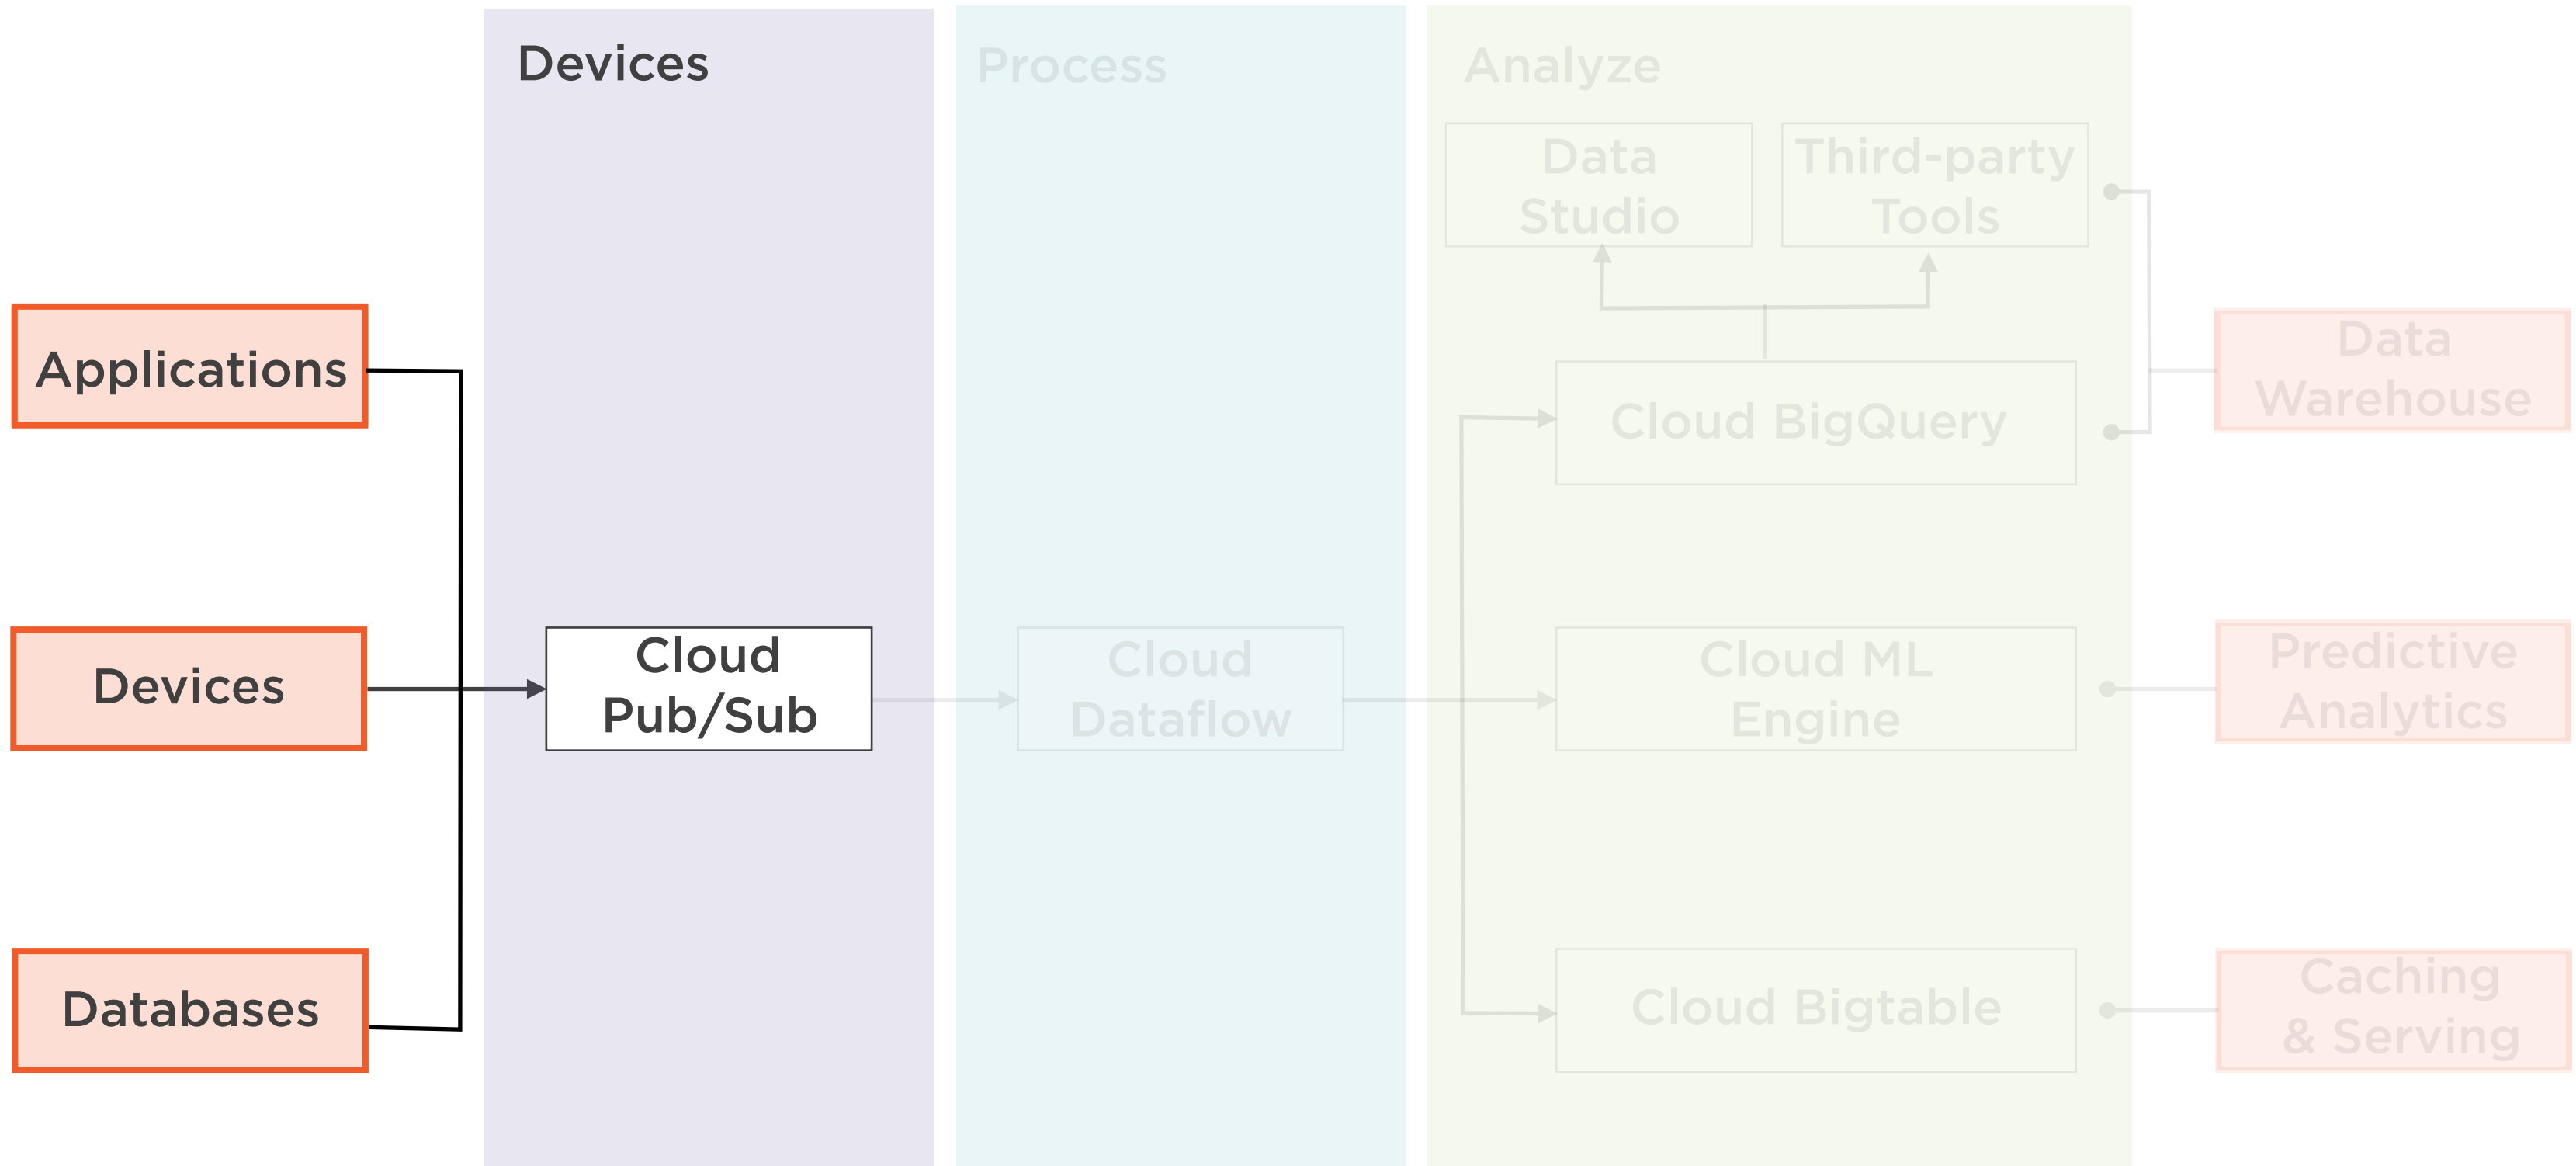
# Pull Subscriptions



# Streaming Analytics

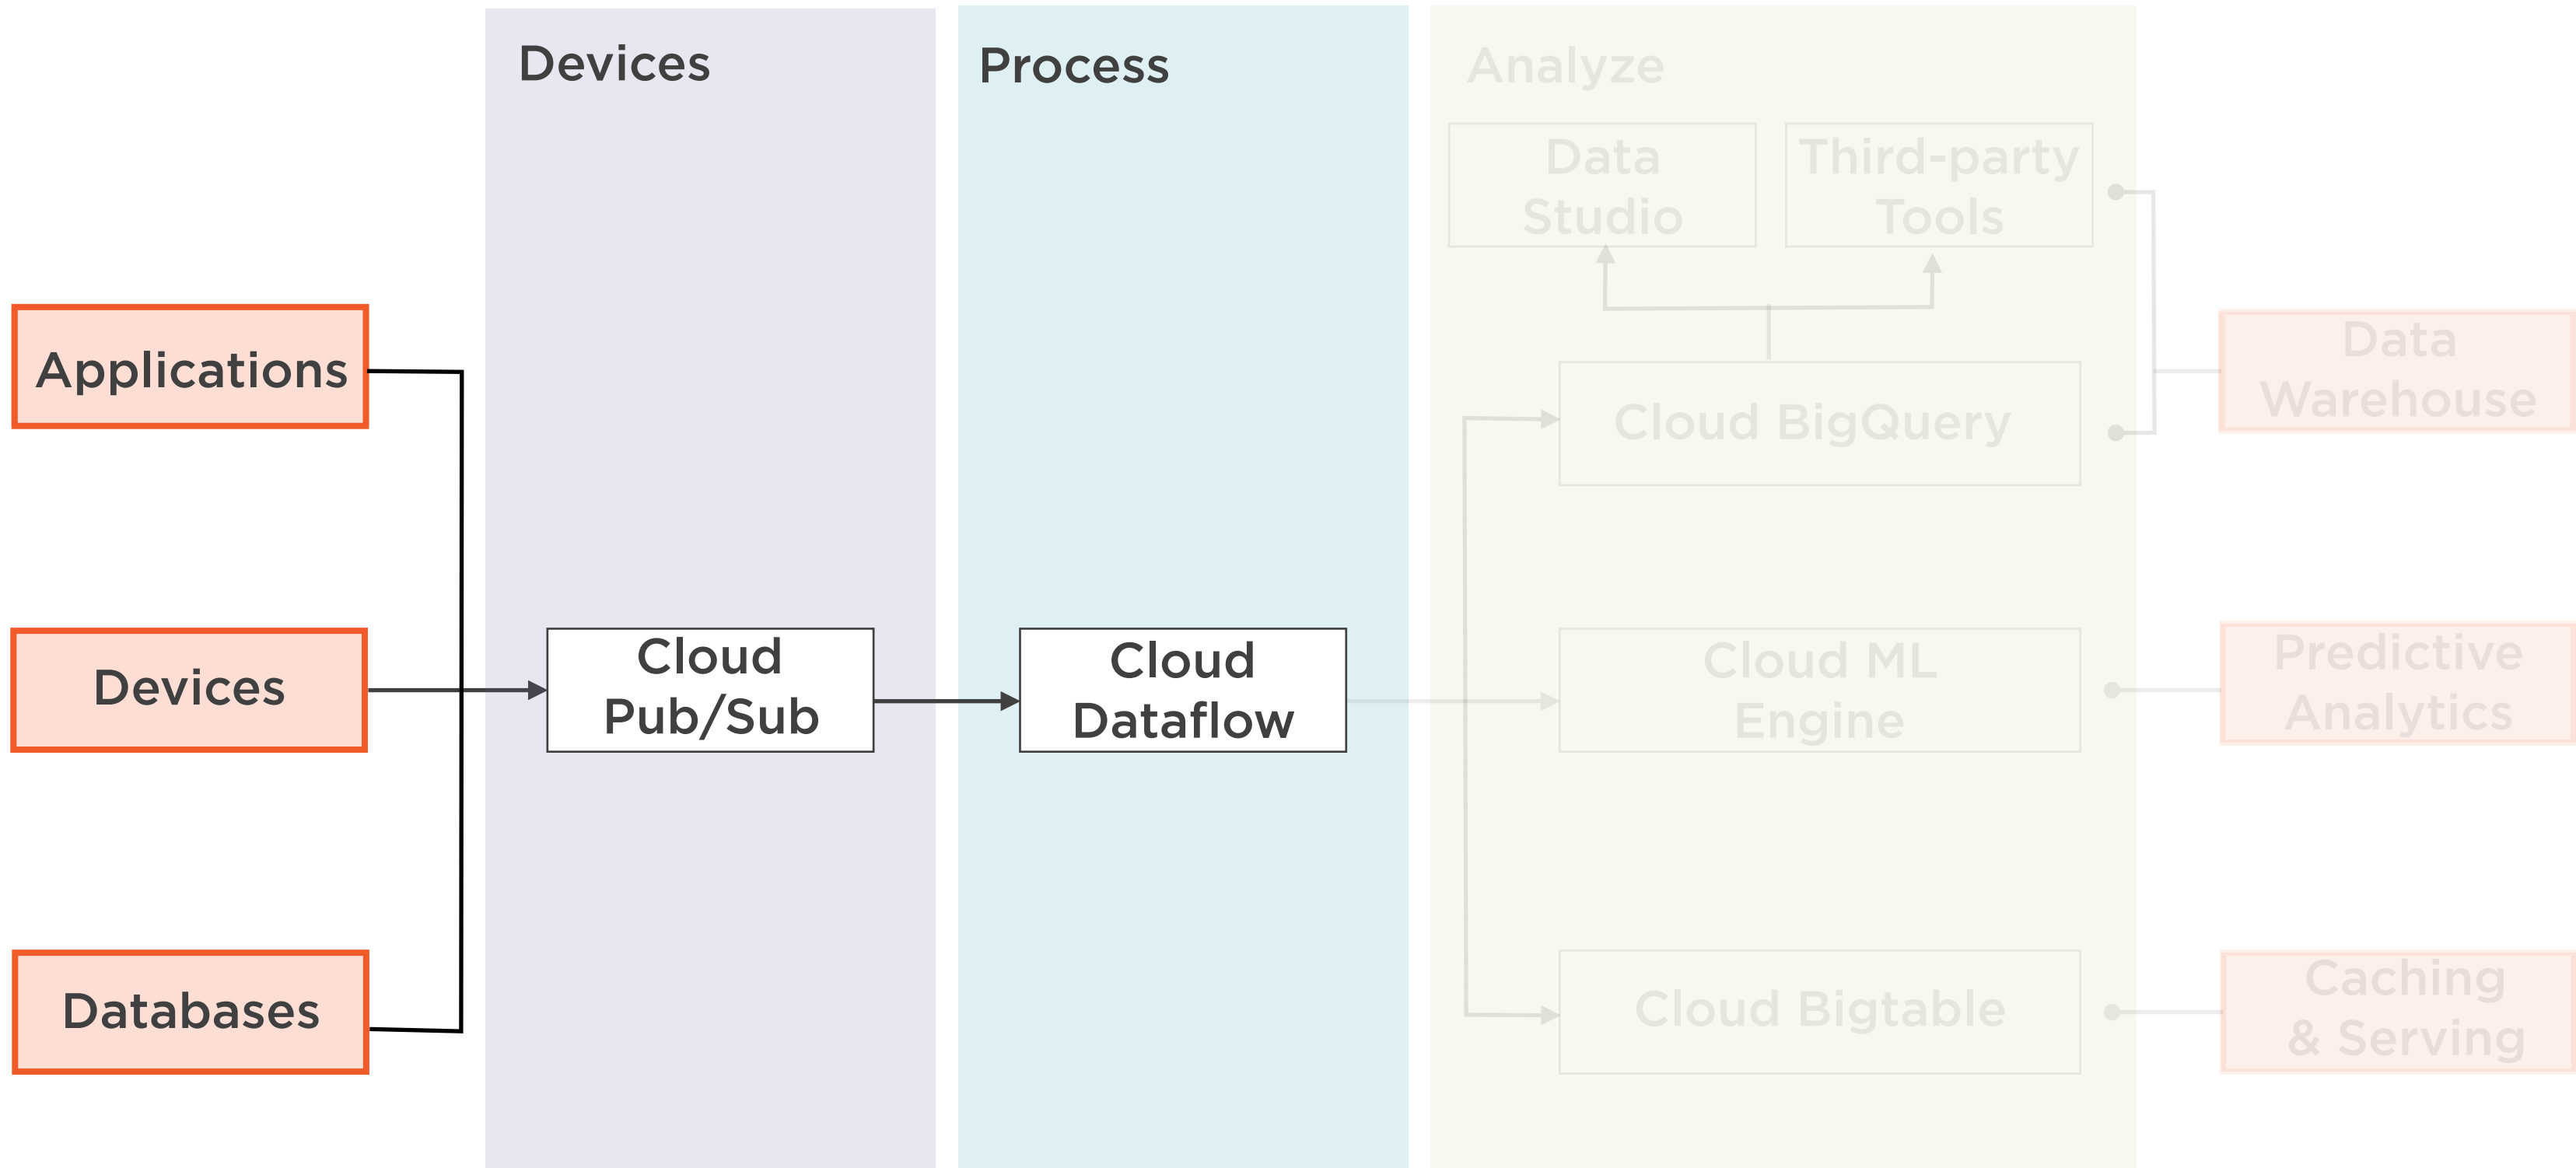


# Stream Into Pub/Sub

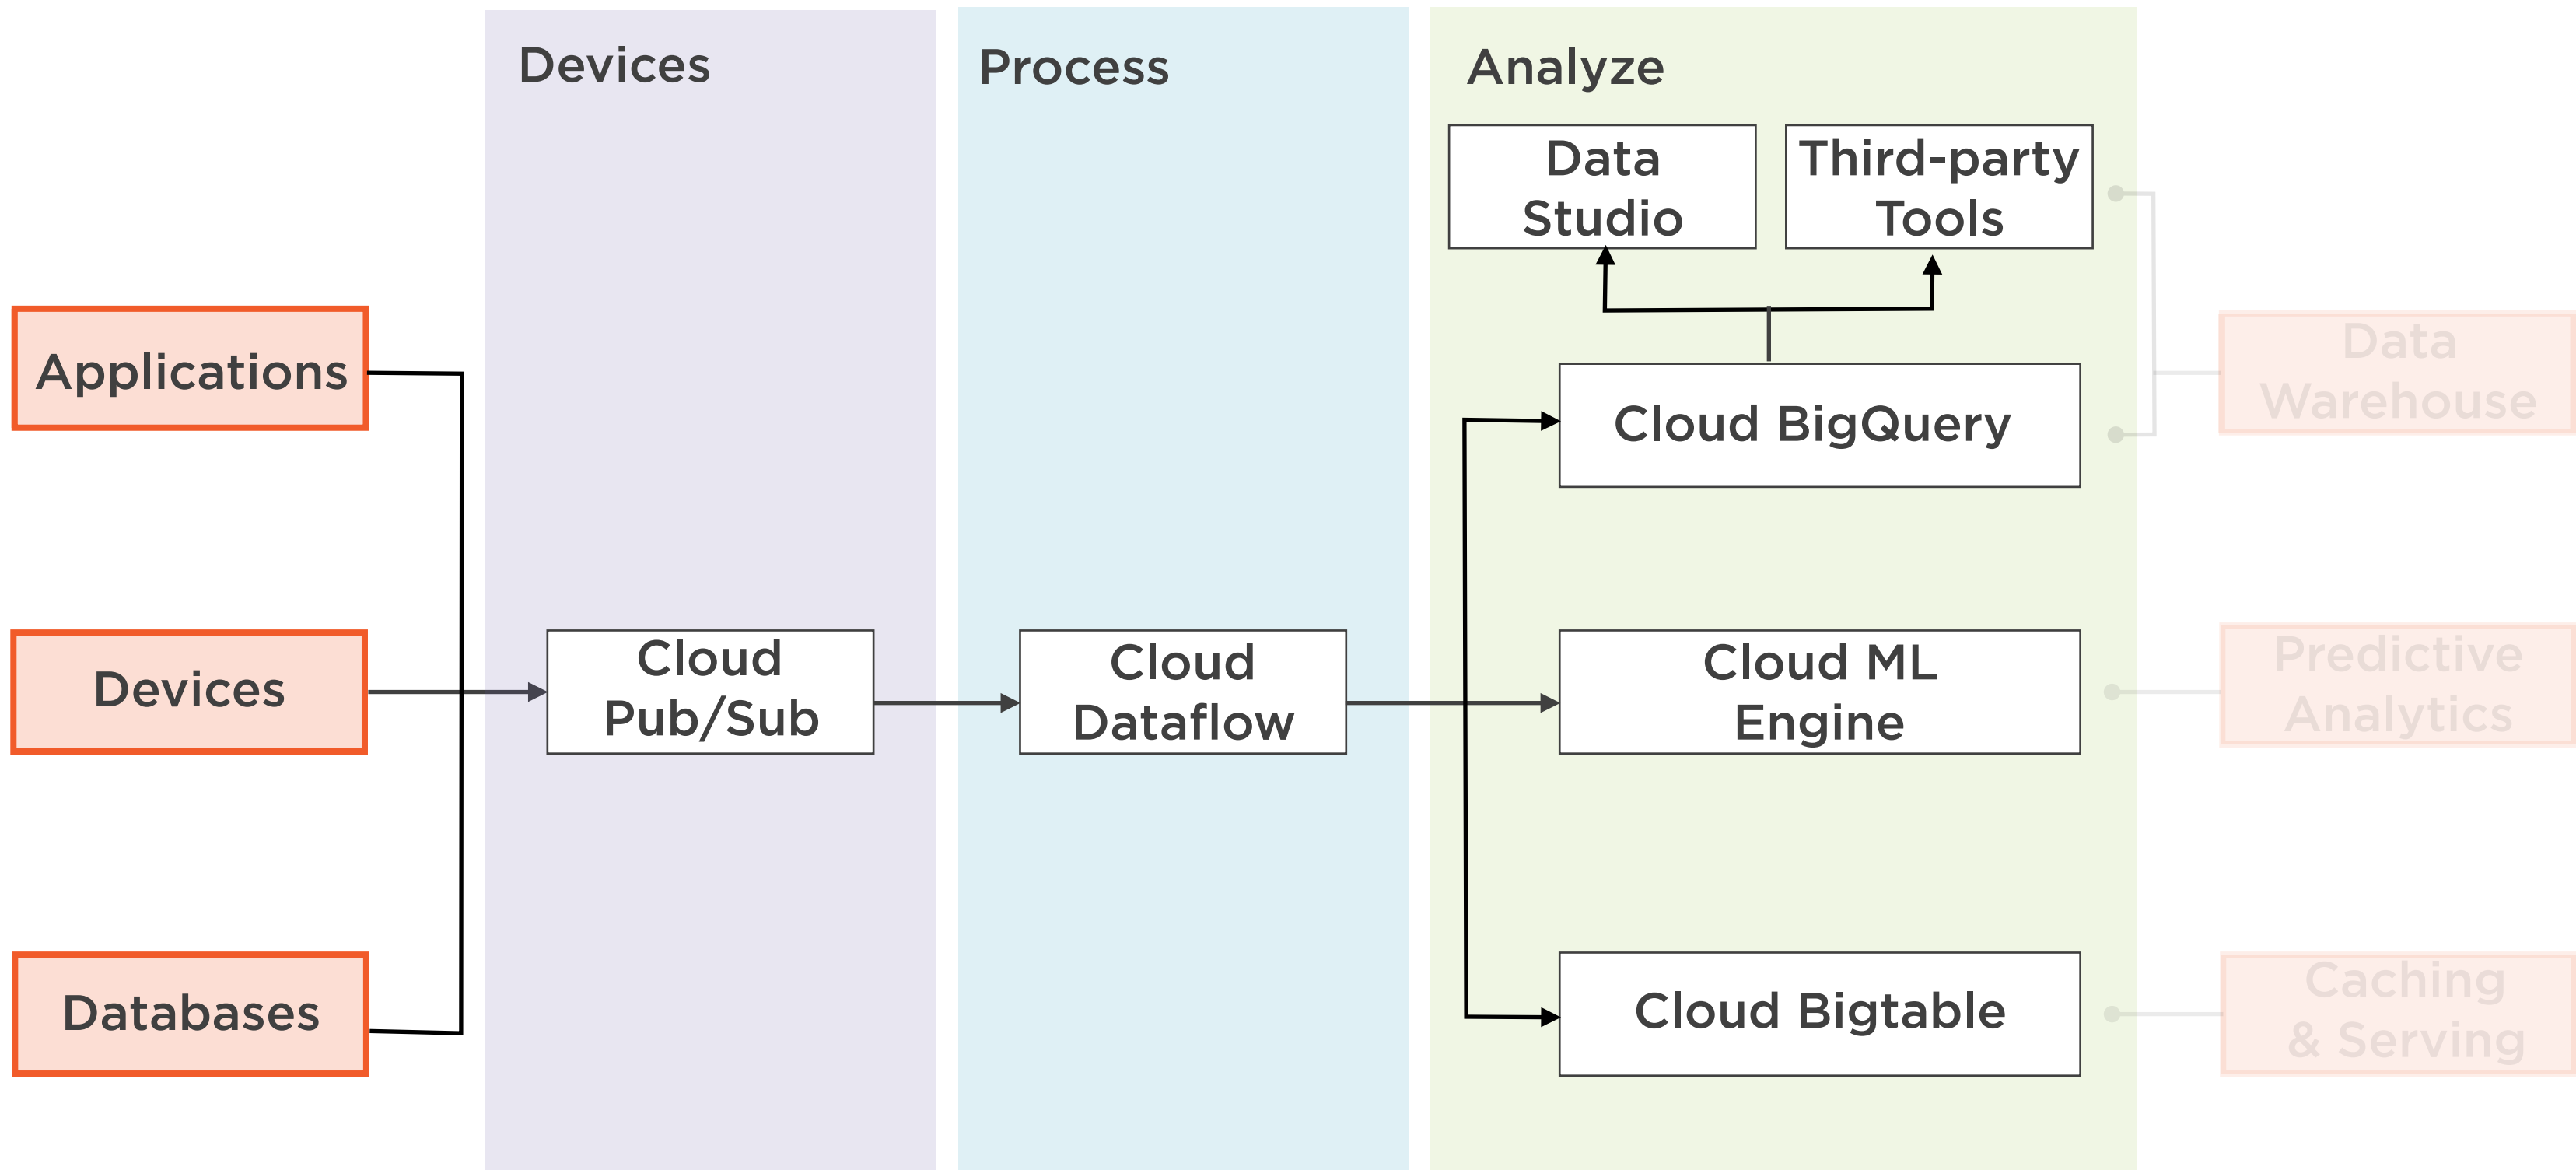




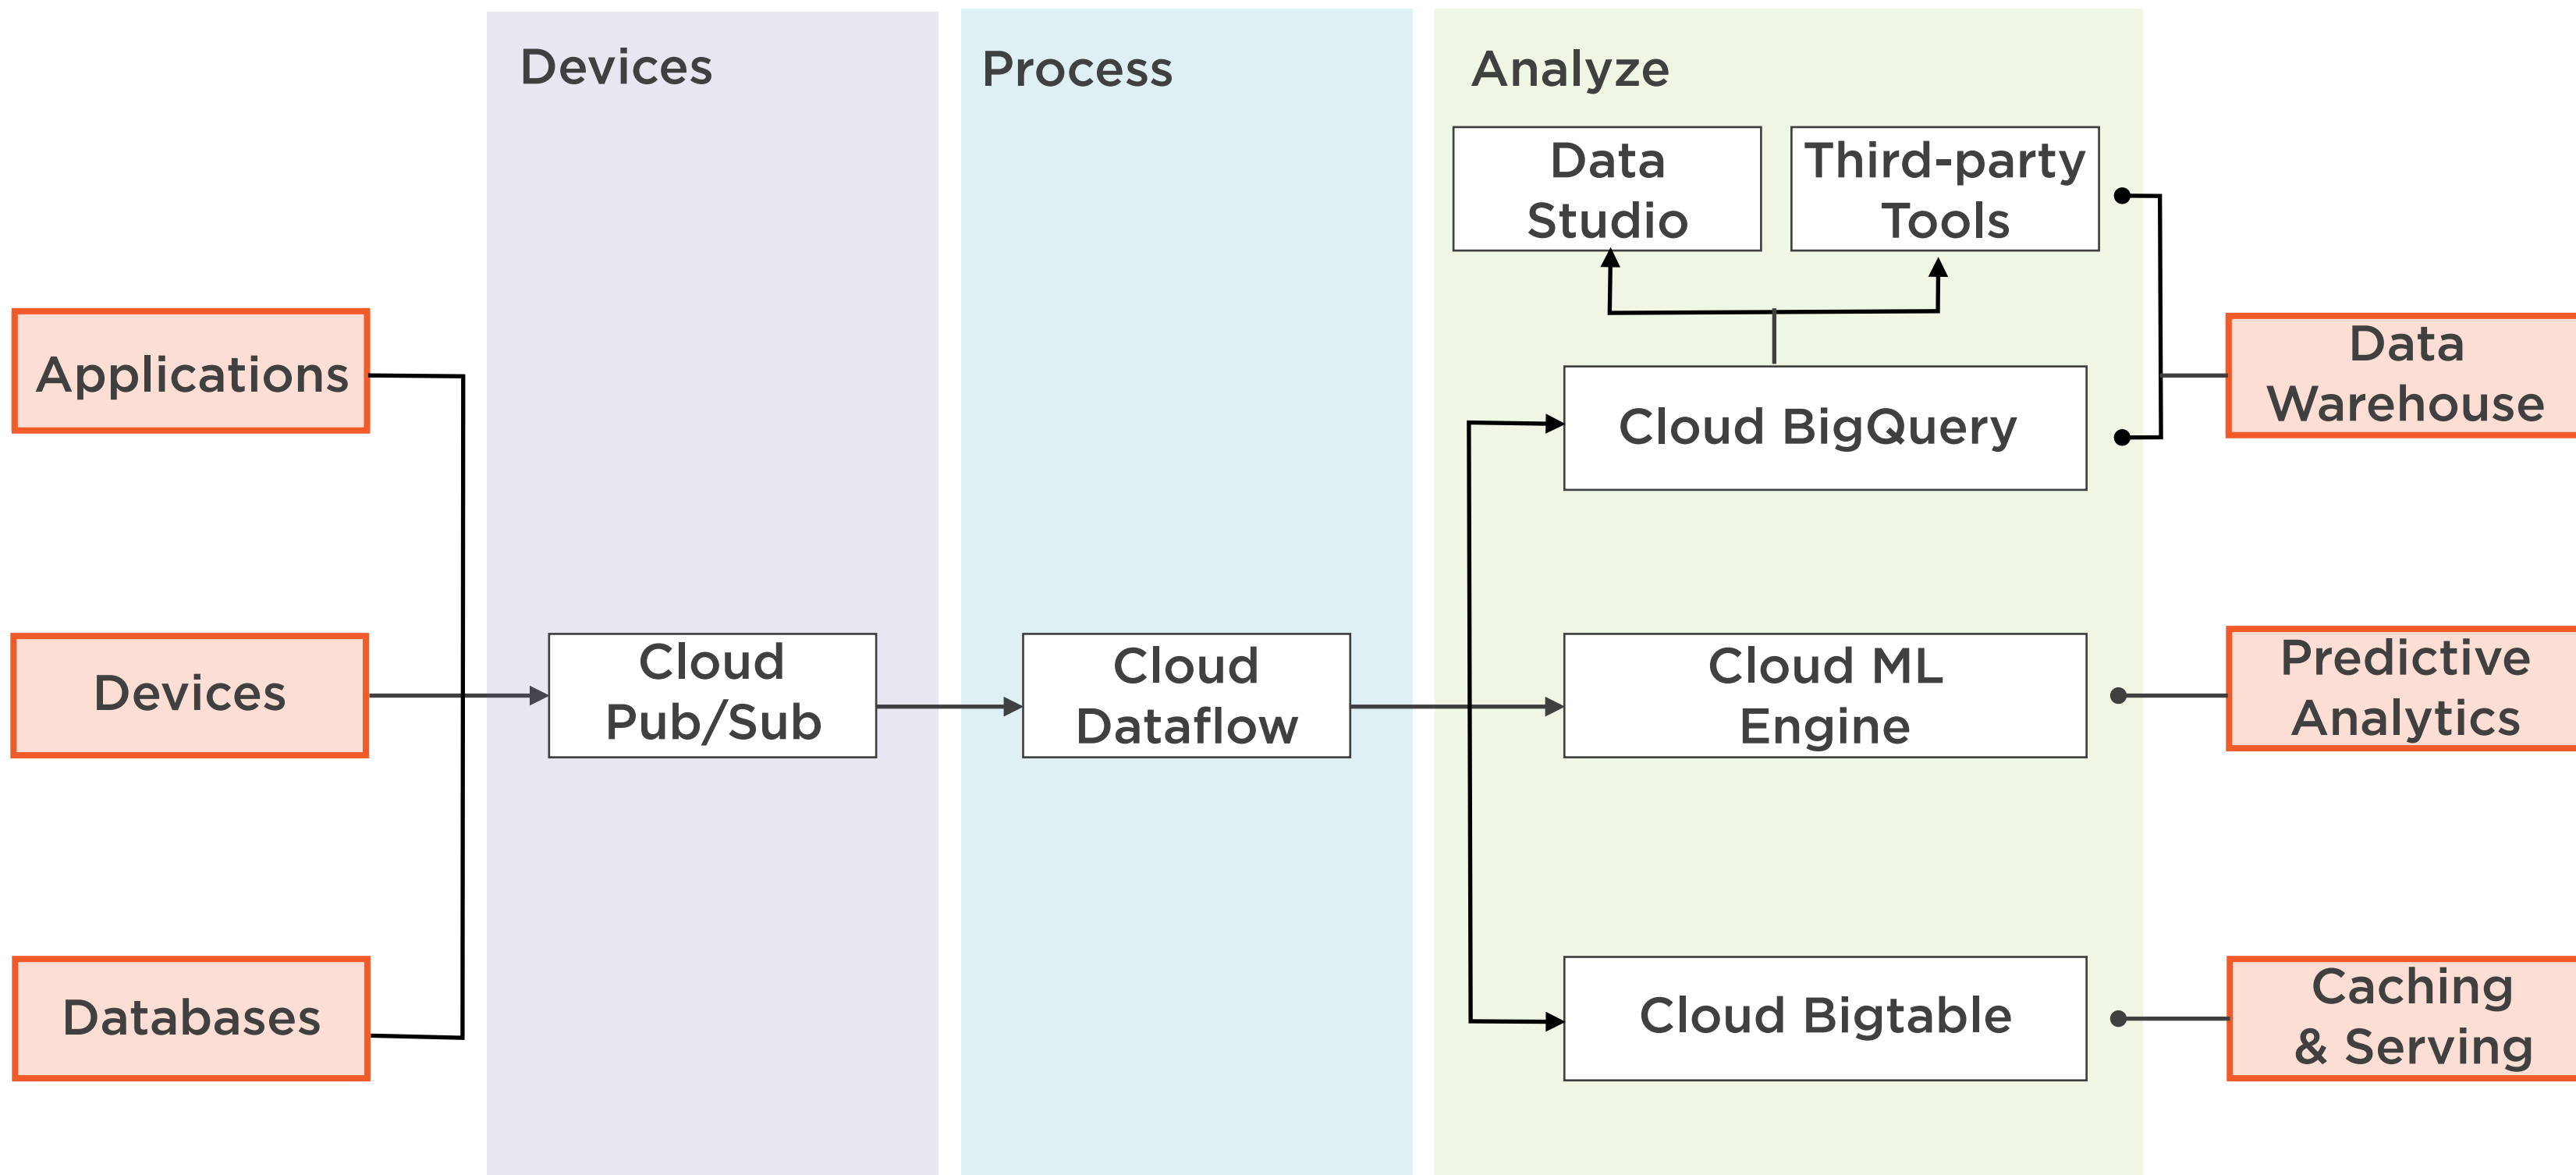
# Transform Using Dataflow



# Analyze and Visualize Data



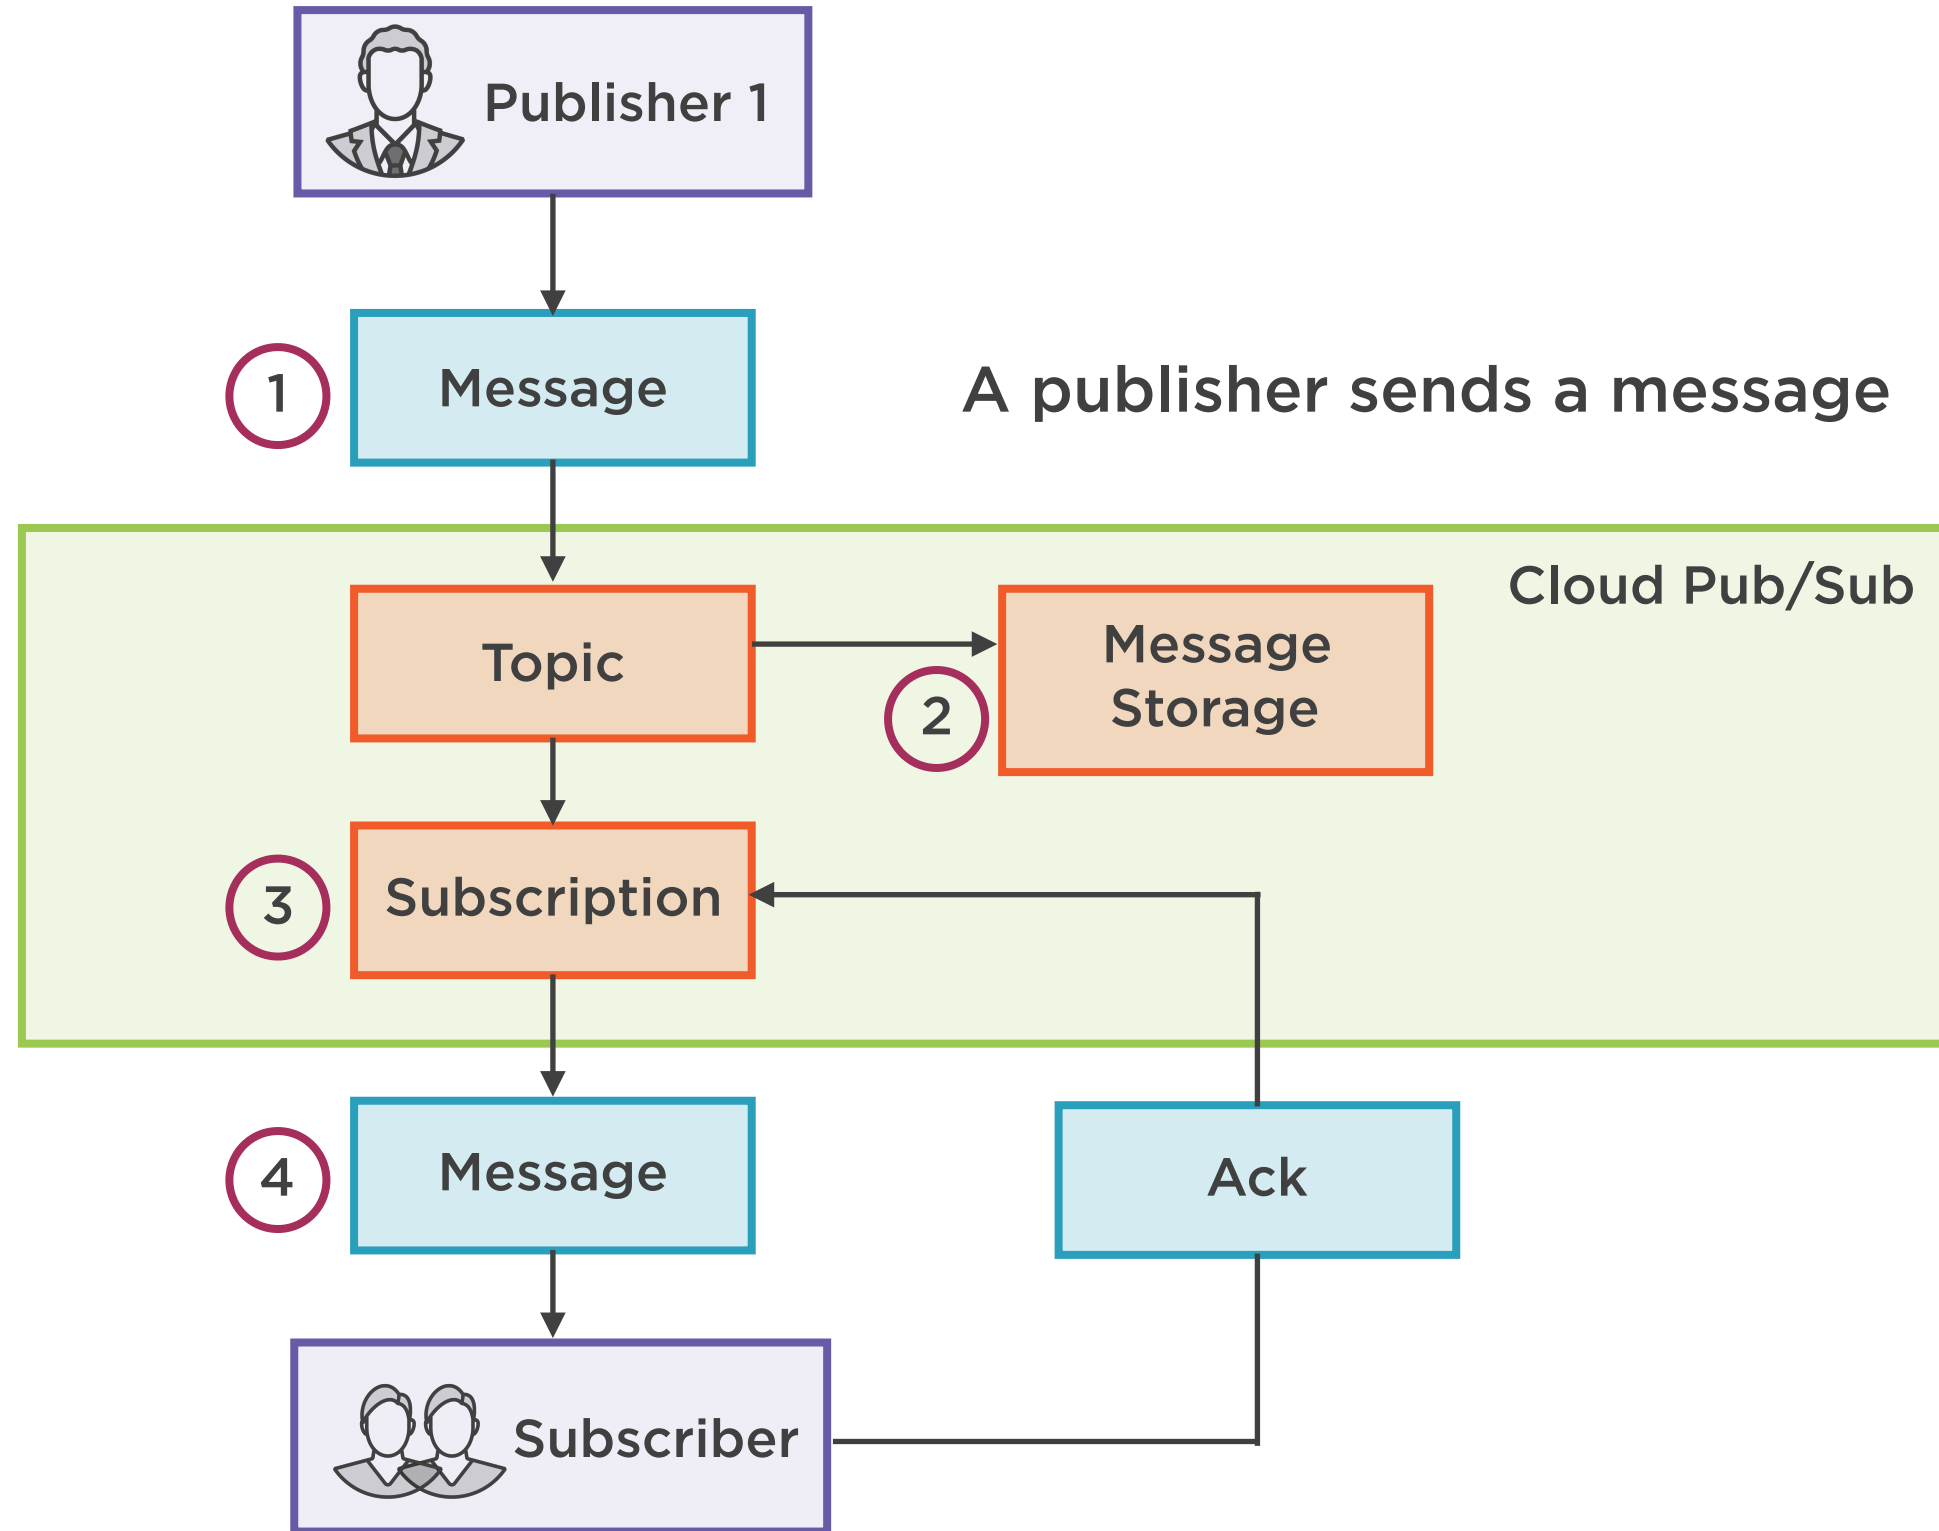
# Store in a Data Warehouse



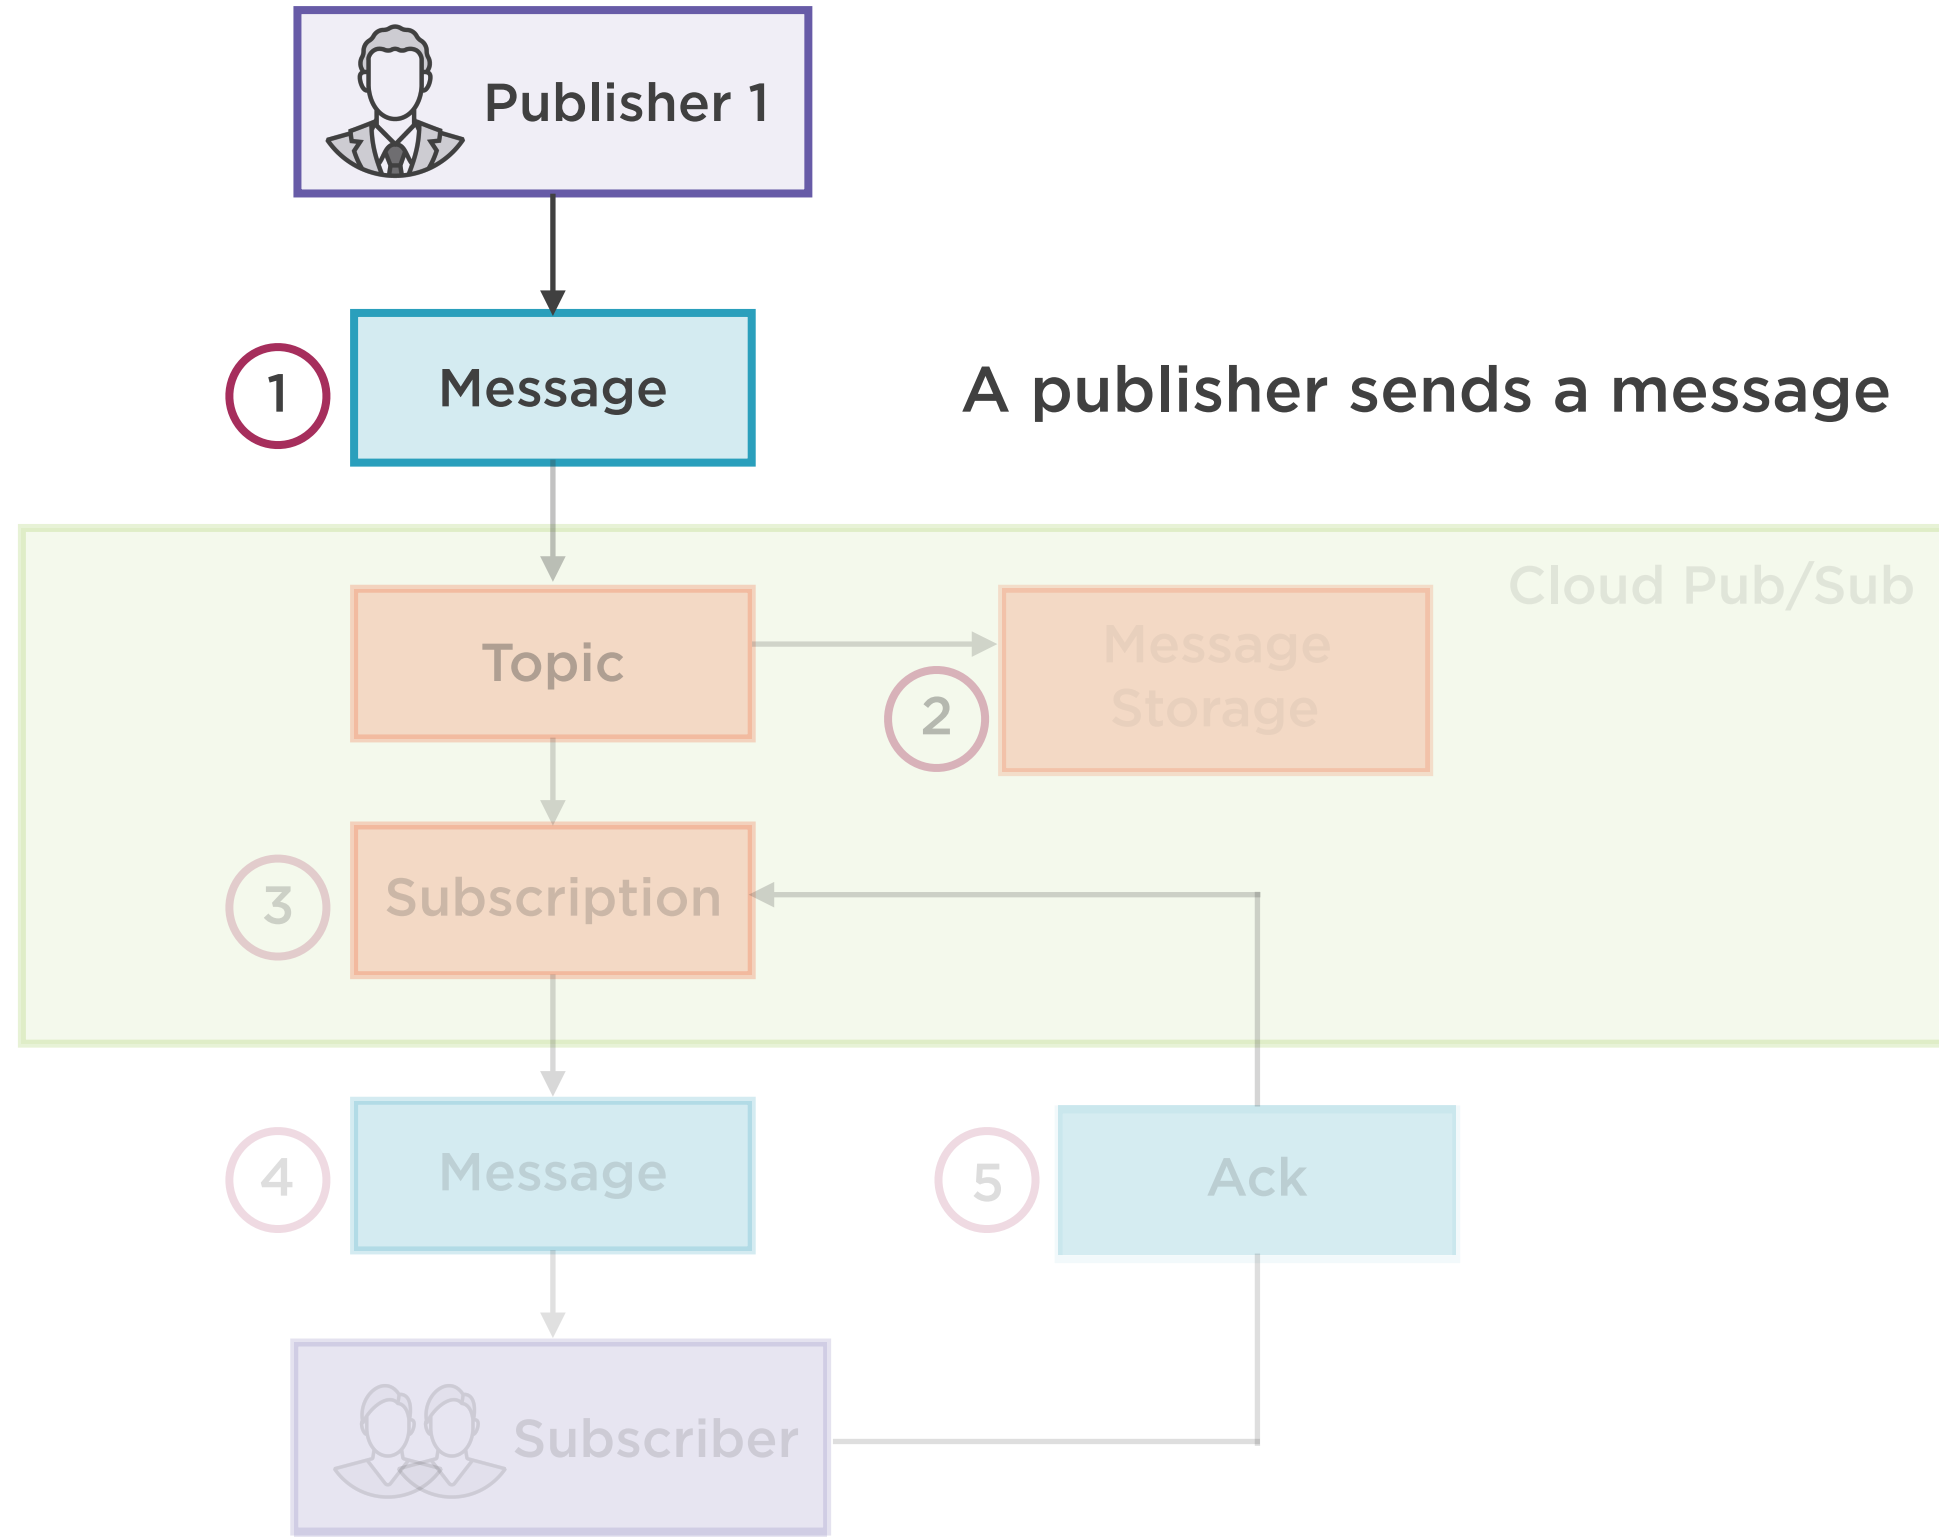
# Pub/Sub Message Flow

---

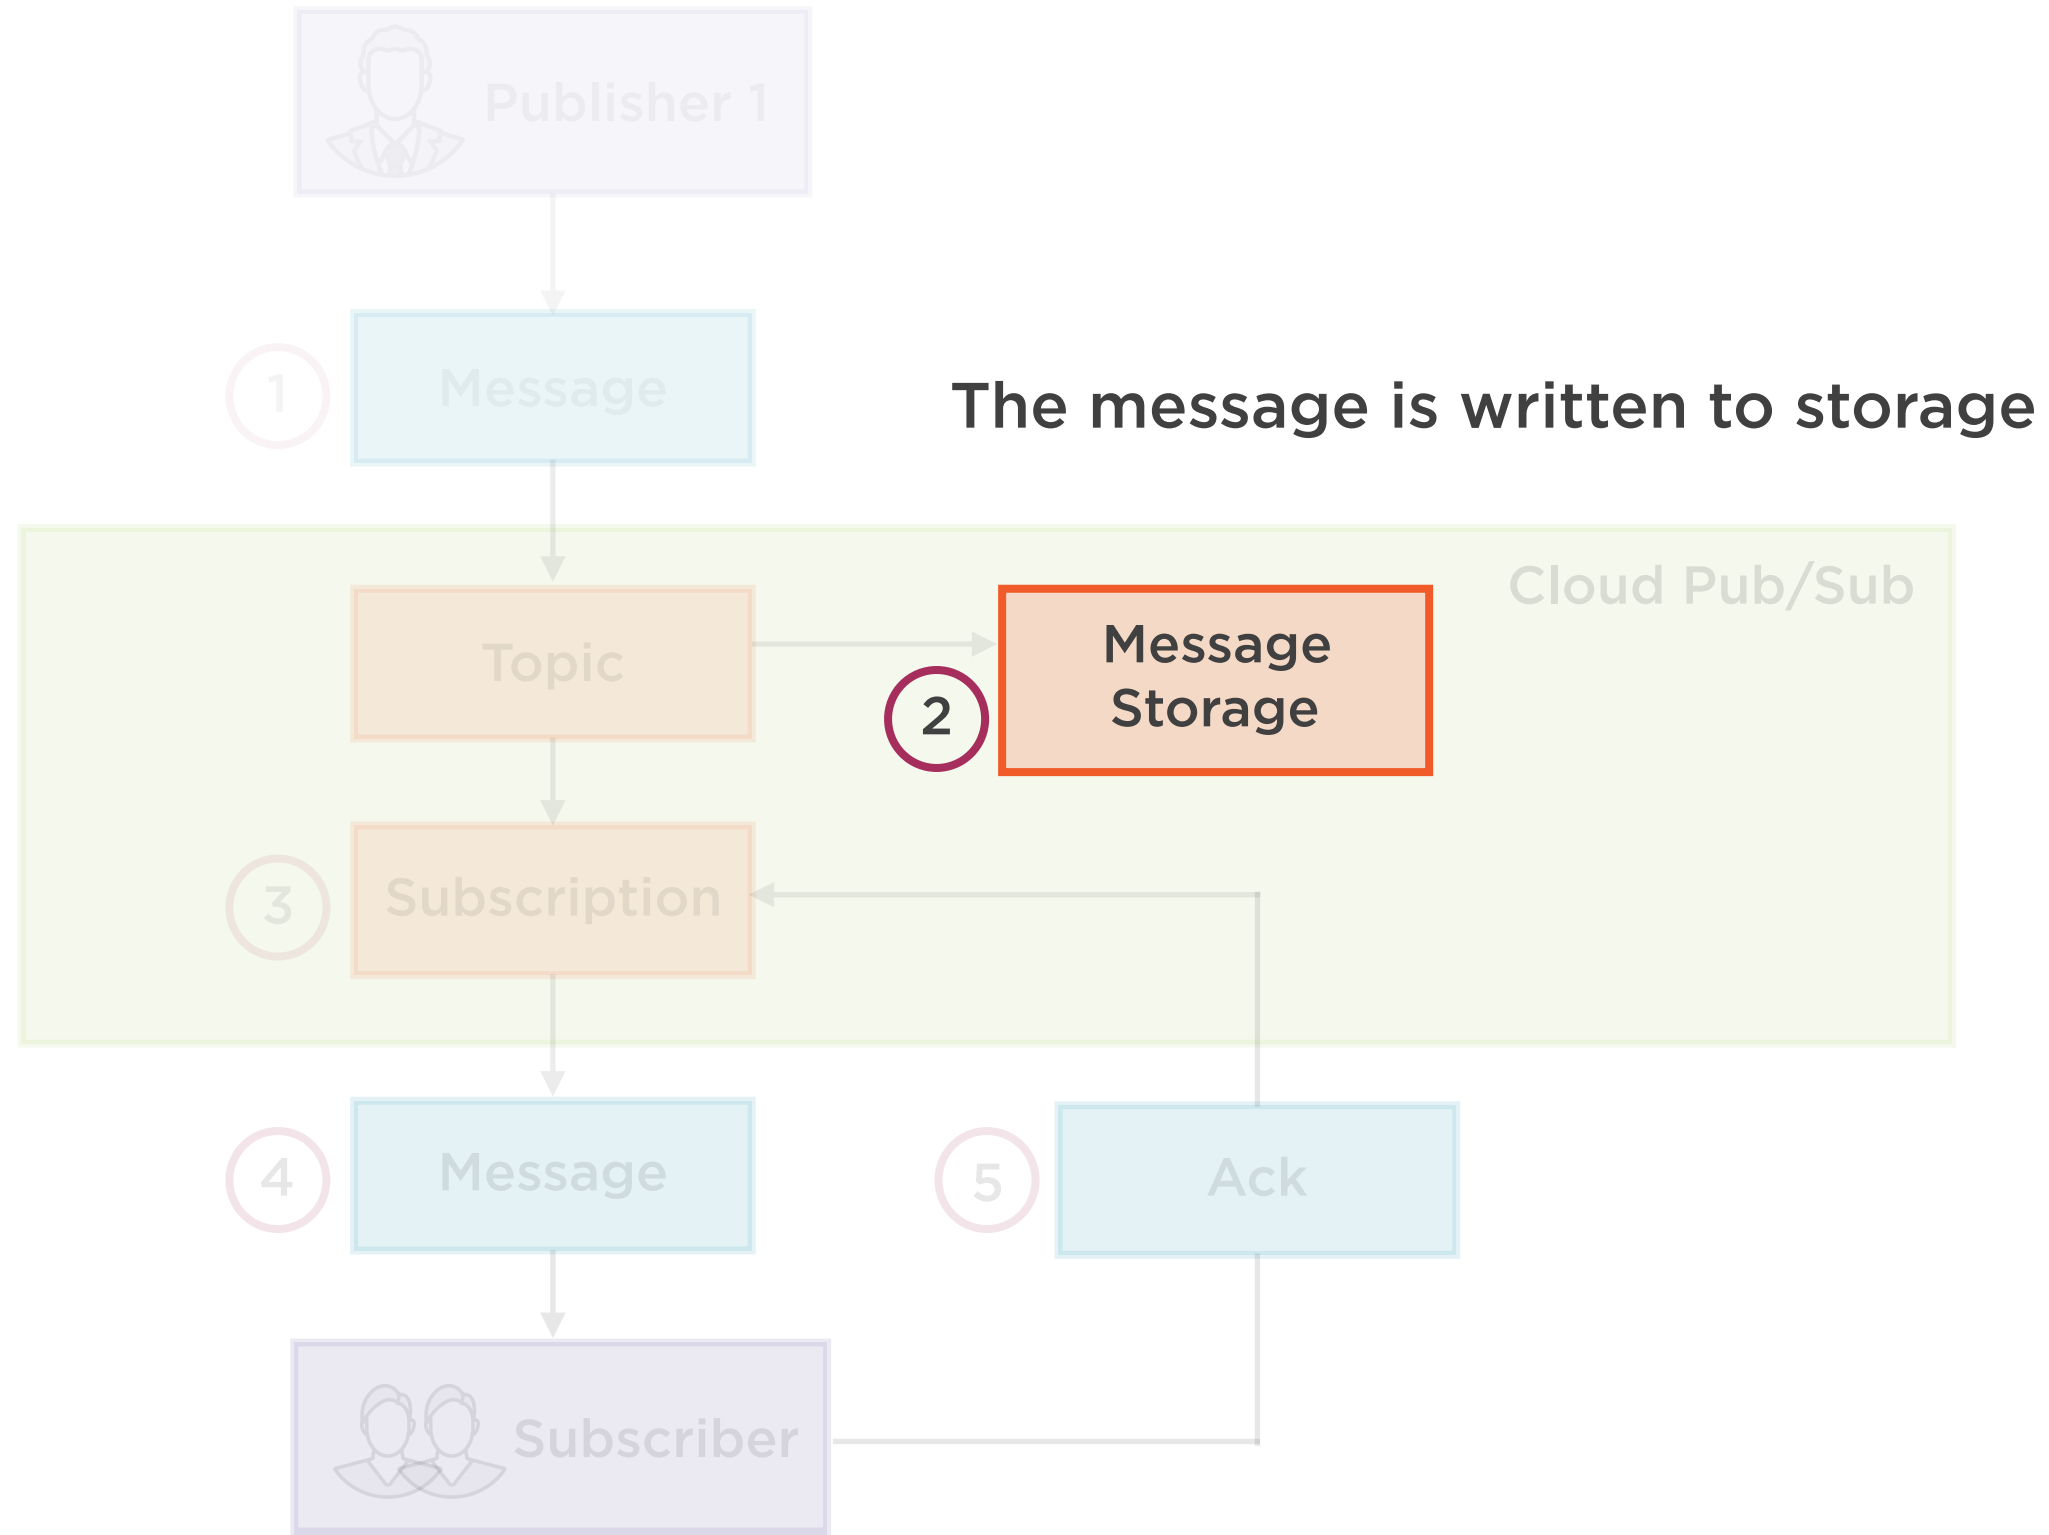
# Pub/Sub Messaging



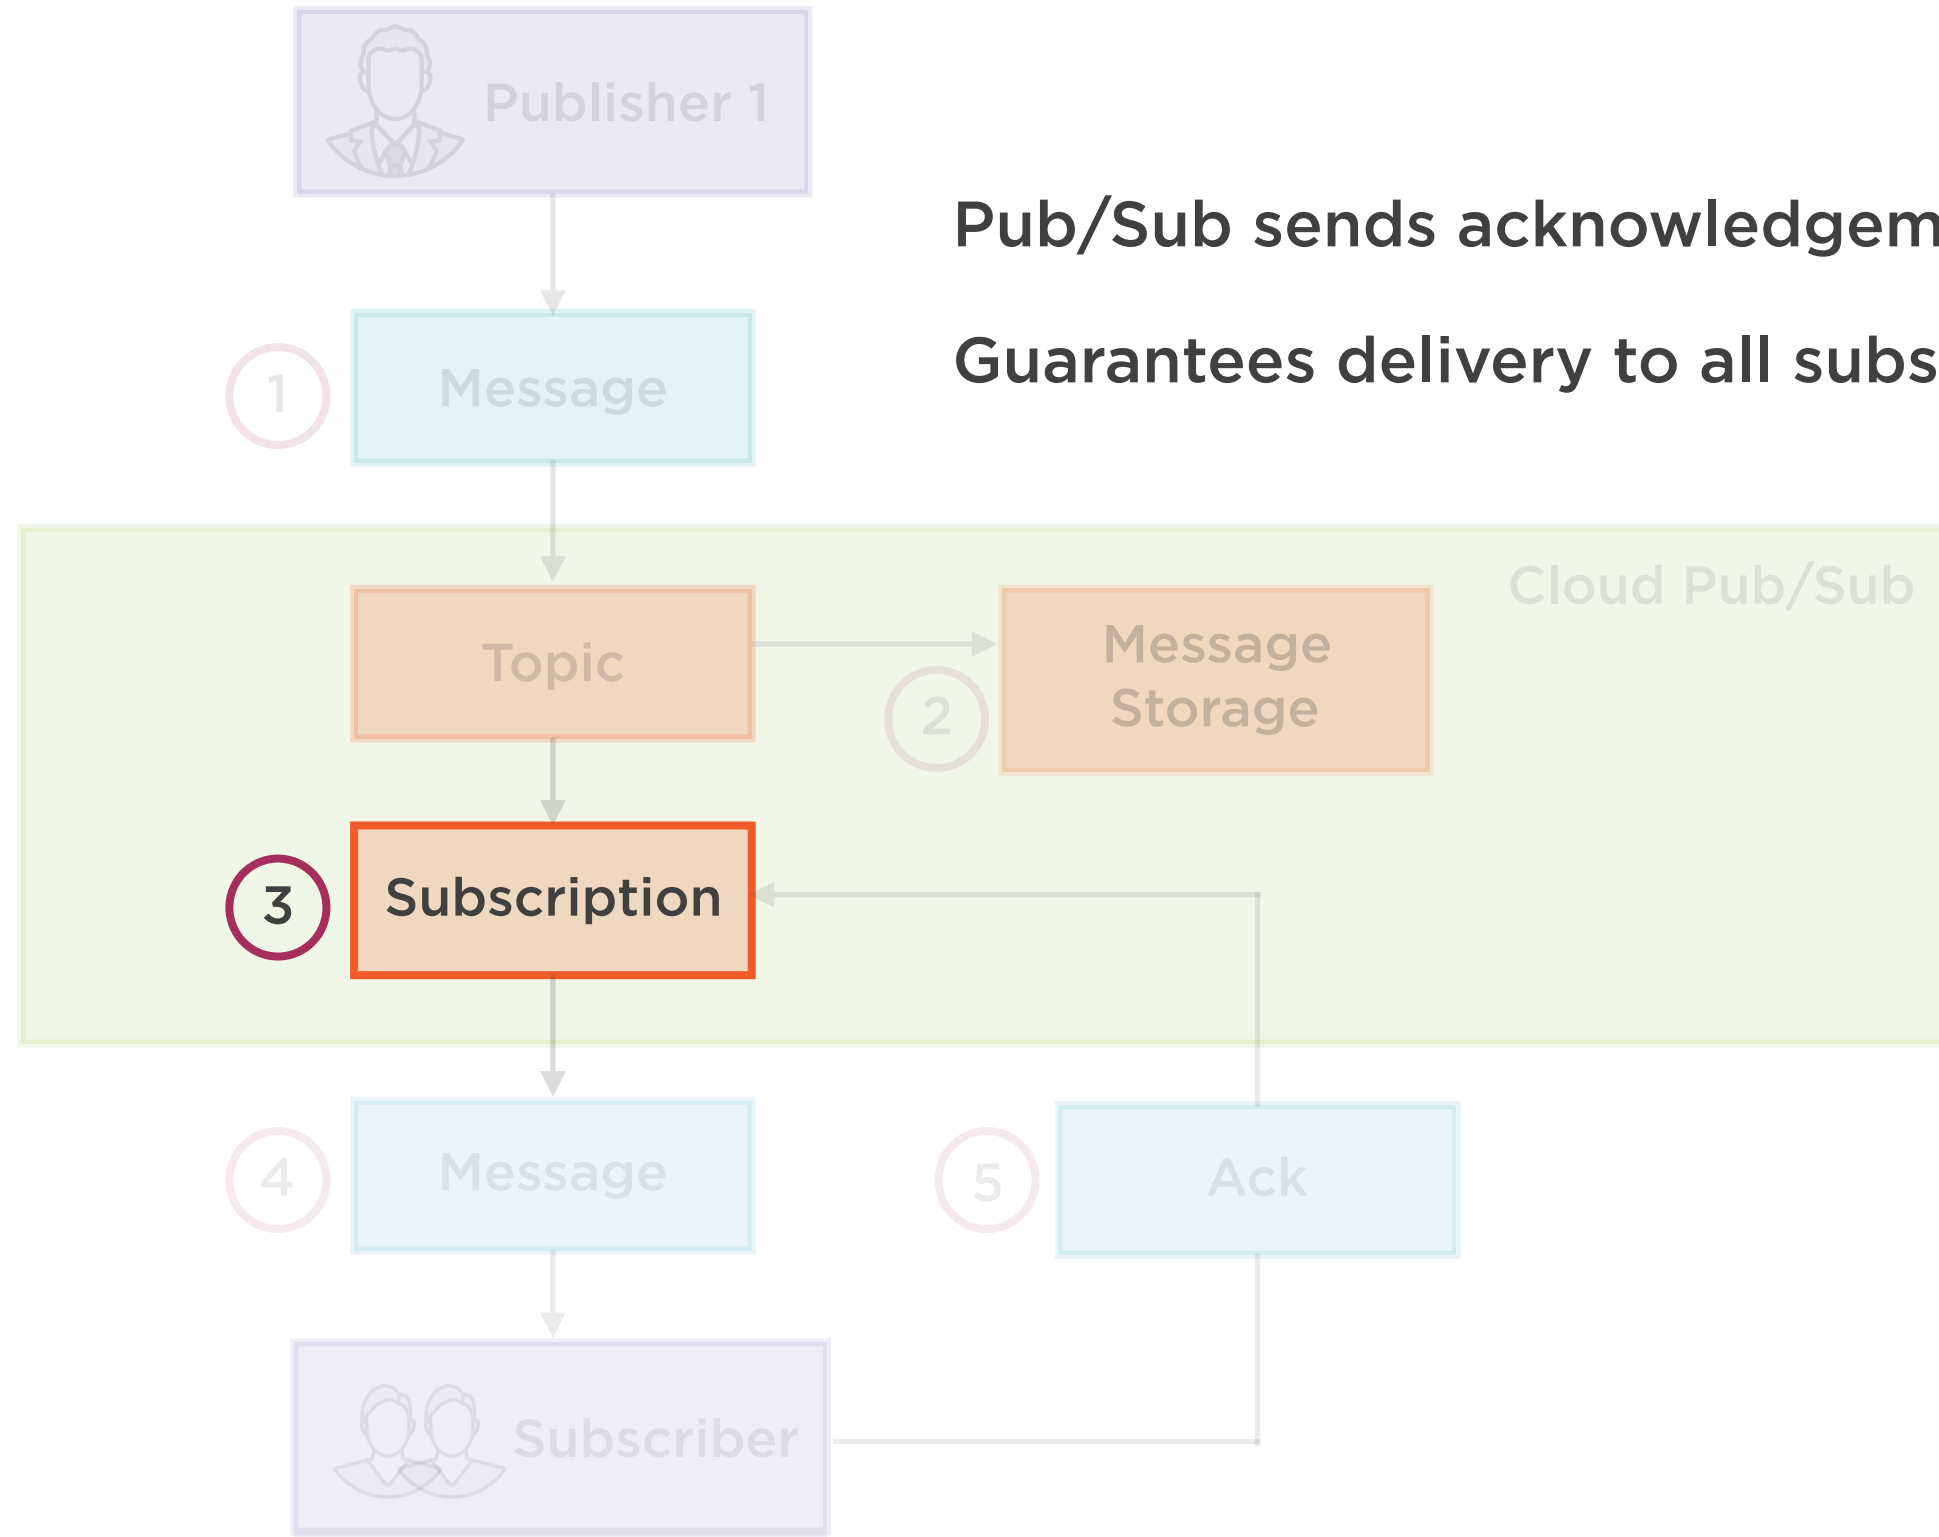
# Pub/Sub Messaging



# Pub/Sub Messaging



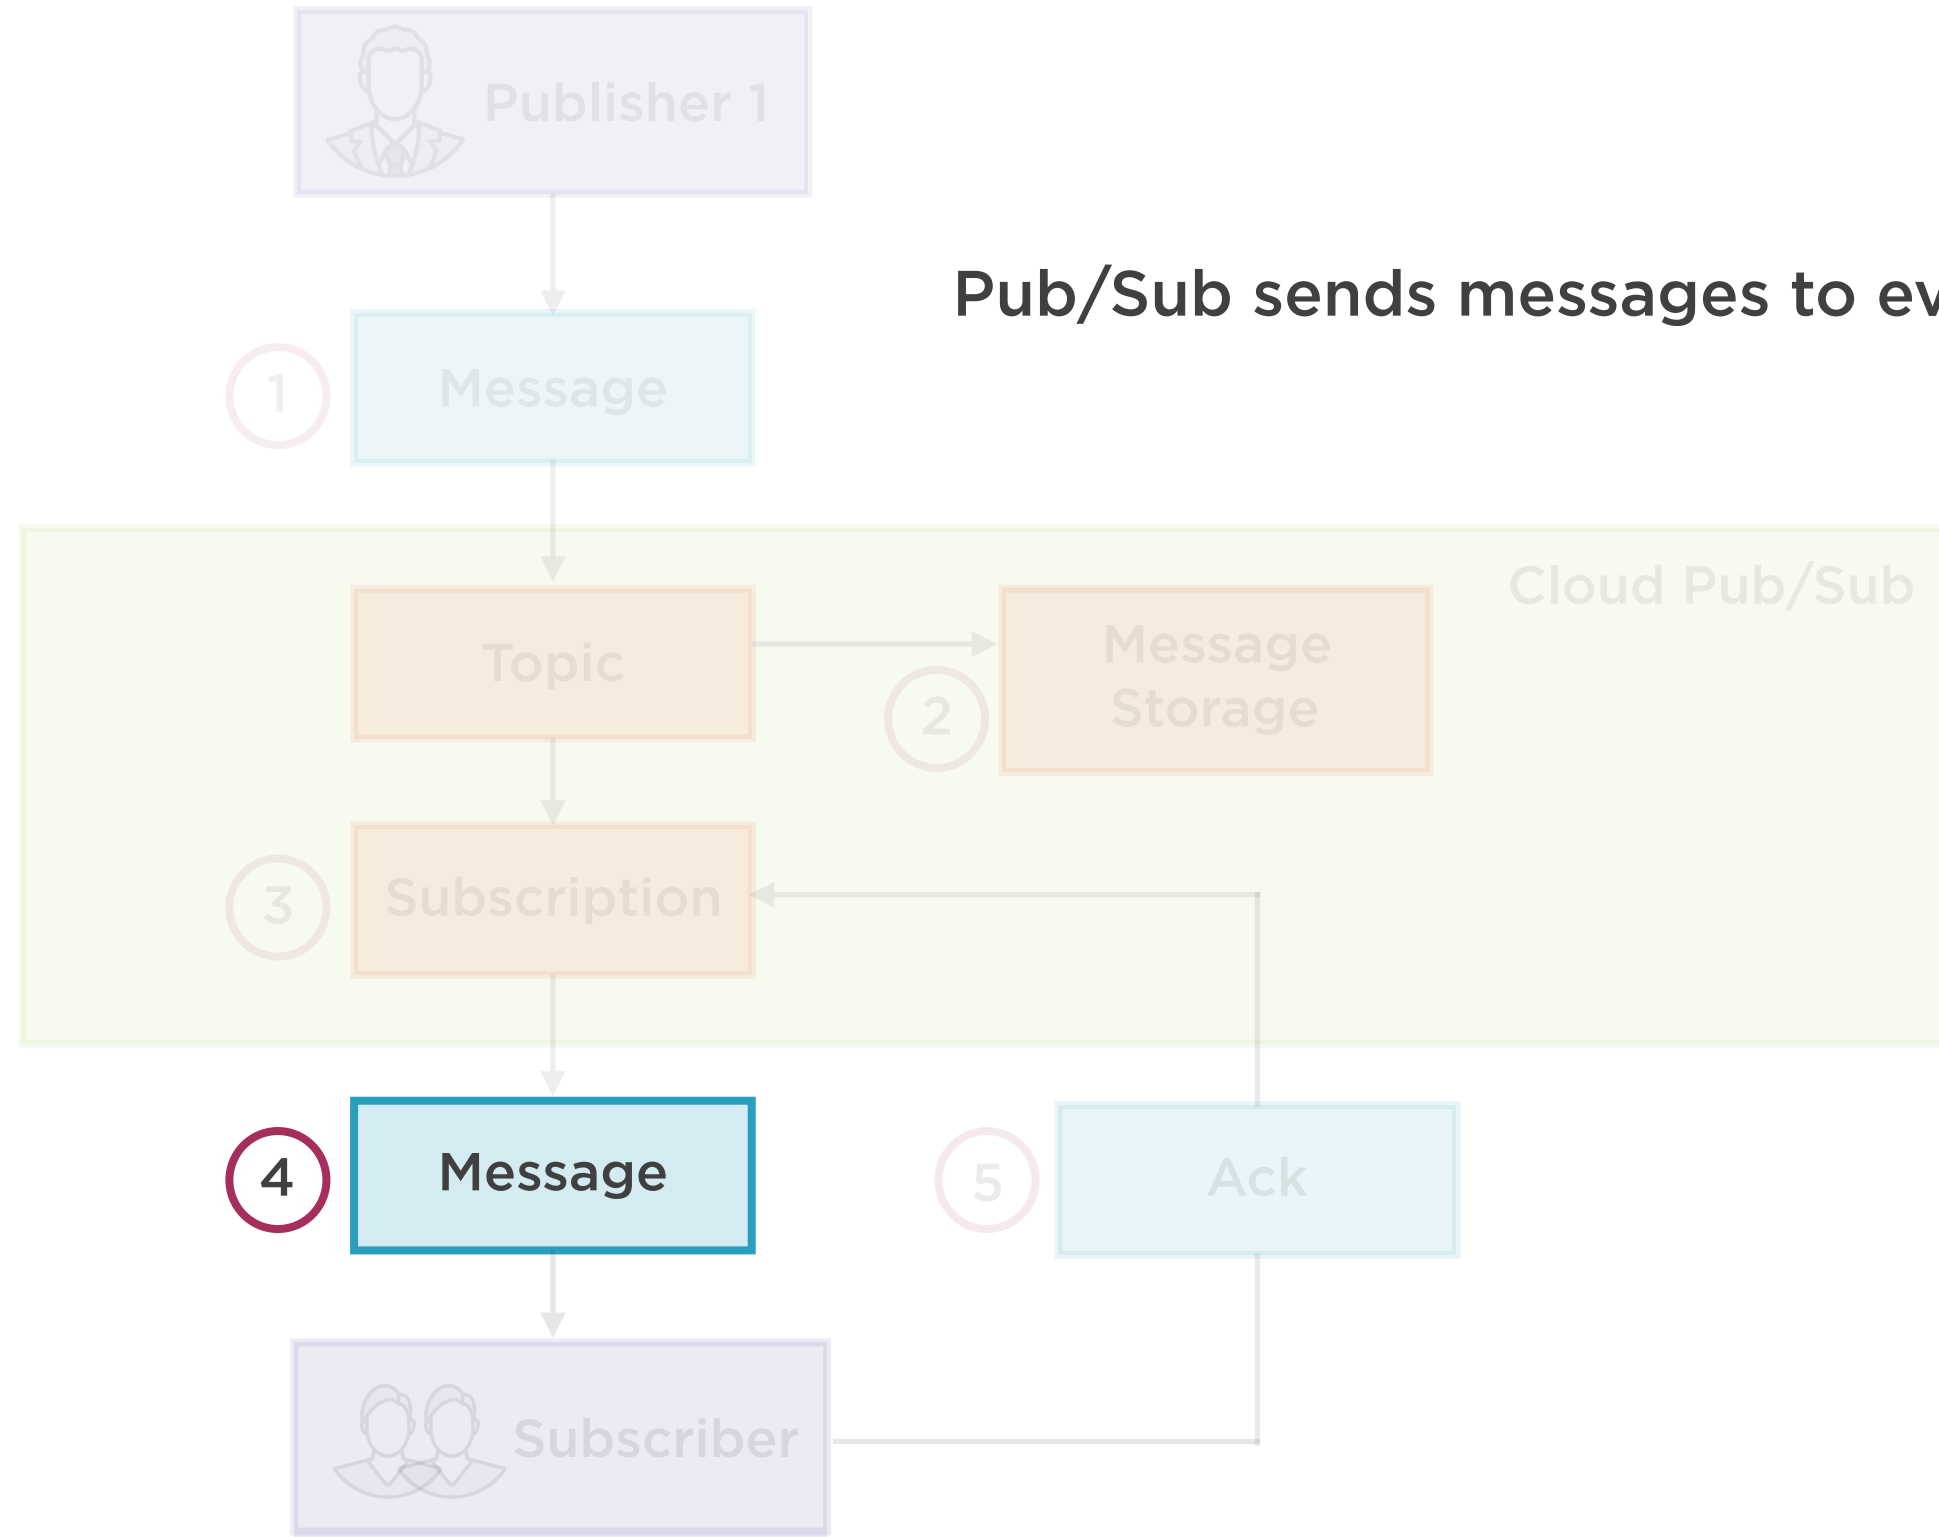
# Pub/Sub Messaging



**Pub/Sub sends acknowledgement to publisher**  
**Guarantees delivery to all subscribers**

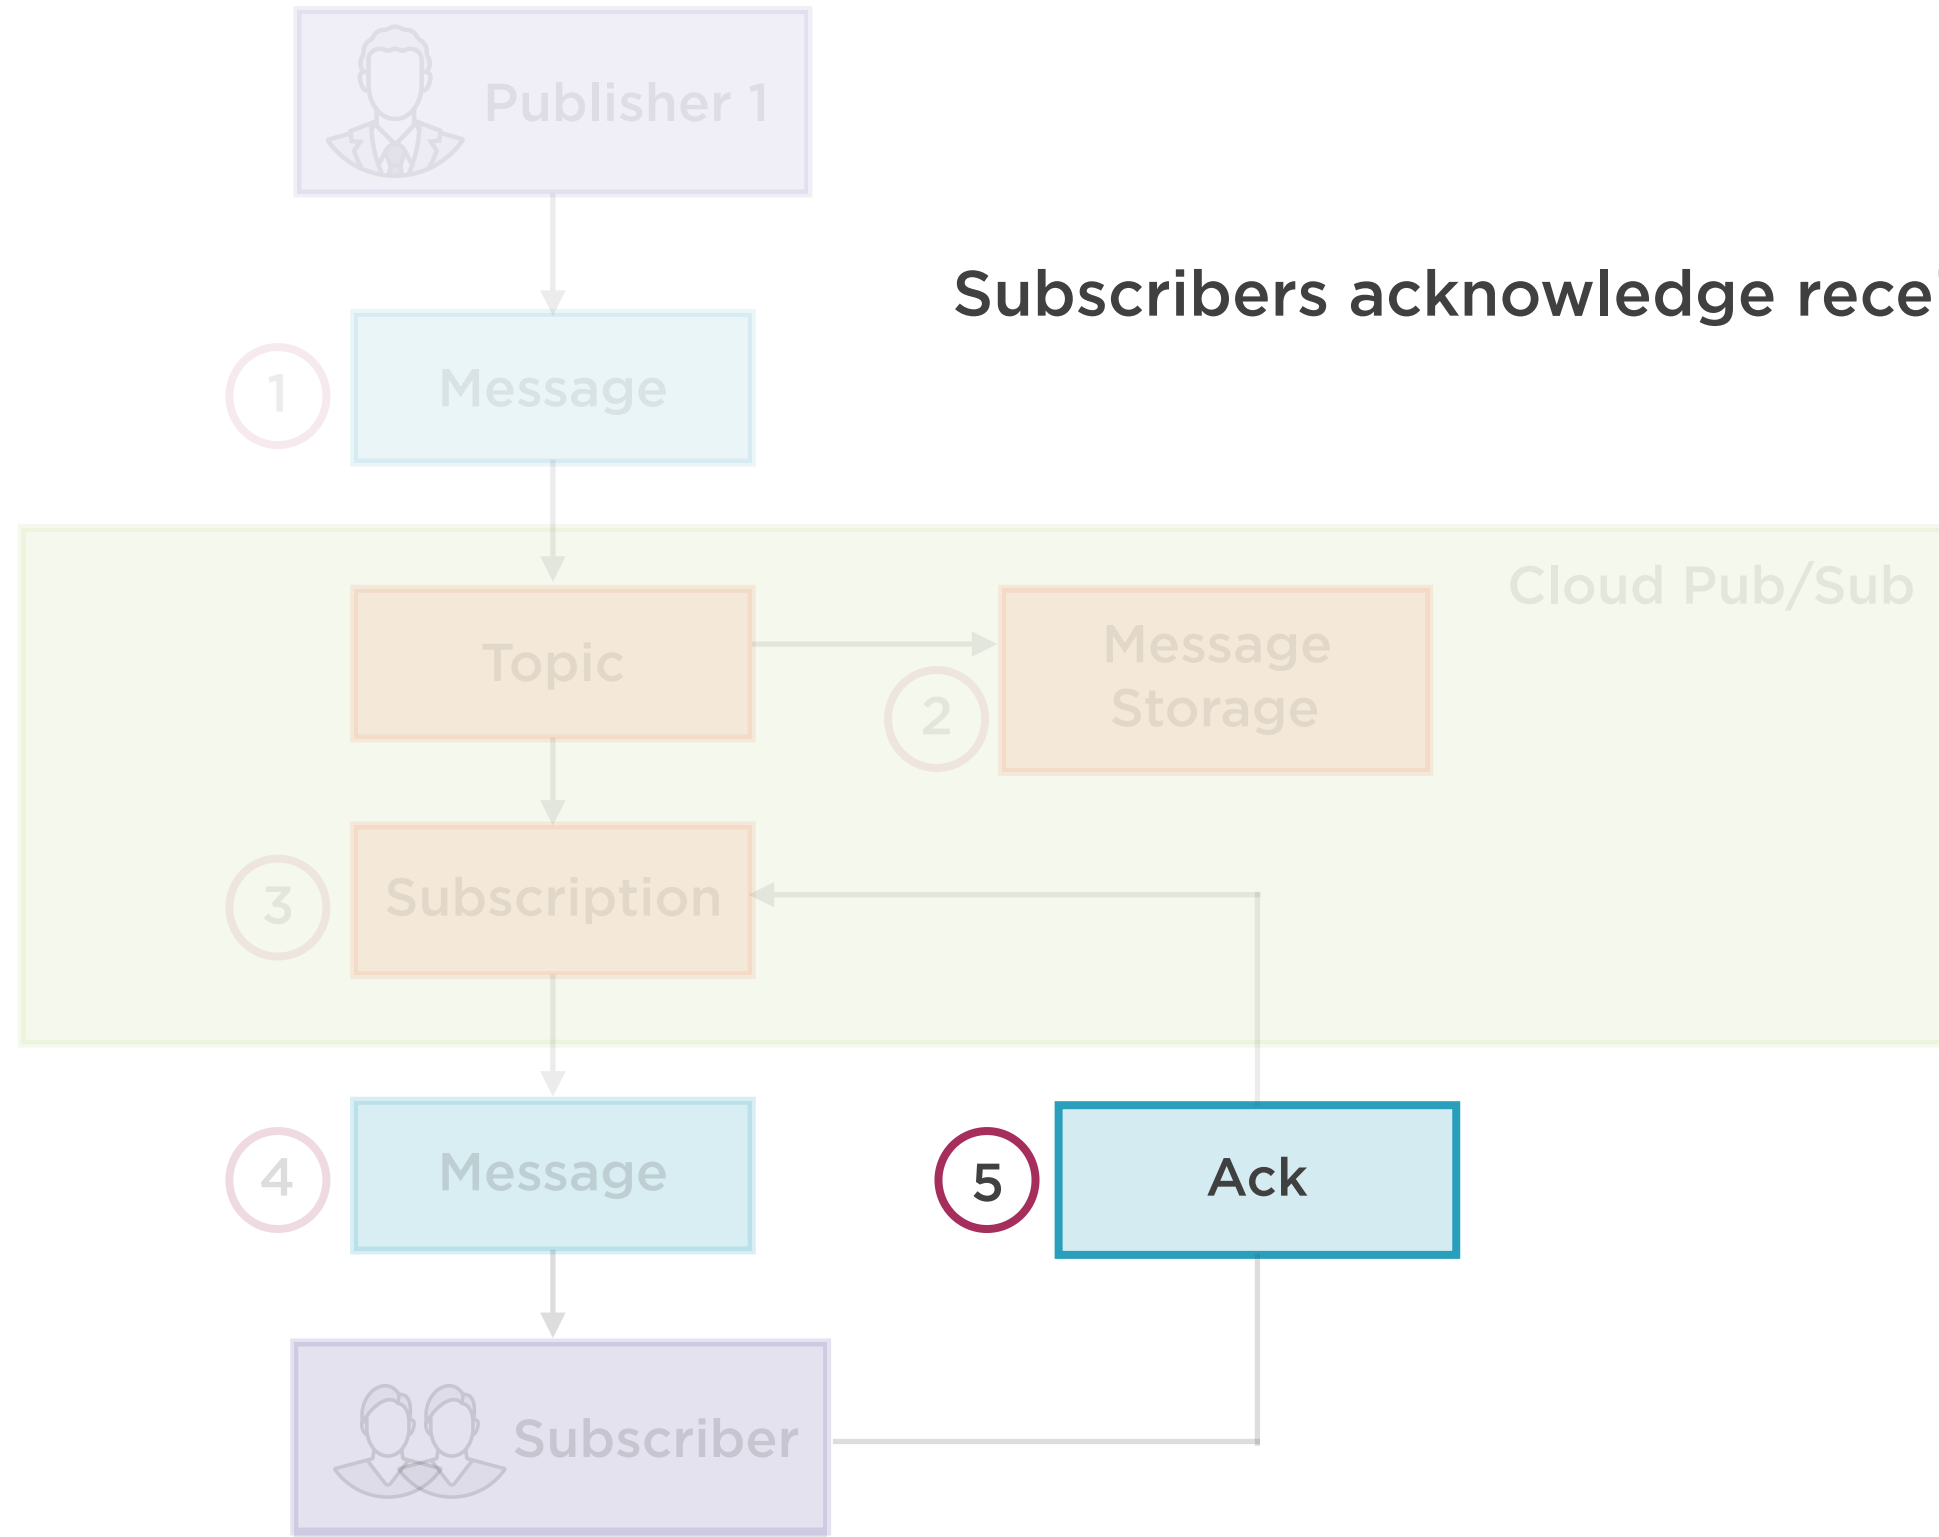


# Pub/Sub Messaging



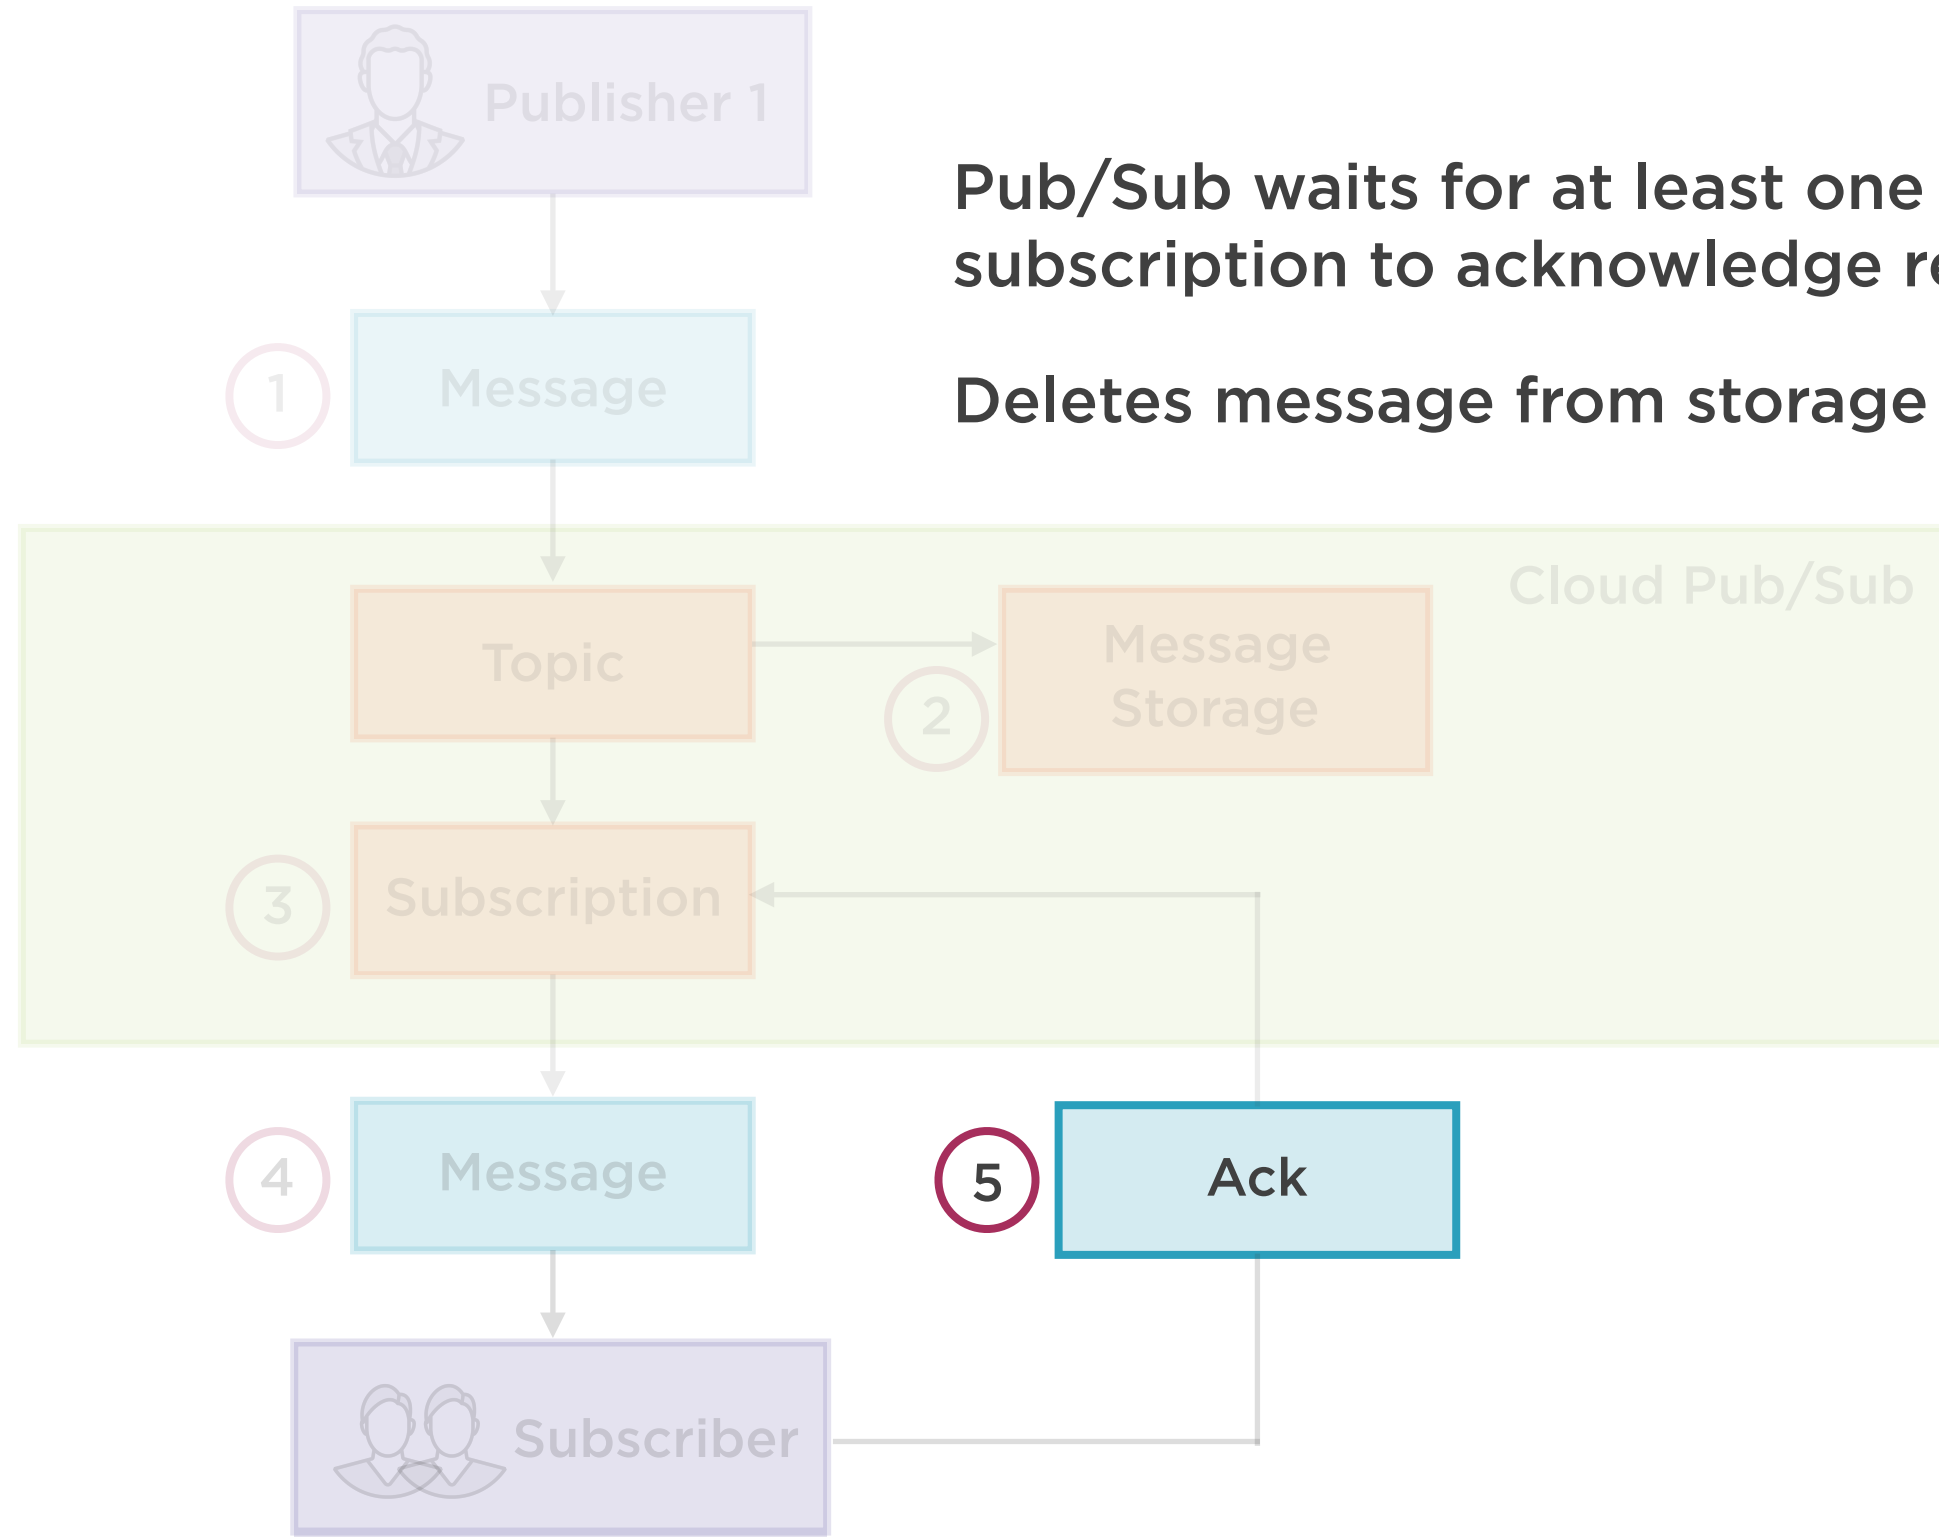
**Pub/Sub sends messages to every subscription**

# Pub/Sub Messaging



Subscribers acknowledge receipt to Pub/Sub

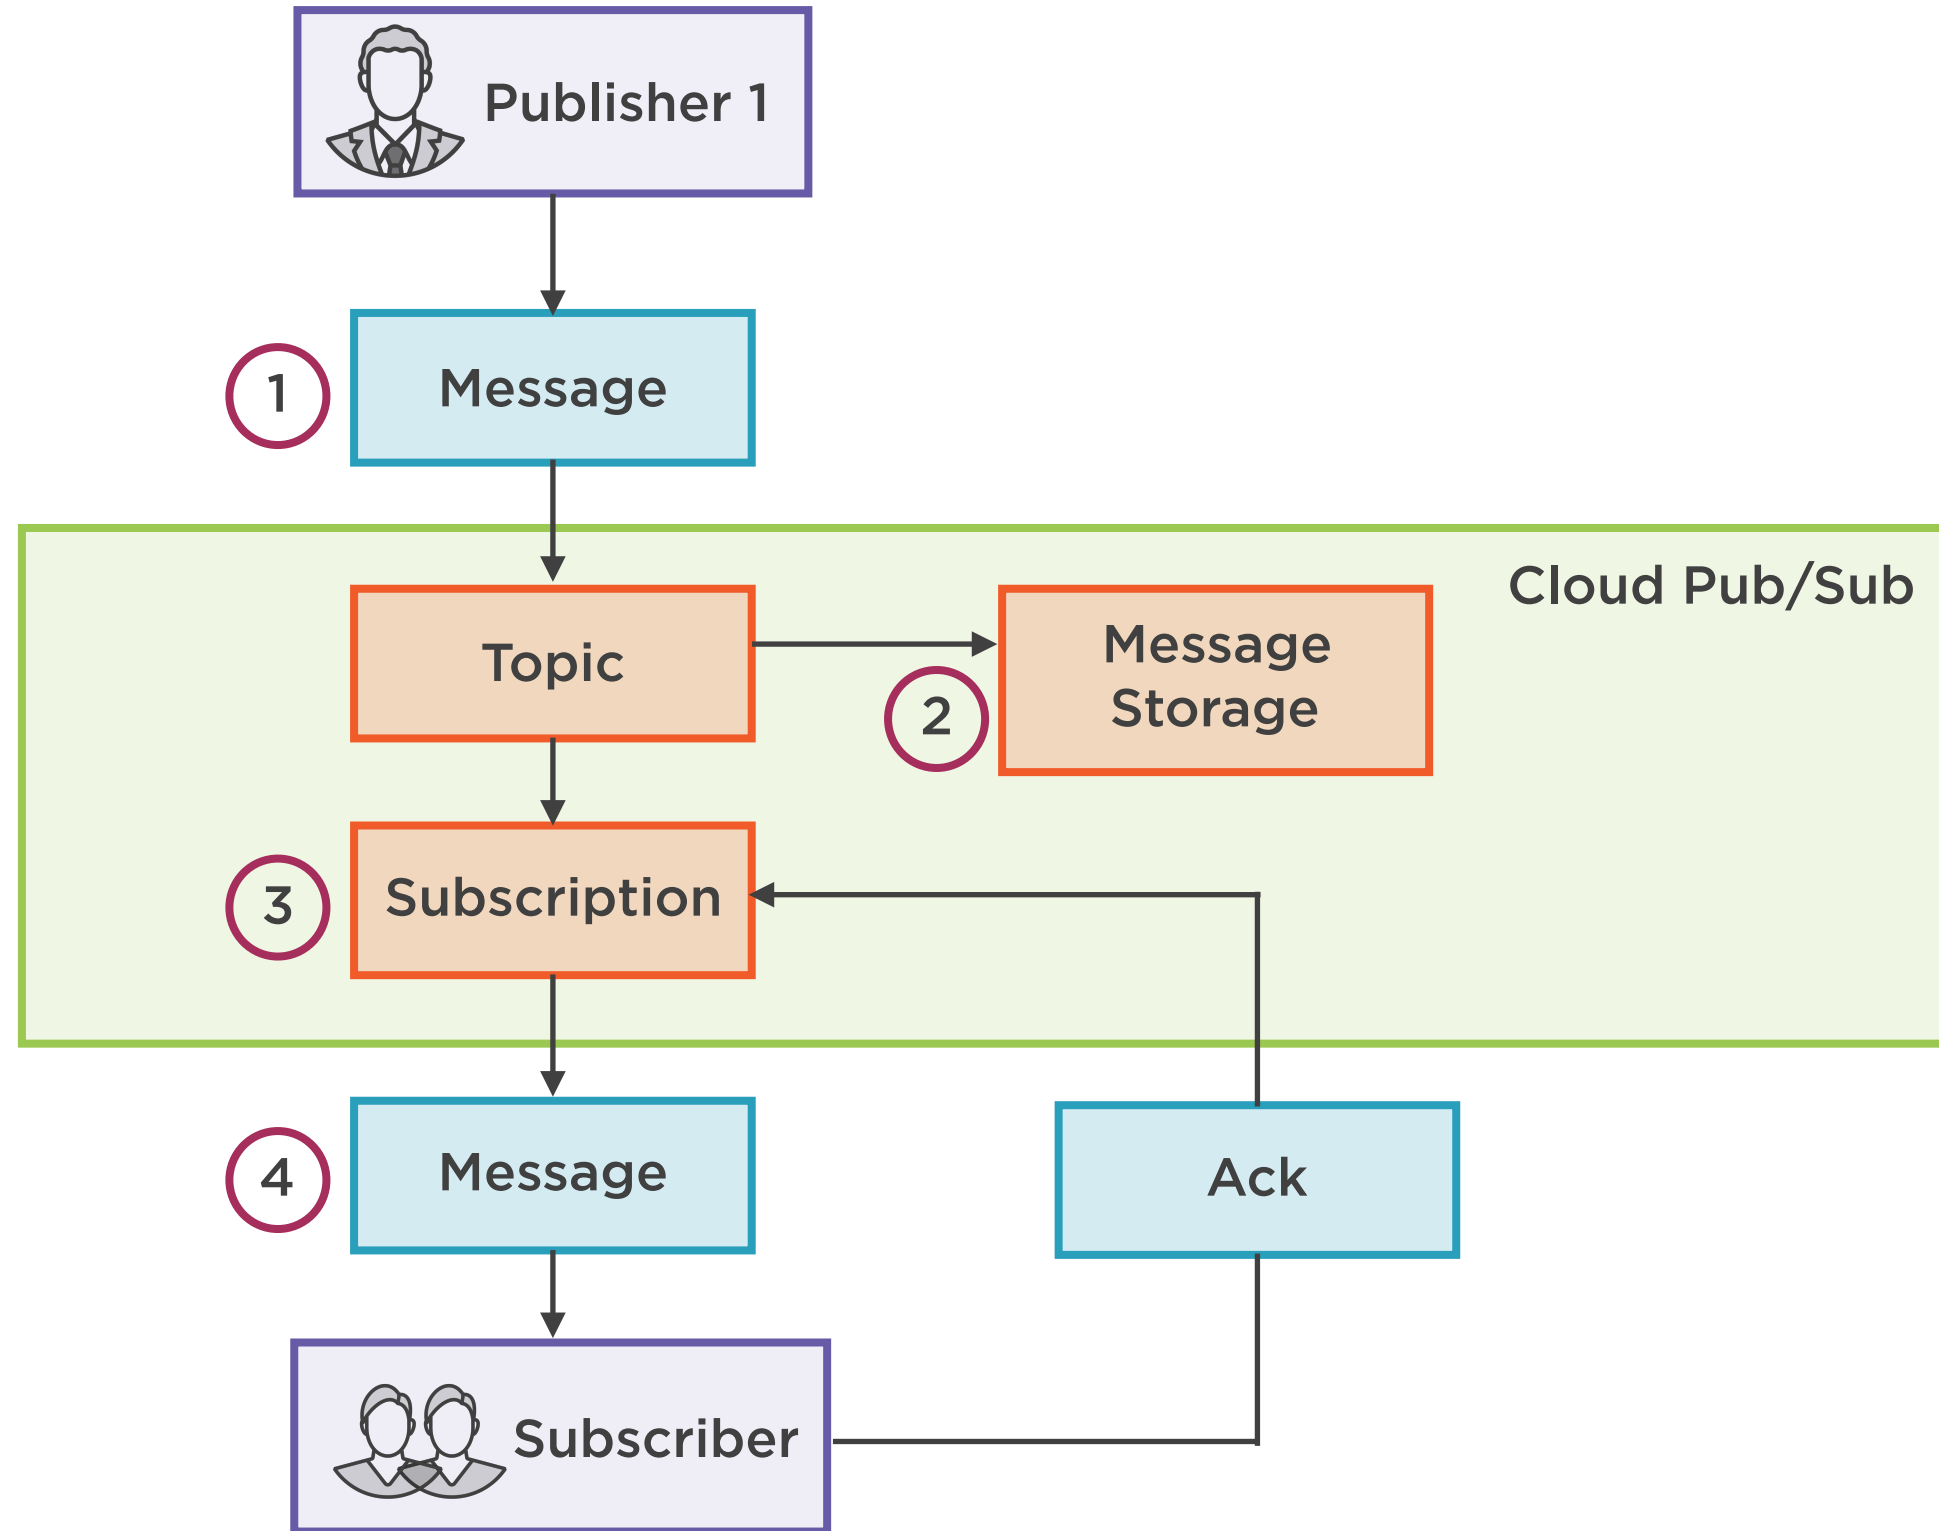
# Pub/Sub Messaging



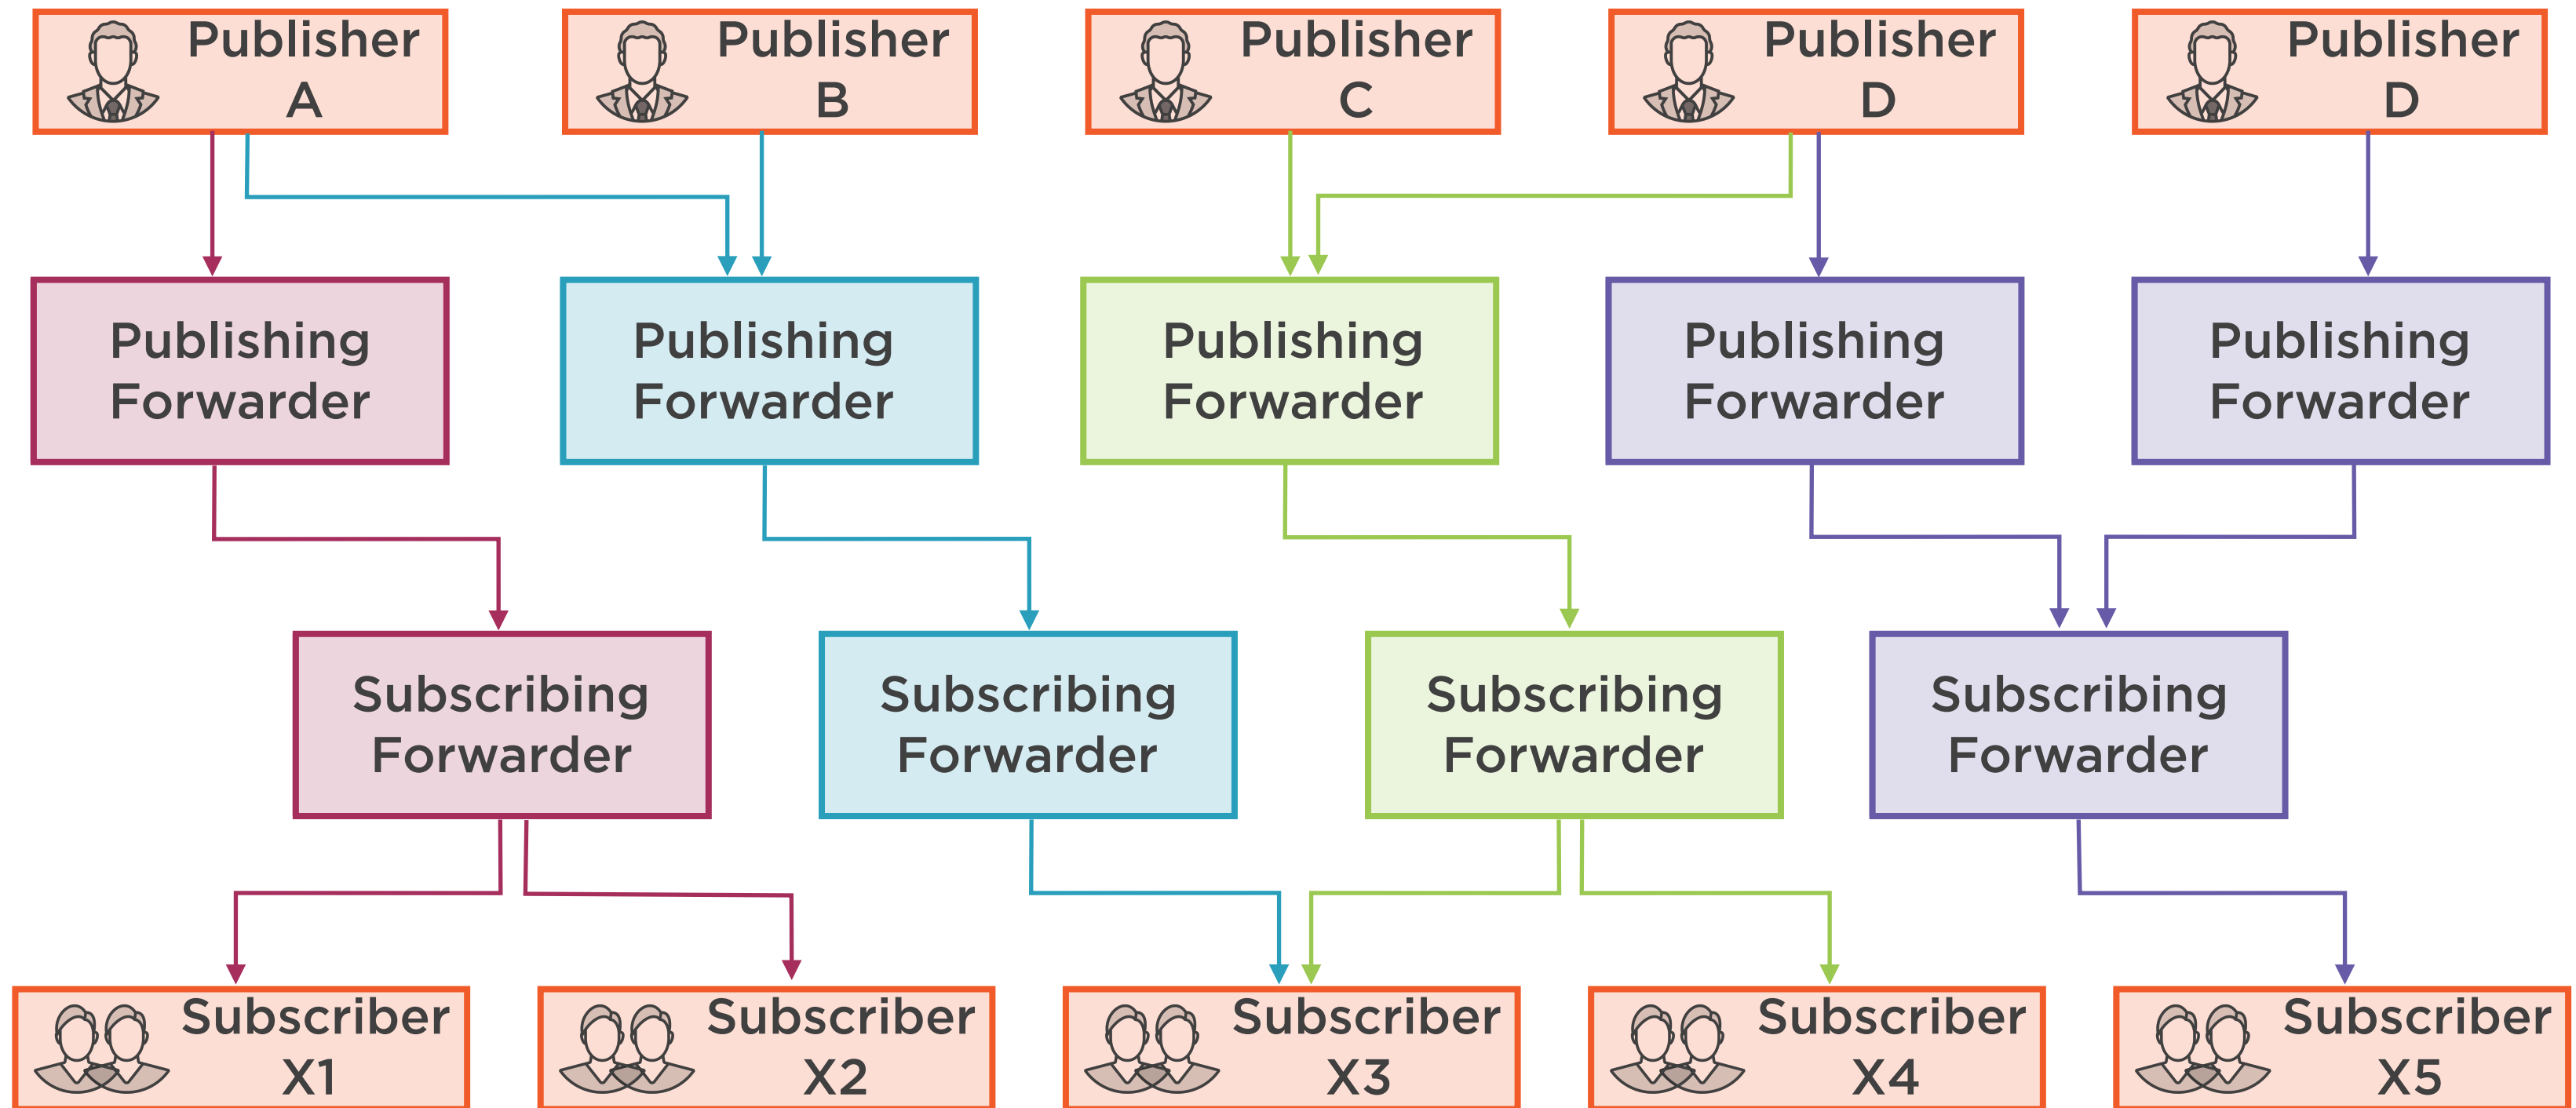
**Pub/Sub waits for at least one subscriber on each subscription to acknowledge receipt**

**Deletes message from storage**

# Pub/Sub Messaging



# Pub/Sub Messaging



# Architecture



Data plane

Control plane

# Data Plane



**Data plane** moves messages between publishers and subscribers

- Servers in data plane are called **forwarders**
- Publishing and subscribing forwarders

# Control Plane



**Control plane** assigns publishers and subscribers to servers on the data plane

- Servers in control plane are called **routers**



# Replaying Messages

---

# Seek

Extends subscriber functionality by allowing you to alter the acknowledgement state of messages in bulk.

# Seek to a Timestamp




Marks every message received **before the timestamp** as acknowledged

Marks every message received **after the timestamp** as unacknowledged

To seek to timestamp in past, must enable **retain acknowledged messages**

To discard messages seek to a time in the future

# Seek to a Snapshot




Captures the message acknowledgement state of a subscription

Retains **all unacknowledged messages** in the subscription at the time of snapshot

Also retains **all messages received after the snapshot was created**

# Use Cases



Update subscriber code safely

Recover from unexpected subscriber problems

Bulk acknowledge unneeded messages to avoid unnecessary processing

Testing subscribers on known data

# Pricing

---

# Pricing



**Message ingestion**

**Delivery**

**Seek-related storage**

- Snapshots
- Retained acknowledged messages

<https://cloud.google.com/pubsub/pricing>

# Pricing Schedule

Monthly data volume	Price
First 10GB	Free
Next 50TiB	\$60/TiB
Next 100TiB	\$50/TiB
Beyond 150TiB	\$40/TiB



**1024 KB messages + 2 subscriptions at rate of 1 MiB/second**

---

## Pricing Scenario

**Message ingestion and delivery**

**Ingesting 1 MiB/second + Delivery 2MiB/second = 3MiB/second**

---

Per-second Usage

**Message ingestion and delivery**

$$3 \text{ MiB/s} \times 3600 \times 24 \times 30 = 7.416 \text{ TiB/Month}$$

---

Per-month Usage

**Message ingestion and delivery**

**7.416 TiB - 10 GB (free tier) = 7.406 TiB**

---

Free Tier

**Deduct from data usage**

$$7.406 \text{ TiB} * \$60/\text{TiB} = \$444.36/\text{month}$$

---

Monthly Pub/Sub Bill

**Overall costs of operation**

# Fine Print



**Minimum billable volume per request is 1 KB**

## **Cross-project billing**

- Bills payable by project hosting resource
- Subscriptions, topics

# Seek-related Message Storage



**A subscription can be configured to retain acknowledged messages**

**Makes it possible to re-process messages using seek**

# Seek-related Message Storage



**A snapshot of a subscription is created**

**Storage fees are charged for storing all snapshot's unacknowledged messages**

**Message storage fees, at a rate of \$0.27/  
GiB-month**



**Steady 1 MiB/second incoming data with negligible backlog and 7 day retention**

---

Pricing Scenario: Retained Acknowledgements

$$1\text{MiB/second} \times 3600 \text{ seconds/hour} \times 24 \text{ hours/day} = 86.4 \text{ GiB/day}$$

---

## Pricing Scenario

**Daily data usage**

$$86.4 \text{ GiB/day} * 7 \text{ days} = 605 \text{ GiB}$$

---

## Pricing Scenario

**Weekly data usage**

**605 GiB-month \* \$0.27/GiB-month = \$163**

---

Retained Acknowledged Messages

**Total monthly bill incurred for retained messages**

# Snapshot Fees



When snapshot is created, one-time fee applied to backlog of unacknowledged messages

Fee equivalent to storing that backlog for the full seven days is charged

In addition to incremental charge for new messages

**60 MiB published in first minute stored for 7 days - 1 minute**

**60 MiB published in second minute stored for 7 days - 2 minutes**

**...**

**Over 7 days, backlog size reaches 605GiB**

---

## Pricing Scenario: Snapshots

**Backlog size reaches 605 GiB when snapshot expires after 7 days**

**Average size of backlog at any point in time is  $1/2 \times 605$  GiB**

**$1/2 \times 605$  GiB  $\times$  7 days = 2118 GiB-days**

---

Pricing Scenario: Snapshots

**Weekly backlog size calculation**

**2118 GiB-days x (1/30 months/day) x \$0.27/GiB-month = \$19**

---

## Snapshot Message Storage Fees

**Snapshots potentially cheaper than acknowledged message retention**



# Fees for Unacknowledged Messages



When snapshot is created, one-time fee applied to backlog of unacknowledged messages

Fee equivalent to storing that backlog for the full seven days is charged

In addition to incremental charge for new messages

**10 GiB unacknowledged message backlog**

---

Snapshot message storage fees

**Assumption for pricing scenario**

$$10 \text{ GiB} \times 7 \text{ days} / 30 \text{ days/month} \times 0.27/\text{GiB-month} = \$0.63$$

---

One-time Fee for Unacknowledged Backlog

**Total monthly fee**

# Summary

**Messaging middleware**

**Decouples senders and receivers**

**Global, replicated, autoscaling**

**Publishers publish to topics**

**Subscribers listen on subscriptions**

**Complex pricing**