

# Axon Training

Module 5 – Snapshotting and Event Processors

# Agenda

## Week 1

1. DDD and CQRS Fundamentals
2. Command Model
3. Event Handling & Projections
4. Sagas and Deadlines


## Week 2

1. **Snapshotting and Event Processors**
2. Preparing for Production
3. CQRS and Distributed Systems
4. Monitoring, Tracing, Advanced Tuning

Compacting the event stream

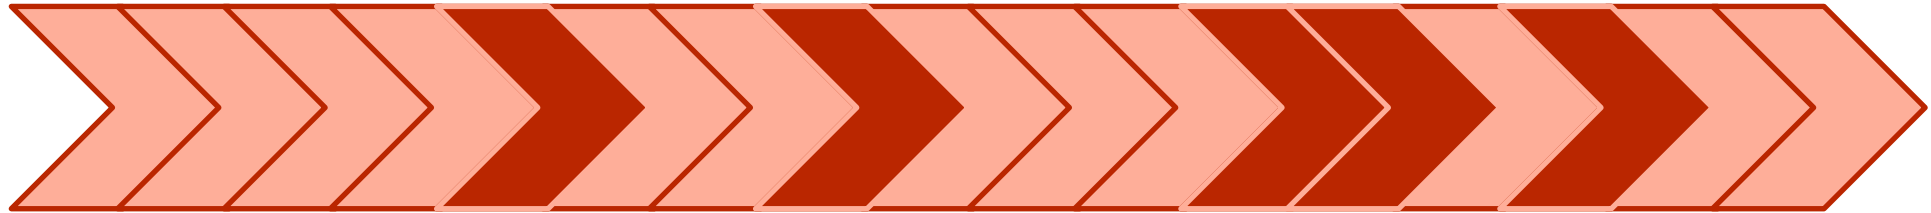
# Snapshotting

# Snapshotting




# Event Store operations

- Read aggregate's events



# Snapshotting

Snapshots



Event Store



# Snapshotting

- Snapshots are a (temporary) replacement for a set of historical events
- Snapshotting may be an asynchronous process
  - Regular intervals
  - After x events
  - When loading takes  $\geq x$  ms

# Snapshot Event

- Snapshot Events contain all relevant information needed to reconstruct an Aggregate's state at that point in time.
- Axon's default: Use the actual Aggregate's state as snapshot
- Note: Snapshot Events are *never* published



# Configuring Snapshots

// in the configuration of the aggregate

```
AggregateConfigurer<Flight> flightConfig = AggregateConfigurer.defaultConfiguration(Flight.class);
```

// we define the trigger

```
flightConfig.configureSnapshotTrigger(  
    c -> new EventCountSnapshotTriggerDefinition(c.getComponent(Snapshotter.class), 100)  
);
```

// in our main configuration, we provide an AggregateSnapshotter that we can reuse

```
Configurer config = DefaultConfigurer.defaultConfiguration()  
    .registerComponent(Snapshotter.class, c -> AggregateSnapshotter.builder()  
        .aggregateFactories(new GenericAggregateFactory<>(Flight.class))  
        .eventStore(c.eventStore())  
        .transactionManager(c.getComponent(TransactionManager.class))  
        .build());
```

# Configuring Snapshots - Spring

// in the configuration of the aggregate

@Aggregate(snapshotTriggerDefinition = "myTriggerDefinition")

public class Flight {

...

}

// we define the trigger. The snapshotter is automatically configured

@Bean

public SnapshotTriggerDefinition myTriggerDefinition(Snapshotter snapshotter) {

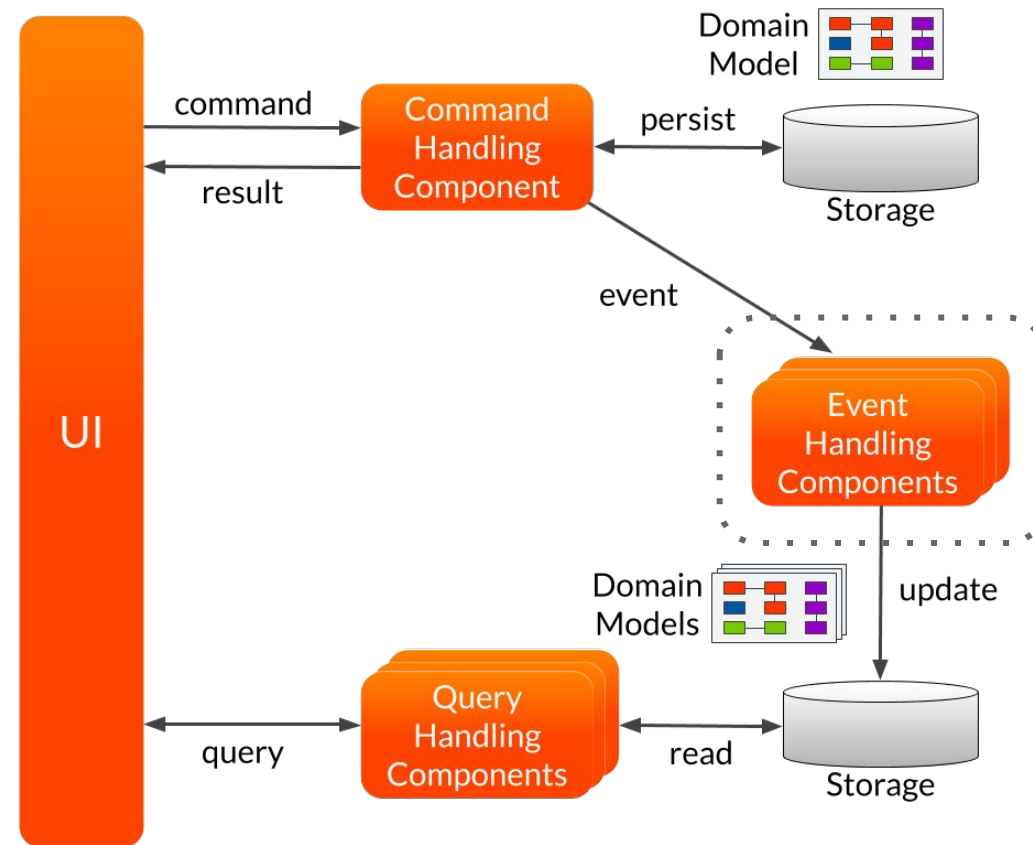
return new EventCountSnapshotTriggerDefinition(snapshotter, 100);

}

Processing what happened...

# Event Processors & Replays

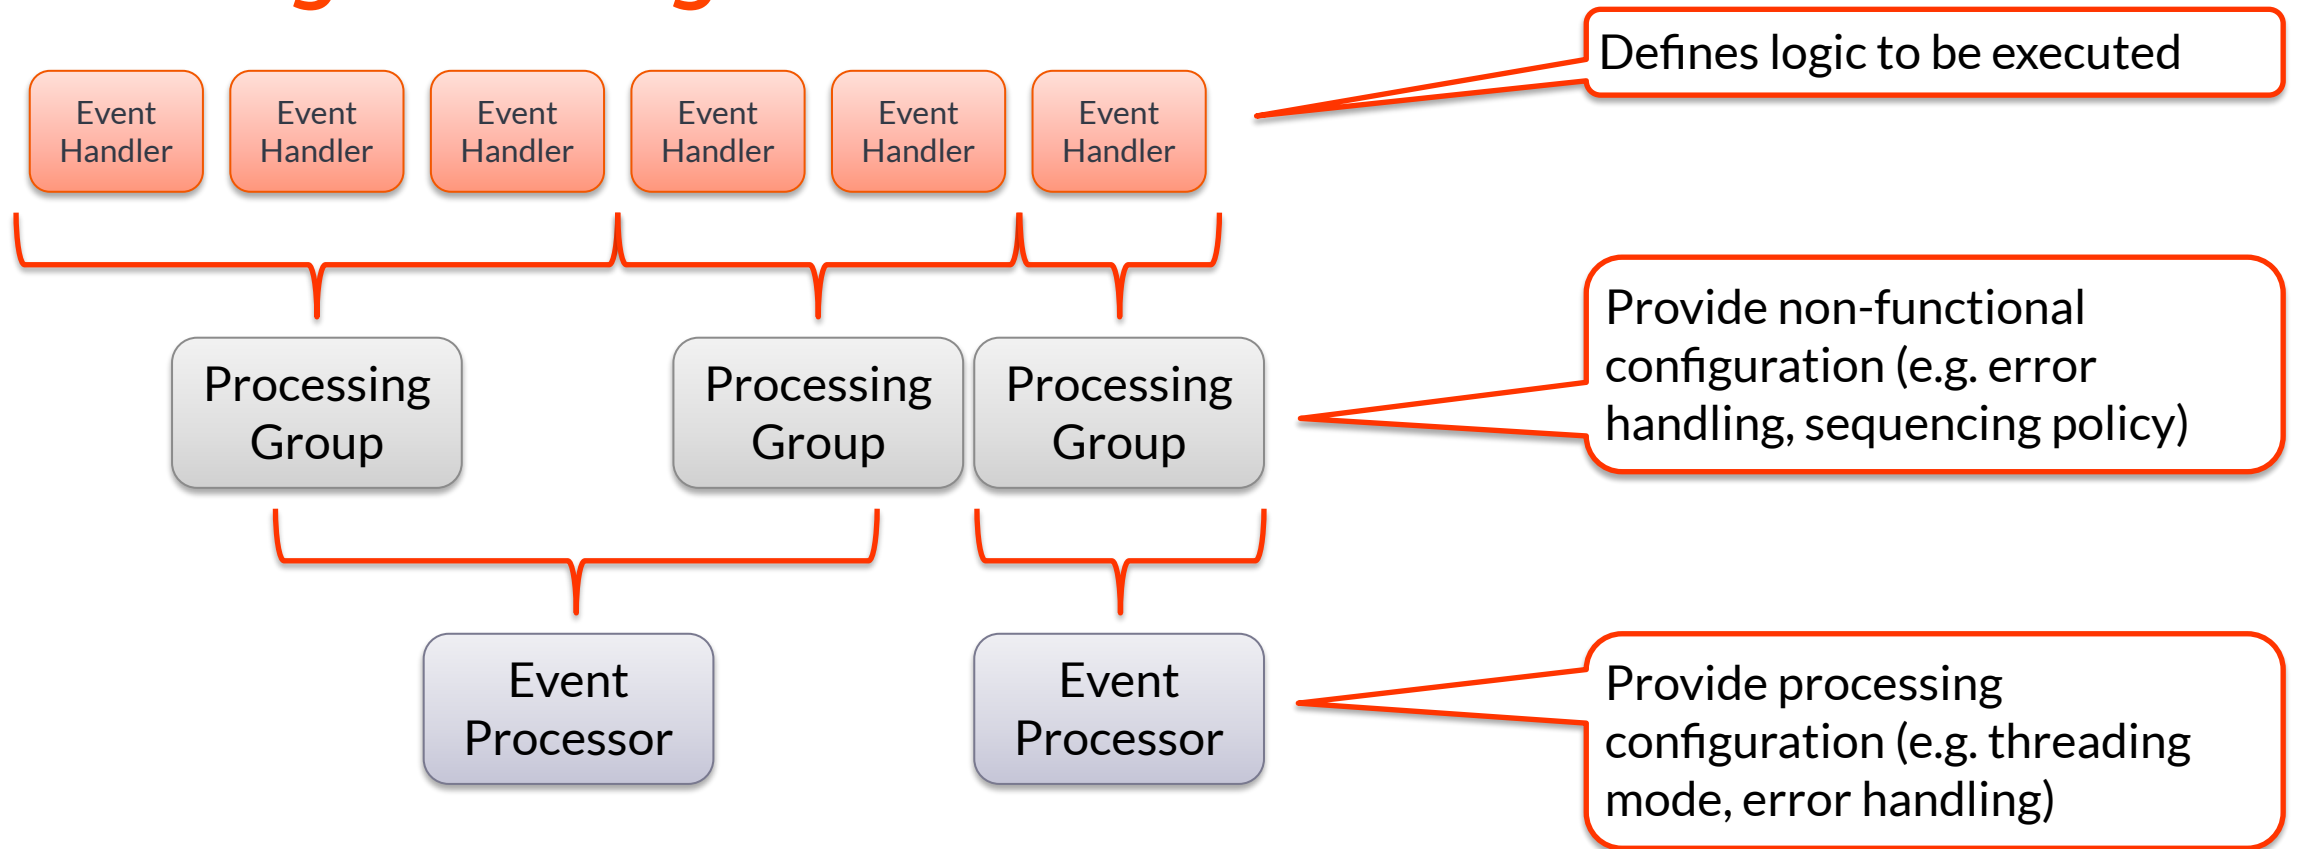
# Event Processing



# Organizing Event Handlers

- Event Processor
  - Responsible for managing the technical aspect of processing an Event
  - Starts and Commits 'Unit of Work'
  - Invokes handler methods
- Each handler is assigned to a single Processor
  - `@ProcessingGroup` on Event Handler class
  - Assignment rules in `EventProcessingConfigurer` (part of Configuration API)


# Organizing Event Handlers



# Event Processors

- SubscribingEventProcessor
  - Receives messages as they are published, in the thread that publishes the messages
  - Requires a Subscribable Message Source
- TrackingEventProcessor (default \*)
  - Uses its own thread(s) to read EventMessages from a Stream
  - Requires a Streamable Message Source
  - Saves progress using TrackingToken

# Tracking Token





# Event Processor Configuration

```
public void configure(EventProcessingConfigurer configurer) {  
    StreamableMessageSource<TrackedEventMessage<?>> source = ...;  
    configurer.registerTrackingEventProcessor("com.example.viewmodel", c -> source);  
}
```

The name of the processor to explicitly register.  
It is only created when handlers are actually assigned to it.

A function returning the source to  
read from, given Configuration c.

Or in Spring Boot with `application.properties`:

```
axon.eventhandling.processors.processor-name.source=message-source-bean-name  
axon.eventhandling.processors.processor-name.mode=tracking
```

# Tracking Event Processor Configuration

- **Batch Size** - The number of events that are processed in a single transaction
- **Initial Token** - The position at which a processor must start when initializing
- **Initial Segment Count** - The number of segments to create when initializing
- **Thread Count** – The maximum number of Threads the processor may start
- **Event Availability Timeout** – Time to wait for events before updating the claim
- **Token Claim Interval** – How long to wait between attempts to claim a segment

# Error Handling

- Exceptions thrown while handling an Event
  - `ListenerInvocationErrorHandler`
  - Defined on Processing Group
  - Default: log error and proceed
  - Rethrow to trigger ErrorHandler
- Exceptions that fail the transaction
  - `ErrorHandler`
  - Defined on Event Processor
  - Default: rollback, release segment claim, and retry (with incremental back-off)

# Thread Count and Segmentation

- A segment can only be processed by a single thread at a time
- A single thread will process a single segment
- At any time: total thread count  $\geq$  segment count
  - Otherwise: unclaimed segment / partial processing
- Initial Segment Count *only* works when initializing a processor
- At runtime, use Split and Merge to increase/decrease segment count

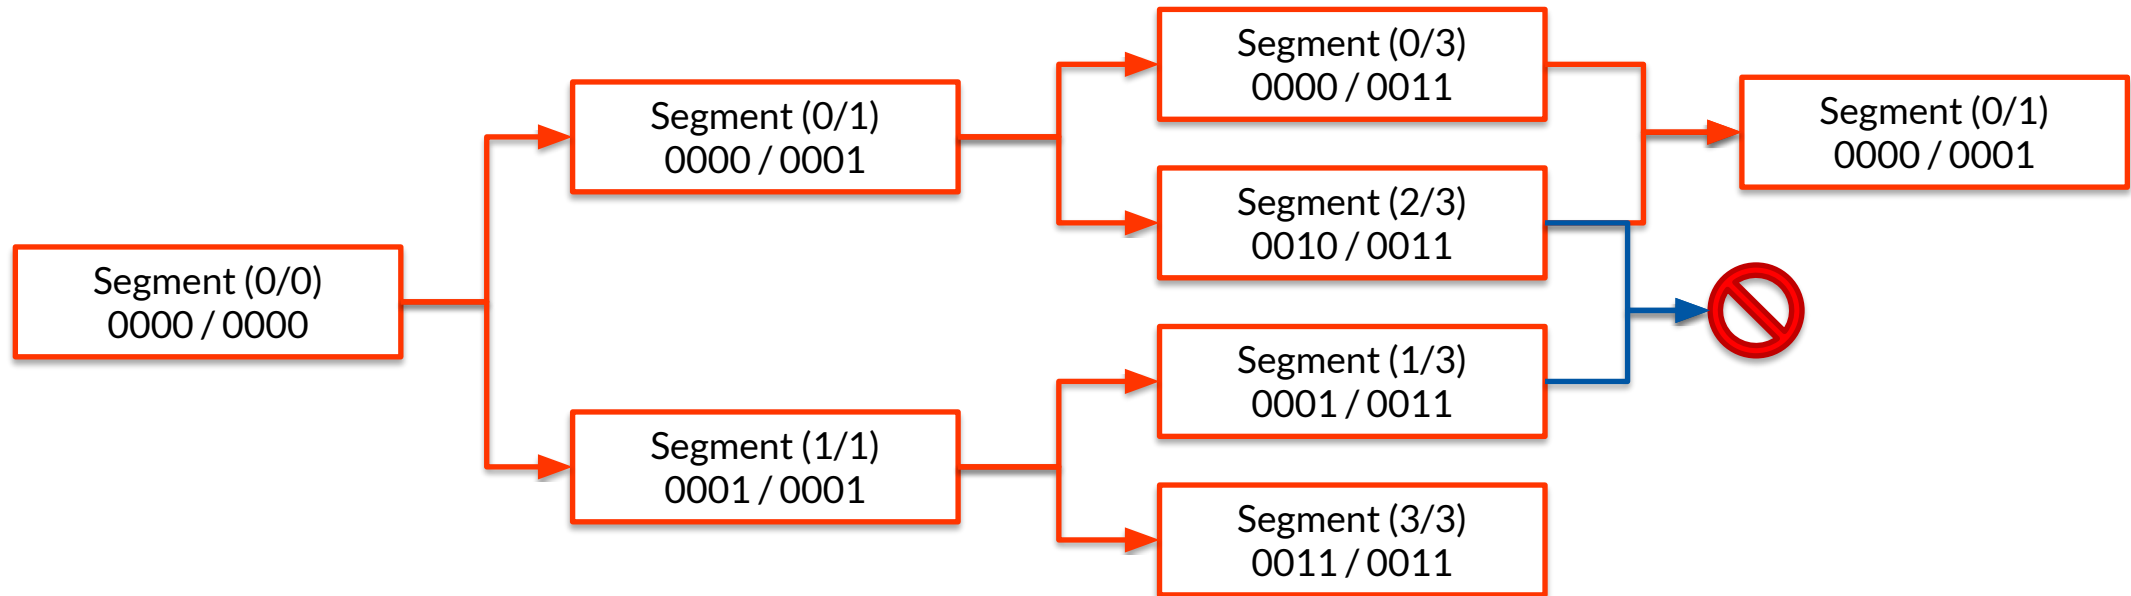
# Tracking Processor Segments

- Multi-threading and/or multi-node
- Each thread “claims” a segment in TokenStore
- SequencingPolicy defines segment
  - the same value for two messages means they ‘belong’ to same segment
  - Message in same segment are always handles sequentially
  - E.g. SequentialPerAggregatePolicy



# Segmentation – Split and Merge

- Segmentation is dynamic
  - Split splits a claimed segment into 2 segments
  - Merge merges two segments into their original combined form



# Replays

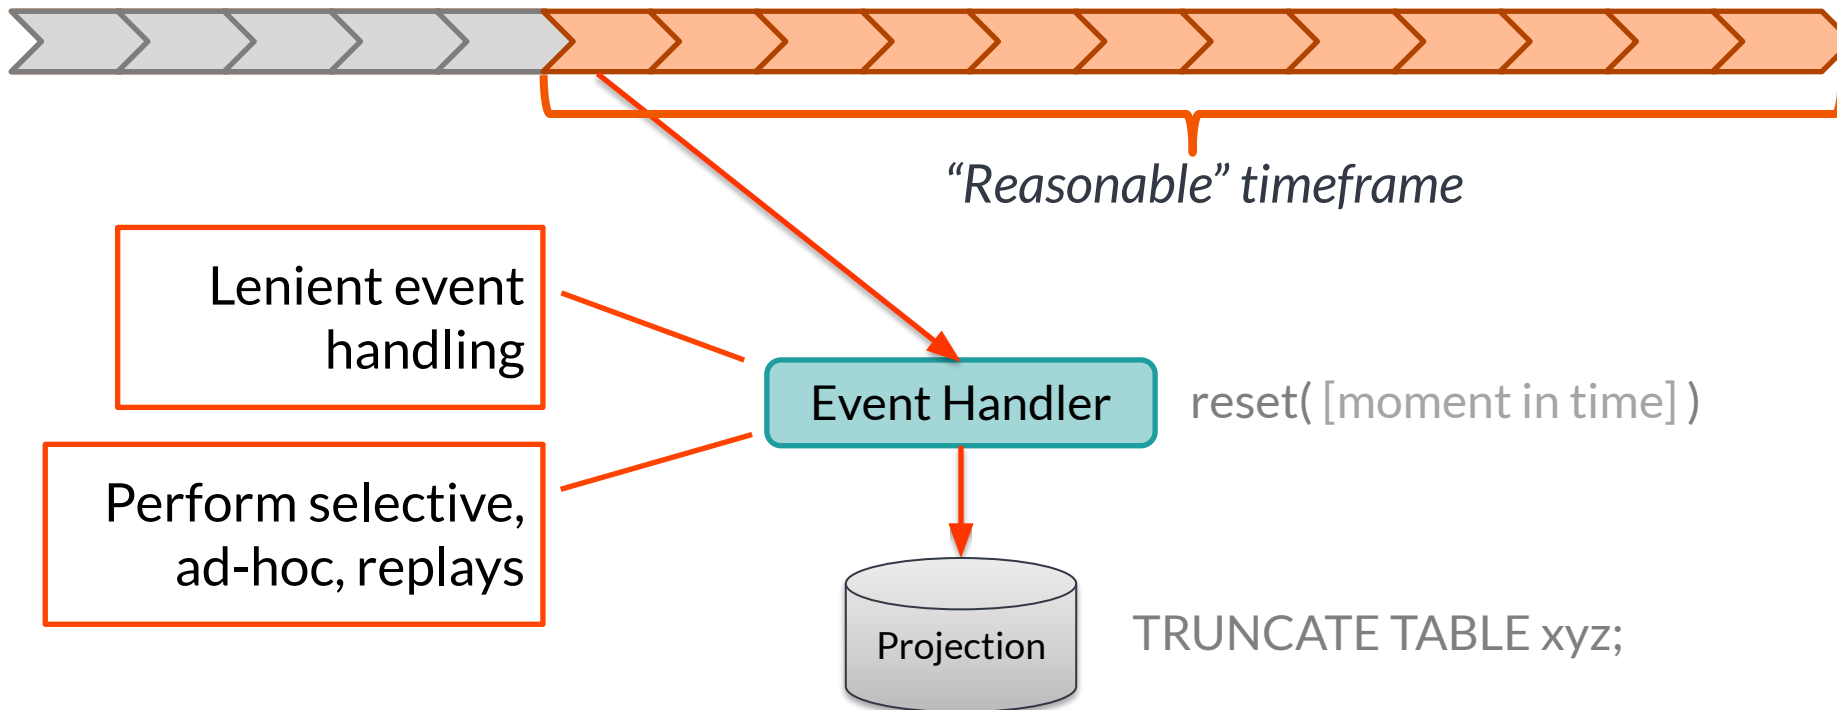
- Tracking Processors can be “reset”
  - Clean up any state their handlers have
  - Reset all tokens for that processor
- Tracking Processor replay status
  - `@AllowReplay` / `@DisallowReplay` – *indicates whether components can deal with replays*
  - `ReplayStatus` – *handler parameter to add conditional logic*
  - `@ResetHandler` – *handler invoked when a replay is triggered*

# Replays - API

```
class FlightStatusProjection {  
    @DisallowReplay  
    @EventHandler  
    public void on(FlightDelayed event) {  
        // This handler is not invoked when replaying  
    }  
  
    @EventHandler  
    public void on(ArrivalTimeChanged event, ReplayStatus replayStatus) {  
        if (replayStatus != ReplayStatus.REPLAY) {  
            // This block is not invoked when replaying  
        }  
        // ...  
    }  
  
    @ResetHandler  
    public void reset() {  
        // Invoked when replay is triggered  
        // e.g. to clear out the view's database  
    }  
}
```



# Partial replays



# Triggering a reset

A reset requires that a single processor can update ***all*** tokens simultaneously. This is only possible when the processor is stopped

1. Stop all processors
2. Ask a single processor to “reset” all tokens
3. Start all processors

In a distributed environment, the AxonServer API/UI can be used to stop all instances of a processor. AxonServer ***does not*** trigger resets.

# Triggering a reset - API

```
// Get access to the processors with the configuration
EventProcessingConfiguration eventProcessingConfiguration = ...;

// and if you know the name of the processor to reset,
String processorName = ...;

// then you can execute a reset
eventProcessingConfiguration.eventProcessor(processorName, TrackingEventProcessor.class)
    .ifPresent(trackingEventProcessor -> {
        trackingEventProcessor.shutdown();
        trackingEventProcessor.resetTokens();
        trackingEventProcessor.start();
    });
```

Whatever else you wanted to know...

# Questions